



Building resilient systems inside the mesh: abstraction and automation of Virtual Service generation

Vladimir Georgiev, Thought Machine



#IstioCon

Sync calls failures inside the mesh

- Everyone says to fail fast and retry quickly, but...
- How fast to timeout?
 - If it's too early then failed the request for no reason.
 - If it's too late then the calling client might be left hanging for too long.
- What errors are retryable?
- Who knows the answer to all the questions?
- How to implement this to be language agnostic?



Virtual Services API

- Solves our problems, but...
- All Service Owners must be aware of the Virtual Services API in order to define their SLOs.
- Potential typing errors when dealing with YAMLS.
- Potential drift between the state of the service API and the Virtual Service config.
- Hard to manage when having hundreds of services.



Abstracting to proto files

```
1 syntax = "proto3";
2
3 package sla;
4 import "google/protobuf/descriptor.proto";
5
6 extend google.protobuf.MethodOptions {
7     // Describes the anticipated SLOs for the RPC and can override the global
8     // Virtual Service Config. See SLO for more description.
9     SLO endpoint_slo = 1;
10 }
11
12 message EndpointSLO {
13     // Defines the global route timeout.
14     string timeout = 1;
15     // Defines the upstream timeout per attempt.
16     string per_try_timeout = 2;
17     // Defines retry attempts amount.
18     int32 retries = 3;
19     // Defines the retryable gRPC errors.
20     repeated RetryOn retry_on = 4;
21 }
22
23 enum RetryOn {
24     CANCELLED = 0;
25     DEADLINE_EXCEEDED = 1;
26     INTERNAL = 2;
27     RESOURCE_EXHAUSTED = 3;
28     UNAVAILABLE = 4;
29 }
```

Annotations API definition

#IstioCon

```
1 syntax = "proto3";
2
3 package hello_world;
4
5 import "demo/virtual_service/sla.proto";
6
7 service Greeting {
8     rpc SayHi(GreetingRequest) returns (GreetingResponse) {
9         option (virtual_service.endpoint_slo) = {
10             timeout: "200ms"
11             per_try_timeout: "150ms"
12             retries: 5
13         };
14     };
15 };
16
17 message GreetingRequest {
18     string msg = 1;
19 }
20
21 message GreetingResponse {
22     string msg = 1;
23 }
```

Greeting service example



Please Build System

- <https://github.com/thought-machine/please>
- Uses BUILD and allows for creation of miscellaneous rules

```
1 def k8s_manifest(  
2     name:str,  
3     srcs:list,  
4     grpc_services:list=None,  
5 ):  
6     """Defines Kubernetes manifests."""  
7     Args:  
8         name: Name of the rule.  
9         srcs: YAML config files containing the Kubernetes config.  
10        grpc_services: The proto files used for creating VirtualService for  
11                       the defined gRPC services.  
12  
13    virtual_service_rule = genrule(  
14        name = f'{name}_virtual_service',  
15        srcs = [grpc_services + srcs],  
16        requires = ['proto', 'yaml'],  
17        cmd = '$TOOLS --yaml $${SRCS_YAMLS}" --protos "$${SRCS_PROTOS}"'  
18    )
```

Misc please rule for autogeneration

```
1 subinclude("//demo/k8s/defs:k8s_manifest")  
2  
3 k8s_manifest(  
4     name = "greeter_k8s",  
5     srcs = [  
6         "greeter.yaml", # K8s Deployment yaml  
7         "greeter-svc.yaml", # K8s Service yaml  
8     ],  
9     grpc_services = [  
10        "//demo/proto:greeter",  
11    ],  
12 )
```

K8s Greeter service example



Building the new rule

```
→ src git:(demo) (* microk8s:default) plz build //experimental/vlad/demo/k8s:greeter_k8s_virtual_service 2>/dev/null
```



Deploying to a cluster

```
→ src git:(demo) (* microk8s:default) plz sef deploy-service //experimental/vlad/demo/k8s:gre  
eter_k8s 2> /dev/null
```



- Easy way to manage Virtual Service configs.
- Virtual Service configs become a release artifact.
- Easy abstraction for defining timeouts and retries in a language agnostic way.
- Application developers using Istio/Envoy for retries and timeouts without knowing it.



Thank you!

vlad@thoughtmachine.net

[#IstioCon](#)

