



Descartes
Labs

Understanding The Data Plane

What Envoy Hears When Istio Speaks




you are here



istio nirvana

(not to scale)



Rob Salmond

SRE Here

Istio in Prod for >2yrs

 @phro

Map of These Territories

Istiod

WTF is an Envoy

Envoy Operation

Envoy
Configuration



istiod

An Envoy “management server”

You had ~~one~~ three jobs!

Galley

Understand
Kubernetes

Pilot

Program
Sidecars

Citadel

Autograph
Certificates

You had ~~one~~ ~~three~~ many jobs!

Galley

Understand
Kubernetes

Pilot

Program
Sidecars

Citadel

Autograph
Certificates

sidecar injection
webhook handler!

cinnamon flavor
crystals!

CRD validation!

Several jobs are typing ...

Galley

Understand
Kubernetes



Pilot

Program
Sidecars

Citadel




SERVICE
DISCOVERY




ENVOY
CONFIGURATION


istiod - galley




istiod - galley




istiod - galley



istiod - galley



istiod - galley



istiod - galley

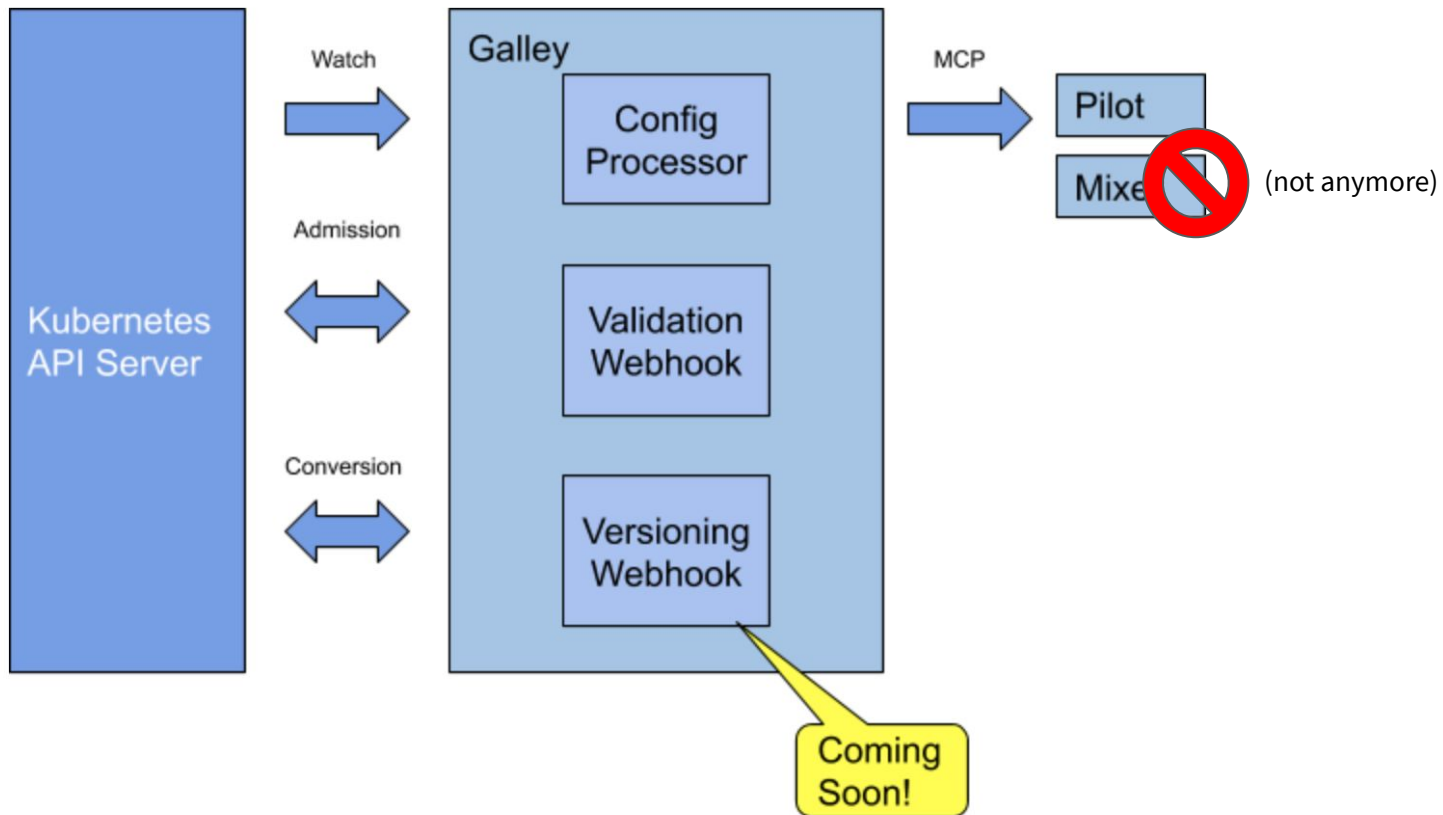
```
rob@ayyyyyyylmao:~$ kubectl port-forward po/istiod-58f84ffddc-r8nqf 8080
Forwarding from 127.0.0.1:8080 -> 8080
Forwarding from [::1]:8080 -> 8080
```

istiod - galley

```
← → ↺ localhost:8080/debug/registryz

[
  {
    "Attributes": {
      "ServiceRegistry": "Kubernetes",
      "Name": "api",
      "Namespace": "master",
      "UID": "istio://master/services/api",
      "ExportTo": null,
      "ClusterExternalAddresses": null,
      "ClusterExternalPorts": {
        "Kubernetes": {
          "8000": 30731
        }
      }
    },
    "ports": [
      {
        "name": "grpc-web-port",
        "port": 8000,
        "protocol": "GRPC-Web"
      }
    ],
    "creationTime": "2021-01-28T20:46:33Z",
    "hostname": "api.master.svc.cluster.local",
    "address": "10.56.3.75",
    "Mutex": {},
    "cluster-vips": {
      "Kubernetes": "10.56.3.75"
    },
    "Resolution": 0,
    "MeshExternal": false
  },
]
```

istiod - galley



istiod - galley - validation

```
apiVersion: admissionregistration.k8s.io/v1
kind: ValidatingWebhookConfiguration
metadata:
  name: istiod-istio-system
webhooks:
- admissionReviewVersions:
  - v1beta1
  clientConfig:
    service:
      name: istiod
      namespace: istio-system
      path: /validate
      port: 443
  failurePolicy: Fail
  matchPolicy: Exact
  name: validation.istio.io
  namespaceSelector: {}
  objectSelector: {}
  rules:
  - apiGroups:
    - config.istio.io
    - rbac.istio.io
    - security.istio.io
    - authentication.istio.io
    - networking.istio.io
    apiVersions:
    - '*'
    operations:
    - CREATE
    - UPDATE
    resources:
    - '*'
    scope: '*'
  sideEffects: None
  timeoutSeconds: 30
```

istiod - galley - sidecar injection

```
apiVersion: admissionregistration.k8s.io/v1
kind: MutatingWebhookConfiguration
metadata:
  name: istio-sidecar-injector
webhooks:
- admissionReviewVersions:
  - v1beta1
  clientConfig:
    service:
      name: istiod
      namespace: istio-system
      path: /inject
      port: 443
  failurePolicy: Fail
  matchPolicy: Exact
  name: sidecar-injector.istio.io
  namespaceSelector:
    matchLabels:
      istio-injection: enabled
  objectSelector: {}
  reinvocationPolicy: Never
  rules:
  - apiGroups:
    - ""
    apiVersions:
    - v1
    operations:
    - CREATE
    resources:
    - pods
    scope: "*"
  sideEffects: None
  timeoutSeconds: 30
```

sidecar injection sidebar

```
$ istioctl kube-inject -f mypod.yaml -o sidecarpod.yaml
```

You had ~~one~~ three jobs!

Galley



Understand
Kubernetes


Pilot

Program
Sidecars

Citadel



istiod - pilot



Mapping Resources to Envoy Configuration

Resource Type	Envoy Configuration	Notes
Kubernetes Services	Listeners Routes Clusters	New listeners if port/protocol combo is unique Add virtual hosts for existing routes One cluster per Service/Port/Subset
Kubernetes Endpoints	Endpoints	
Istio Gateway	Listeners	Apply to Ingress/Egress Gateways
Istio VirtualService	Listeners Routes	Client side proxies TLS/TCP affect listeners HTTP match blocks affect routes
Istio DestinationRule	Clusters Endpoints	Client side proxies Connection/HTTP/TLS settings
Istio ServiceEntry	Clusters Endpoints	Client side proxies
Istio PeerAuthentication	Listeners Clusters	Server side proxies
Istio RequestAuthentication	Listeners	Server side proxies
Istio Authorization Policies	Listeners	Server side proxies
Istio EnvoyFilter	All	Break glass API to directly manipulate Envoy
Istio Sidecar	All	Client or server side proxies Sidecar scope sets config visibility

You had ~~one~~ three jobs!

Galley



Understand
Kubernetes

Pilot ✓

Program
Sidecars

Citadel

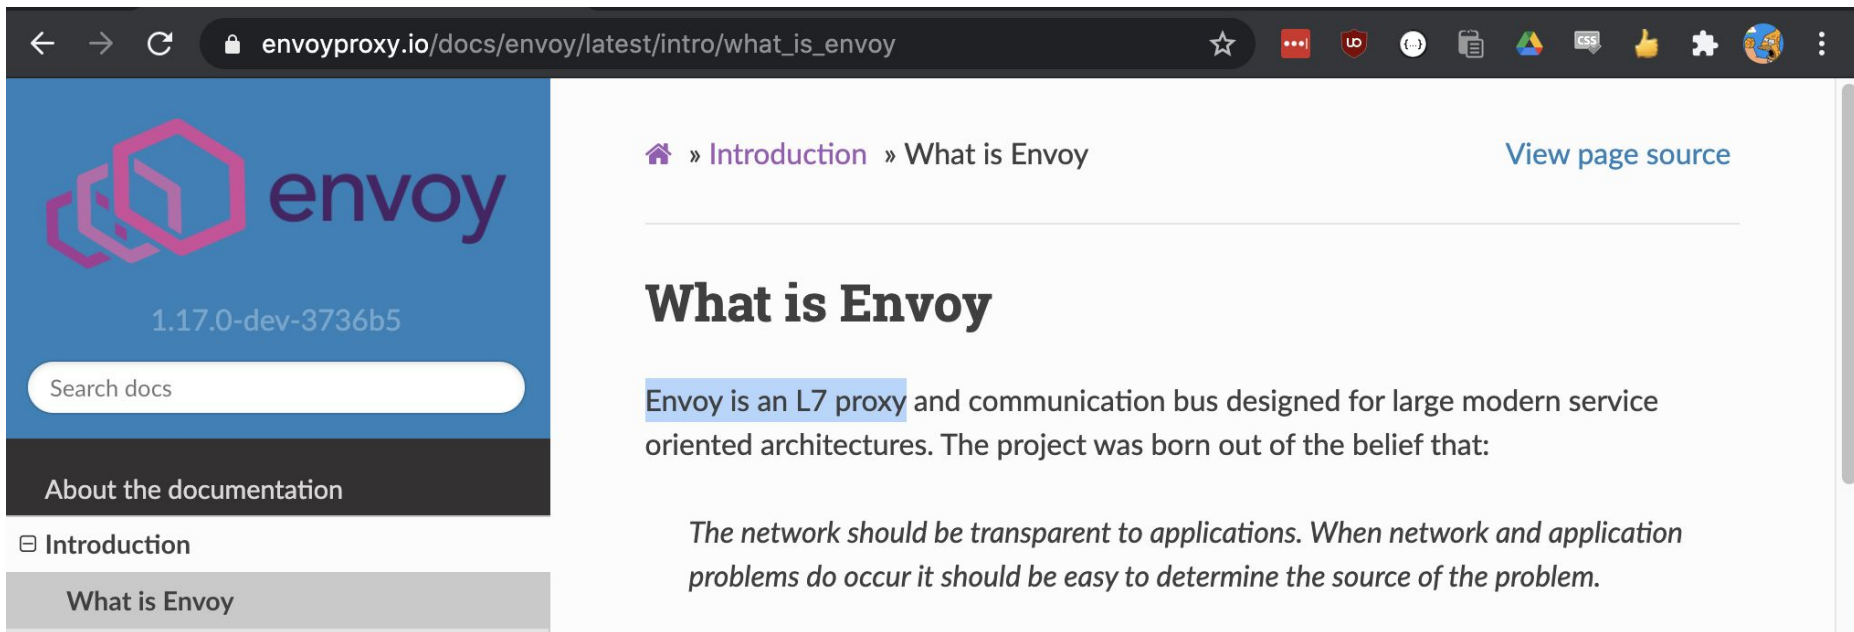


WTF Is An Envoy?

Isn't that the dude who pissed off king Leonidas and got kicked down a well?




envoy - was ist das?



The screenshot shows a web browser displaying the Envoy Proxy documentation. The browser's address bar shows the URL `envoyproxy.io/docs/envoy/latest/intro/what_is_envoy`. The page has a dark blue header with the Envoy logo and version `1.17.0-dev-3736b5`. A search bar is present. The left sidebar contains a navigation menu with 'Introduction' and 'What is Envoy' (the latter is highlighted). The main content area shows the breadcrumb '» Introduction » What is Envoy', a 'View page source' link, and the title 'What is Envoy'. The text describes Envoy as an L7 proxy and communication bus for modern service-oriented architectures, stating it was born from the belief that the network should be transparent to applications.

← → ↺ `envoyproxy.io/docs/envoy/latest/intro/what_is_envoy` ☆ ... LD (c) [Icons] ⌵

 **envoy**
1.17.0-dev-3736b5

Search docs

About the documentation

- ☐ Introduction
- What is Envoy**

» Introduction » What is Envoy [View page source](#)

What is Envoy

Envoy is an L7 proxy and communication bus designed for large modern service oriented architectures. The project was born out of the belief that:

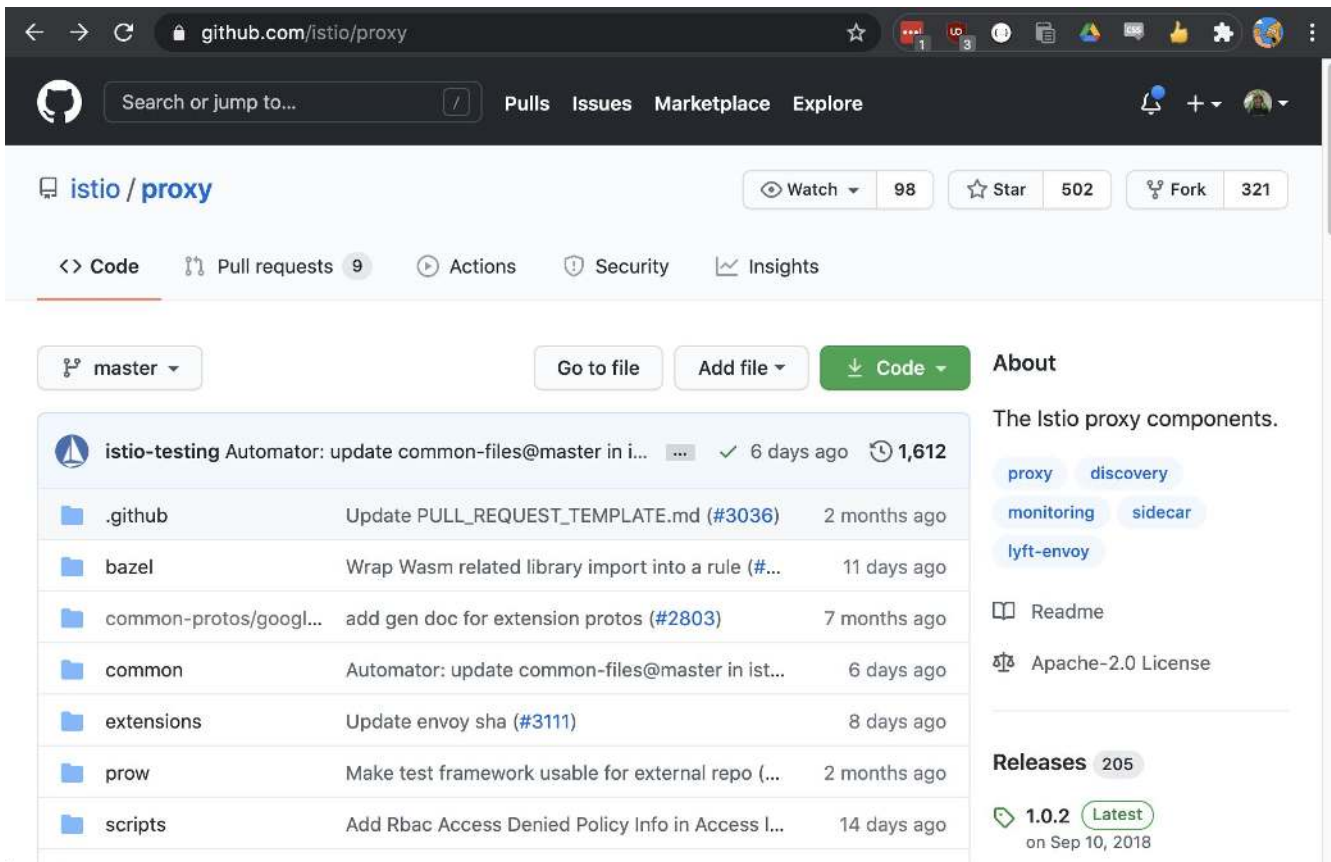
The network should be transparent to applications. When network and application problems do occur it should be easy to determine the source of the problem.

envoy - was ist das?

Envoy refresher

- **Out of process architecture:** Let's do a lot of really hard stuff in one place and allow application developers to focus on business logic.
- **Modern C++11 code base:** Fast and productive.
- **L3/L4 filter architecture:** A byte proxy at its core. Can be used for things other than HTTP (e.g., MongoDB, redis, stunnel replacement, TCP rate limiter, etc.).
- **HTTP L7 filter architecture:** Make it easy to plug in different functionality.
- **HTTP/2 first!** (Including **gRPC** and a nifty gRPC HTTP/1.1 bridge).
- **Service discovery** and **active/passive health checking**.
- **Advanced load balancing:** Retry, timeouts, circuit breaking, rate limiting, shadowing, outlier detection, etc.
- Best in class **observability:** stats, logging, and tracing.
- **Edge proxy:** routing and TLS.

envoy - was ist das?



The screenshot shows the GitHub repository page for `istio/proxy`. The repository has 98 watches, 502 stars, and 321 forks. The main content area displays a list of recent commits, including updates to the pull request template, Wasm library imports, extension protocols, common files, envoy sha, test framework, and Rbac Access Denied Policy Info. The right sidebar contains the 'About' section, which describes the Istio proxy components and lists related repositories like `proxy`, `discovery`, `monitoring`, `sidecar`, and `lyft-envoy`. It also includes a 'Releases' section showing the latest version 1.0.2, released on Sep 10, 2018.

istio / proxy

Watch 98 Star 502 Fork 321

Code Pull requests 9 Actions Security Insights

master

Go to file Add file Code

istio-testing Automator: update common-files@master in i... 6 days ago 1,612

File	Commit Message	Time Ago
.github	Update PULL_REQUEST_TEMPLATE.md (#3036)	2 months ago
bazel	Wrap Wasm related library import into a rule (#...	11 days ago
common-protos/googl...	add gen doc for extension protos (#2803)	7 months ago
common	Automator: update common-files@master in ist...	6 days ago
extensions	Update envoy sha (#3111)	8 days ago
prow	Make test framework usable for external repo (...)	2 months ago
scripts	Add Rbac Access Denied Policy Info in Access I...	14 days ago

About

The Istio proxy components.

proxy discovery monitoring sidecar lyft-envoy

Readme

Apache-2.0 License

Releases 205

1.0.2 Latest on Sep 10, 2018

envoy - was ist das?



Fraser Today at 1:30 AM

Whats the plan for the istio fork of envoy, once envoy-wasm is merged into upstream envoy will istio fork the main envoy? or will it use the upstream envoy?

2 replies



John Howard 1 hour ago

Still fork main envoy until we have no more custom filters



John Howard 1 hour ago

Which will happen eventually once they move to wasm

Upstream

Downstream



Where Envoy
connections
come from

Downstream


Upstream

Where Envoy
connections go
to

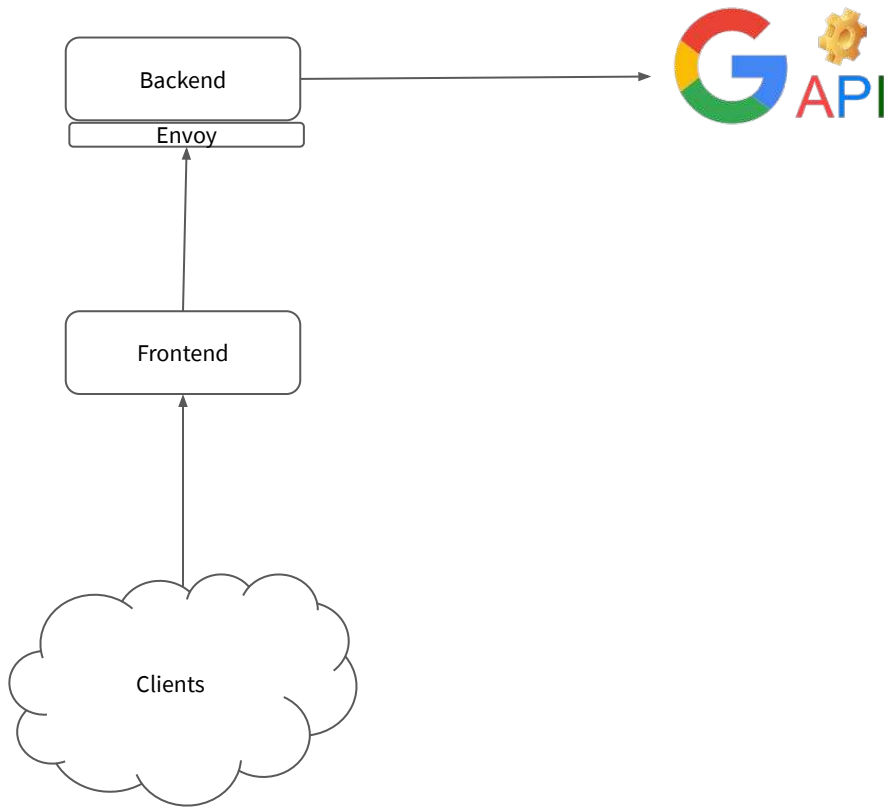
“North-South” Traffic




envoy - upstream and downstream



envoy - upstream and downstream



envoy - upstream and downstream



Clusters
&
Endpoints

Listeners
&
Routes

Clusters
&
Endpoints

≈

Hostnames
&
IP Addresses

envoy - terminology

Clusters & Endpoints

outbound|443||kubernetes.default.svc.cluster.local

10.4.0.10_53

xds-grpc

What to accept
How to process it
&
Where to send it

Listeners
&
Routes




Envoy: Operation

*professional SRE

Sidecars and lions and bears!

envoy - operation



envoy - operation

```
from wsgiref.simple_server import make_server
from cgi import parse_qs

def application (environ, start_response):

    query = parse_qs(environ['QUERY_STRING'])
    foo_val = query.get('foo')

    print "lol, this guy wants his foo like {}".format(foo_val)

    start_response('200 OK', [])
    return ['cool query bro']

httpd = make_server ('localhost', 8051, application)
httpd.handle_request()
```

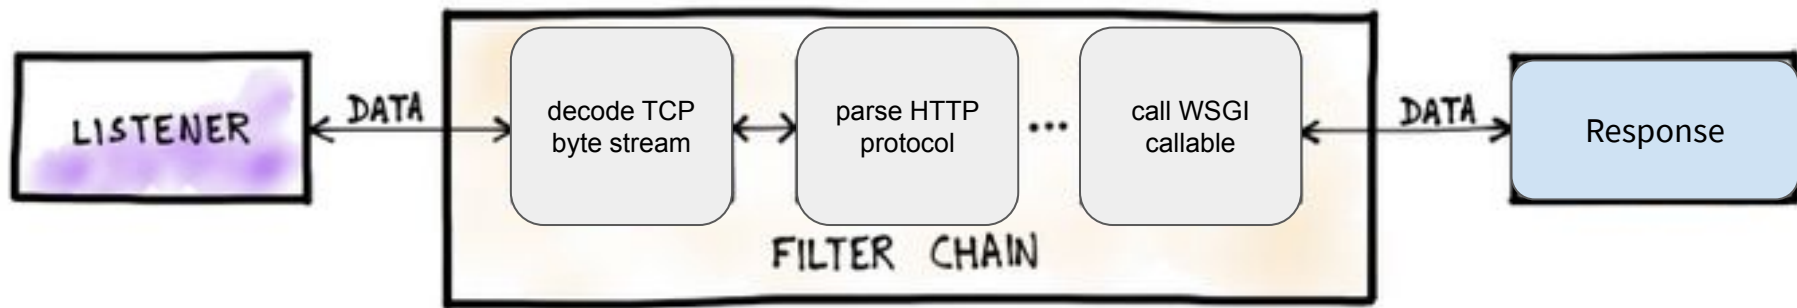
envoy - operation

```
$ echo -e "GET /?foo=bar HTTP/1.1\r\n" | nc localhost 8051
HTTP/1.0 200 OK
Date: Sat, 05 Dec 2020 22:28:28 GMT
Server: WSGIServer/0.1 Python/2.7.18rc1
Content-Length: 14


cool query bro
```

```
$ python web.py
lol, this guy wants his foo like ['bar']
127.0.0.1 - - [05/Dec/2020 15:28:28] "GET /?foo=bar HTTP/1.1" 200 14
```


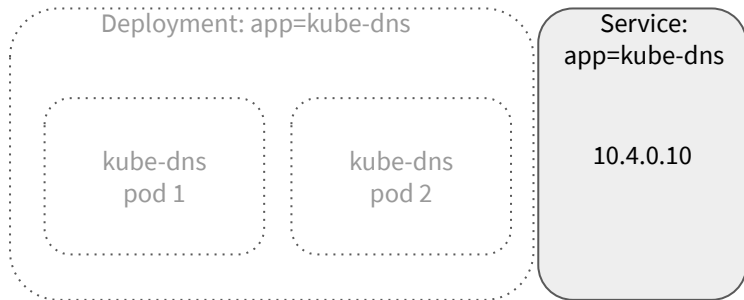
envoy - operation




envoy - operation



k8s networking - classic



K8S Node

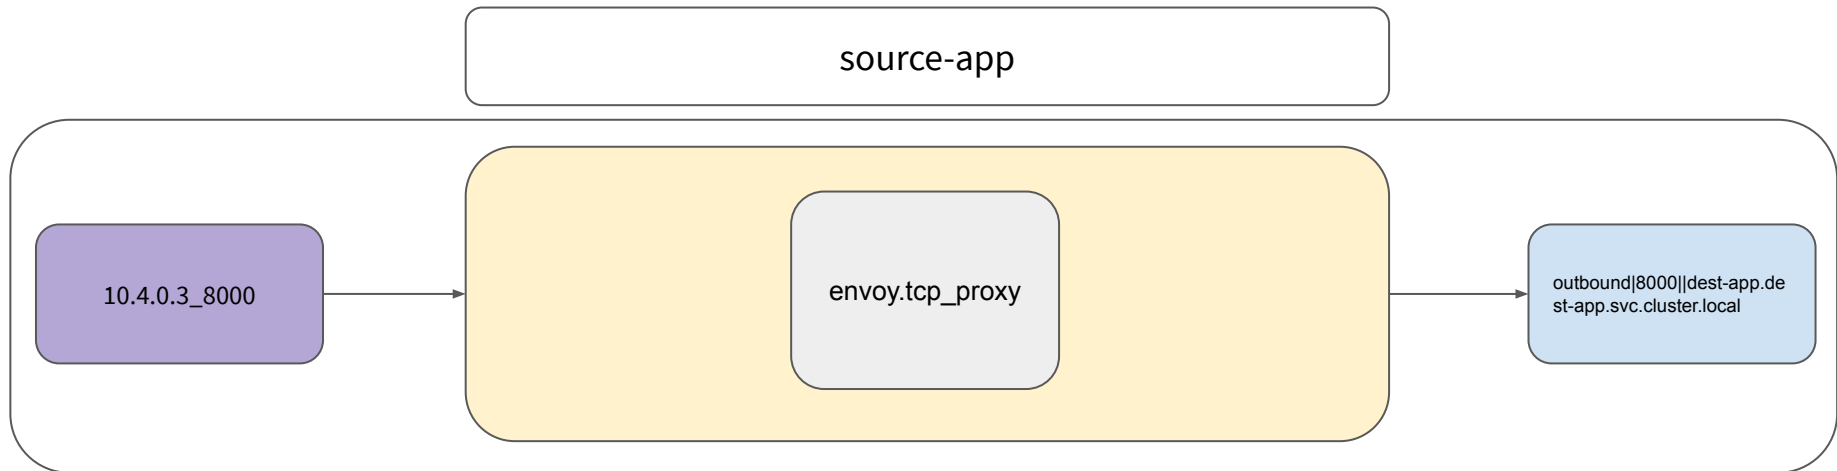


envoy - operation

```
apiVersion: v1
kind: Service
metadata:
  labels:
    app: dest-app
    name: dest-app
    namespace: dest-app
spec:
  clusterIP: 10.4.0.3
  ports:
    - name: tcp-sweetport
      port: 8888
      protocol: TCP
      targetPort: 8888
  selector:
    app: dest-app
  sessionAffinity: None
  type: ClusterIP
```



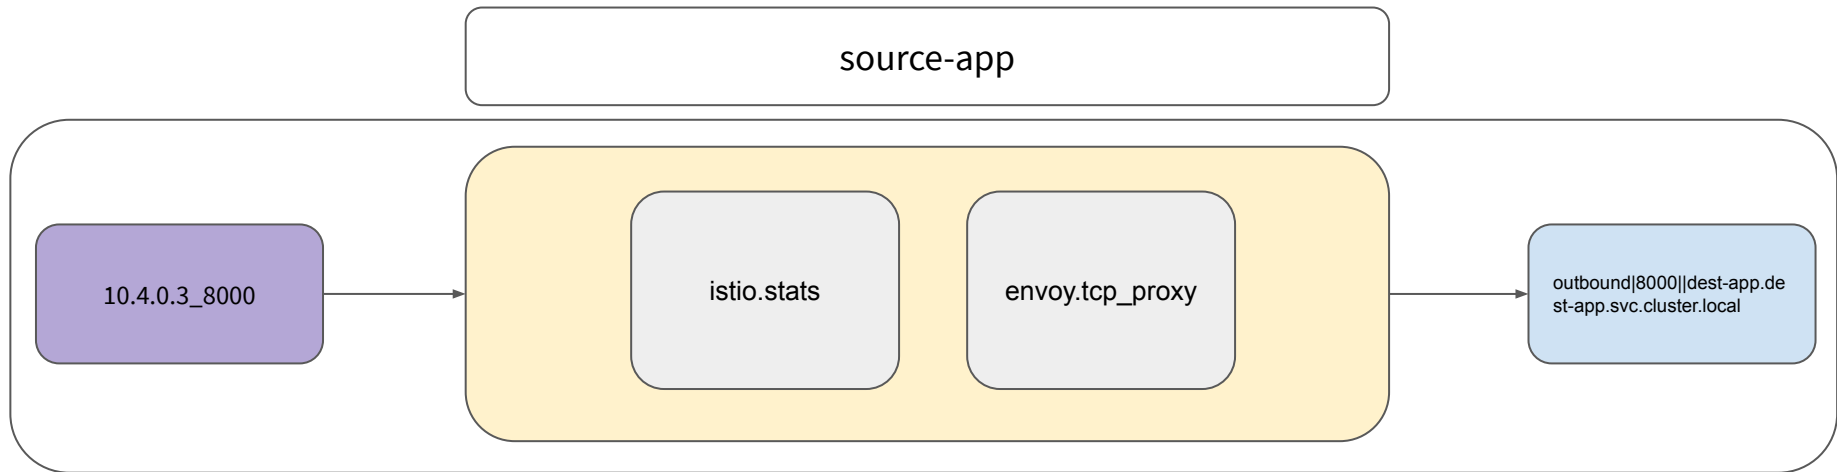
envoy - operation



envoy - operation

```
[
  {
    "name": "10.4.0.3_8000",
    "address": {
      "socketAddress": {
        "address": "10.4.0.3",
        "portValue": 8000
      }
    },
    "filterChains": [
      {
        "filters": [
          {
            "name": "envoy.tcp_proxy",
            "typedConfig": {
              "@type": "type.googleapis.com/envoy.config.filter.network.tcp_proxy.v2.TcpProxy",
              "statPrefix": "outbound|8000||dest-app.dest-app.svc.cluster.local",
              "cluster": "outbound|8000||dest-app.dest-app.svc.cluster.local"
            }
          }
        ]
      }
    ],
    "deprecatedV1": {
      "bindToPort": false
    },
    "trafficDirection": "OUTBOUND"
  }
]
```

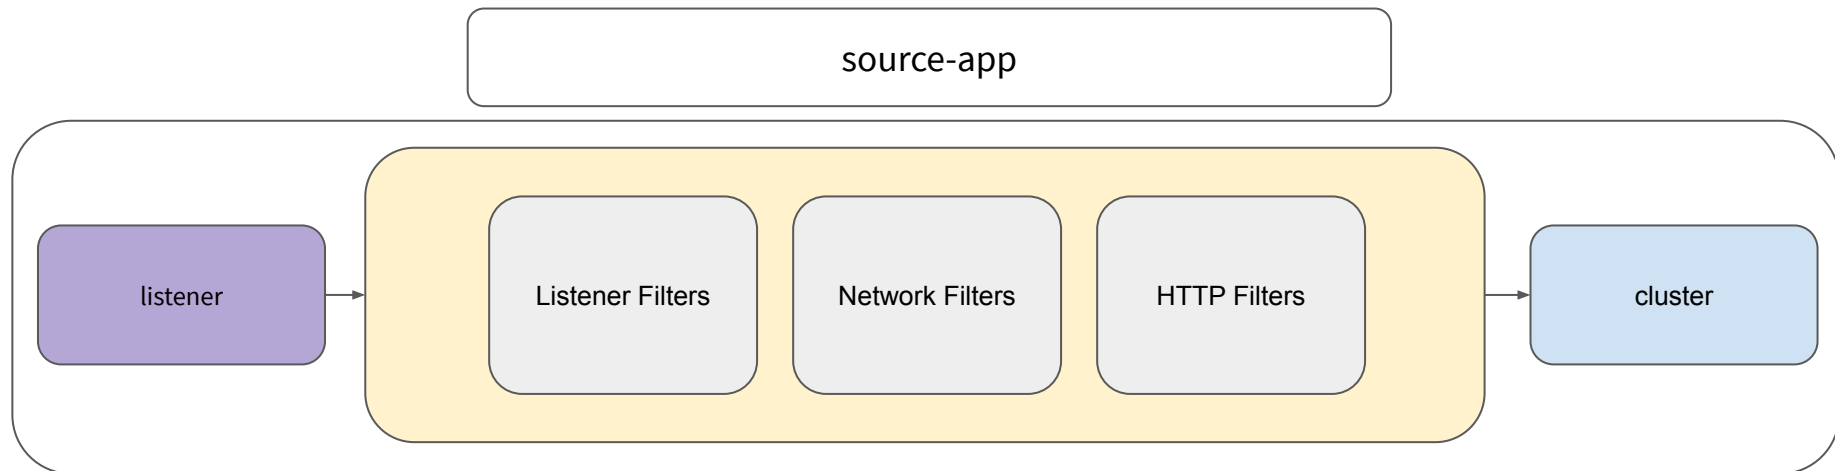
envoy - operation



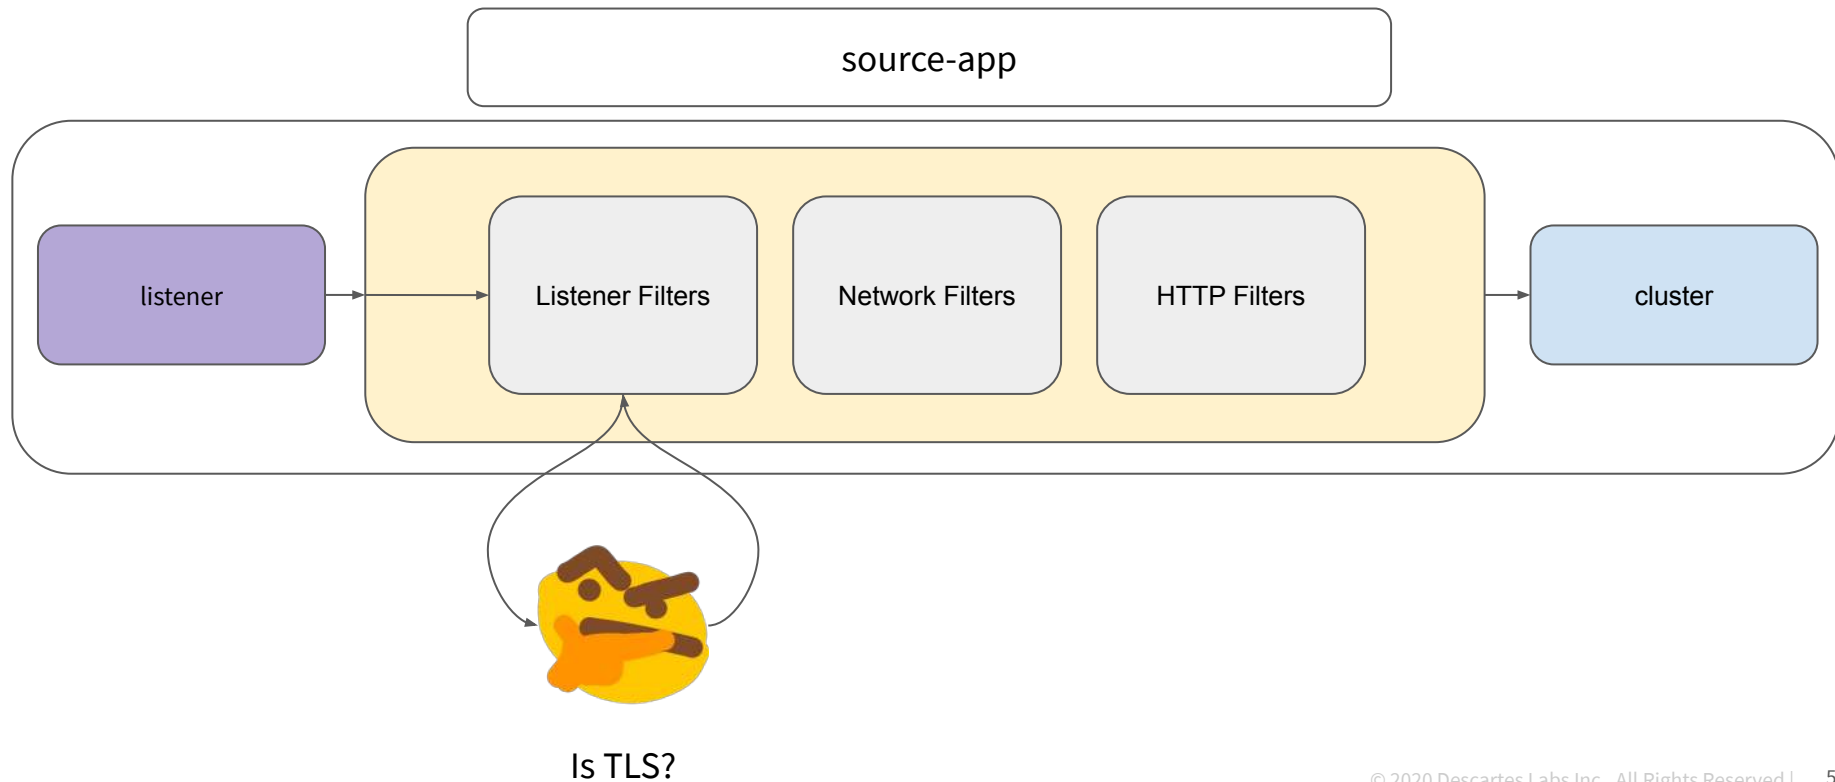
Extensions & Plugins

- Access loggers
- Retry implementations
- Tracers
- Resource Monitors
- Credential Providers

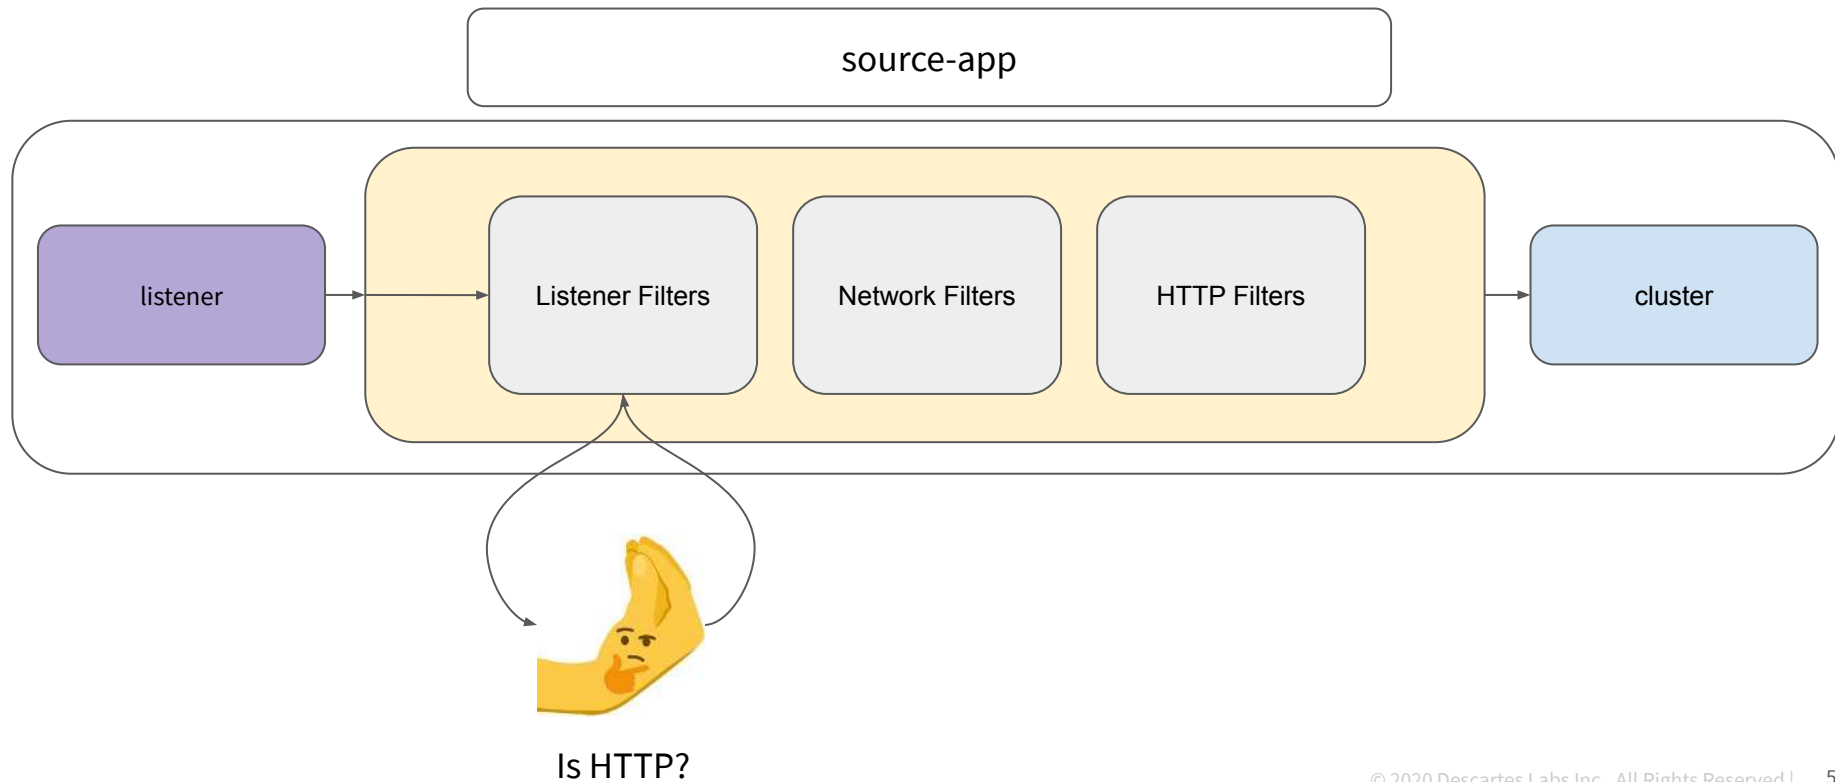
envoy - operation - filters



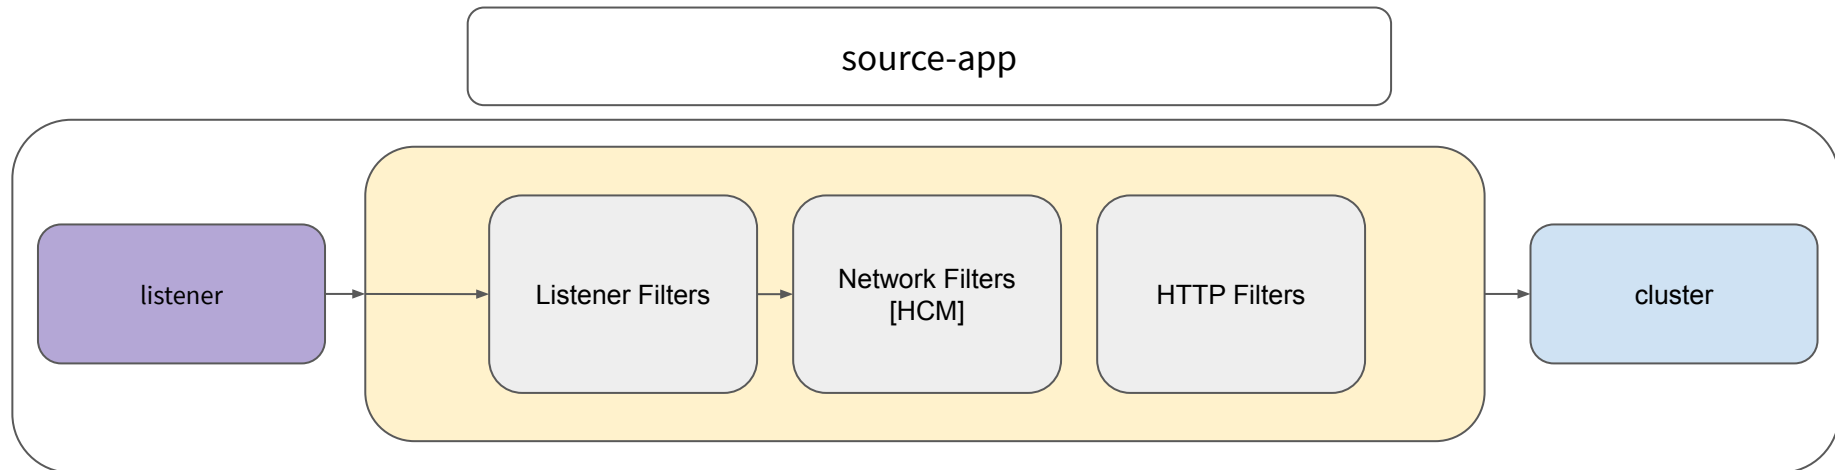
envoy - operation - filter chains



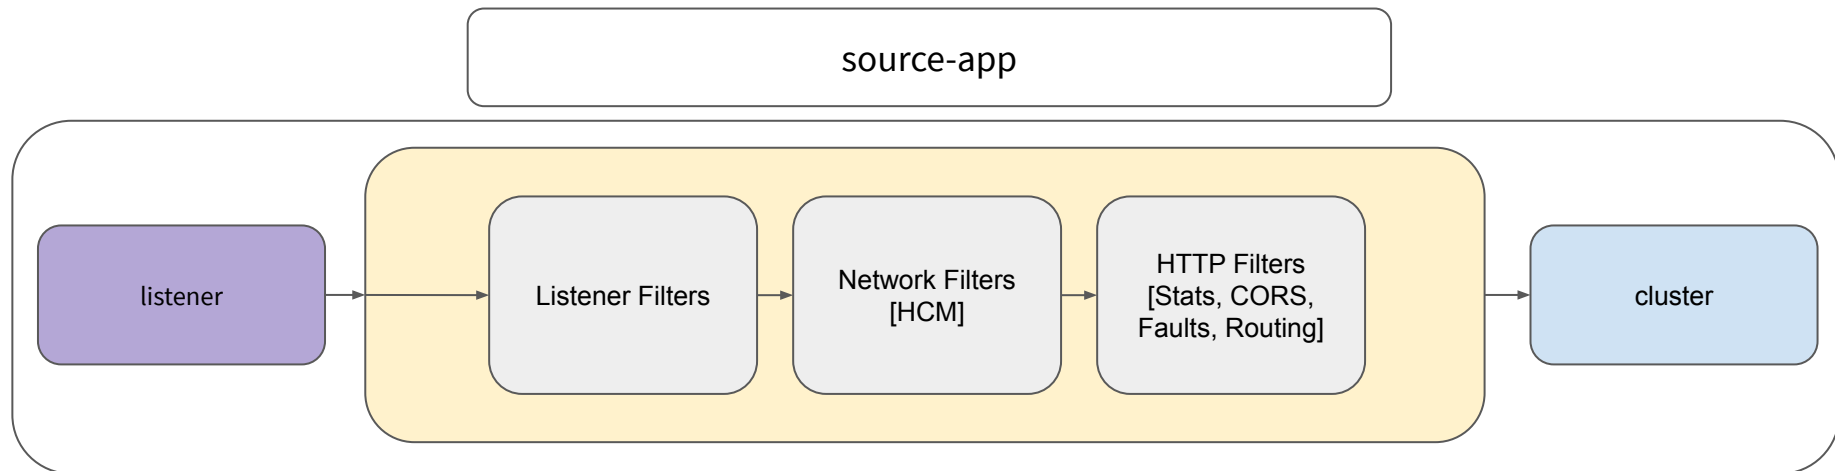
envoy - operation - filter chains



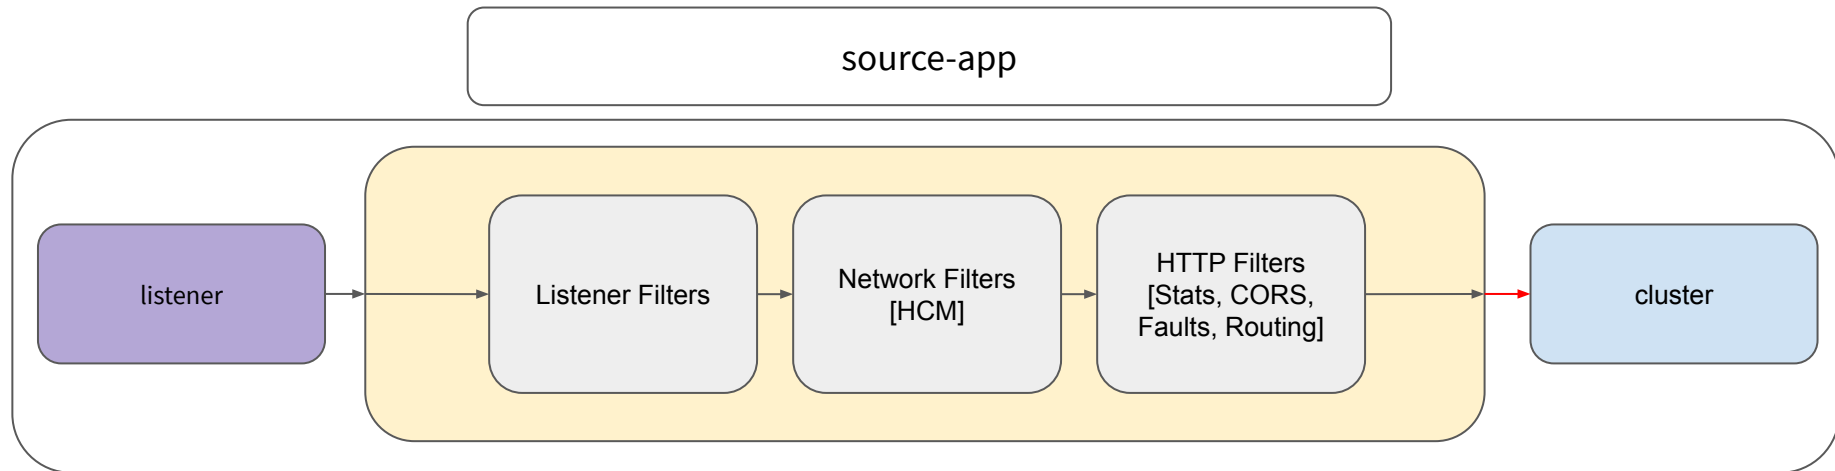
envoy - operation - filter chains



envoy - operation - filter chains

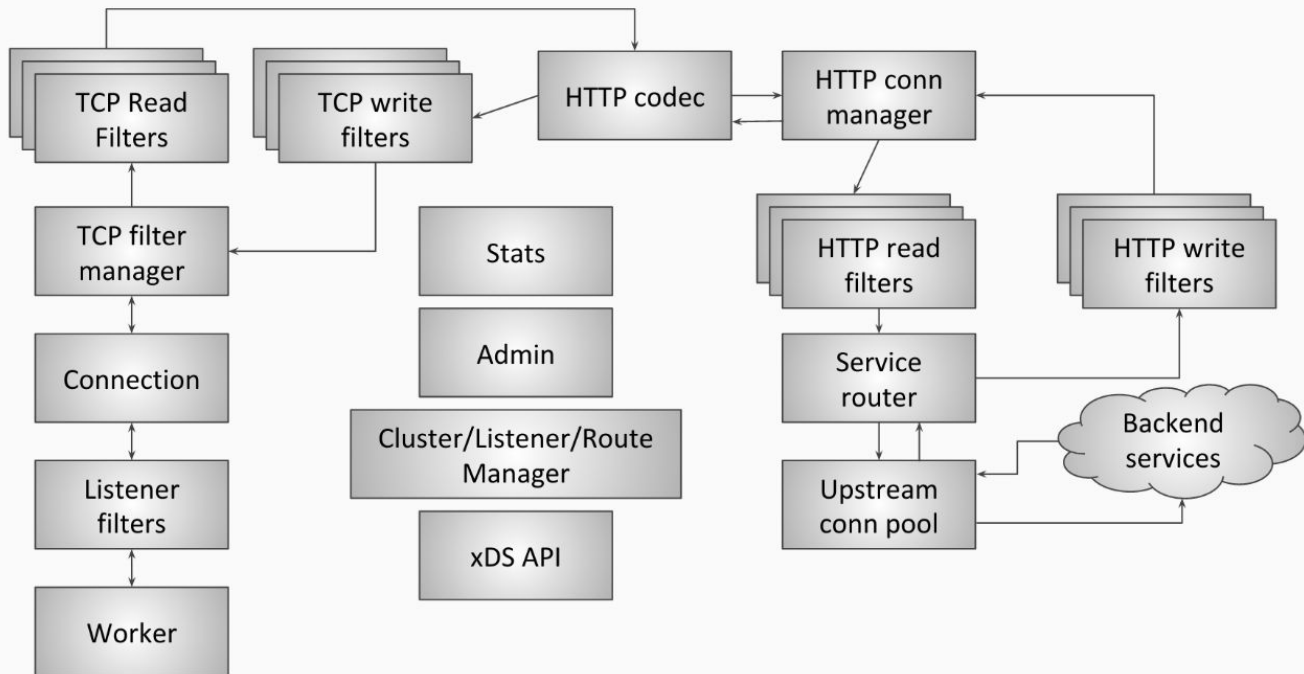


envoy - operation - filter chains



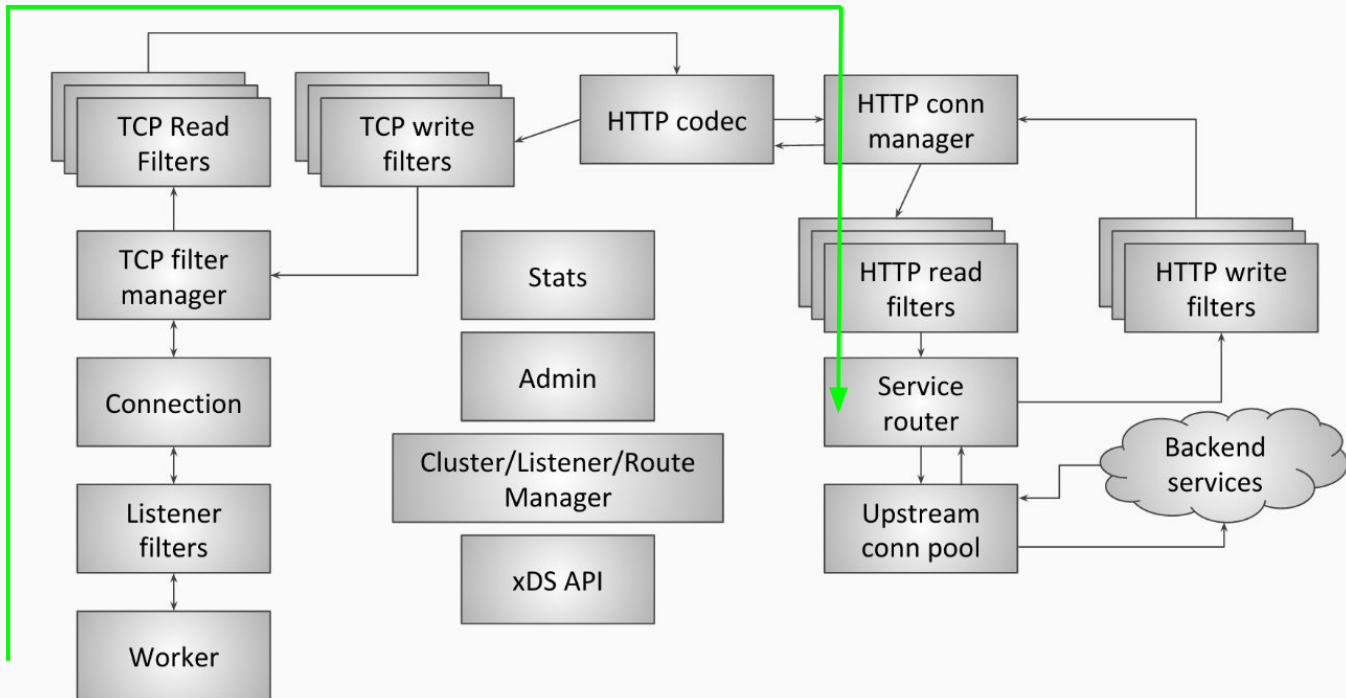
envoy - operation - cluster manager

Envoy architecture diagram



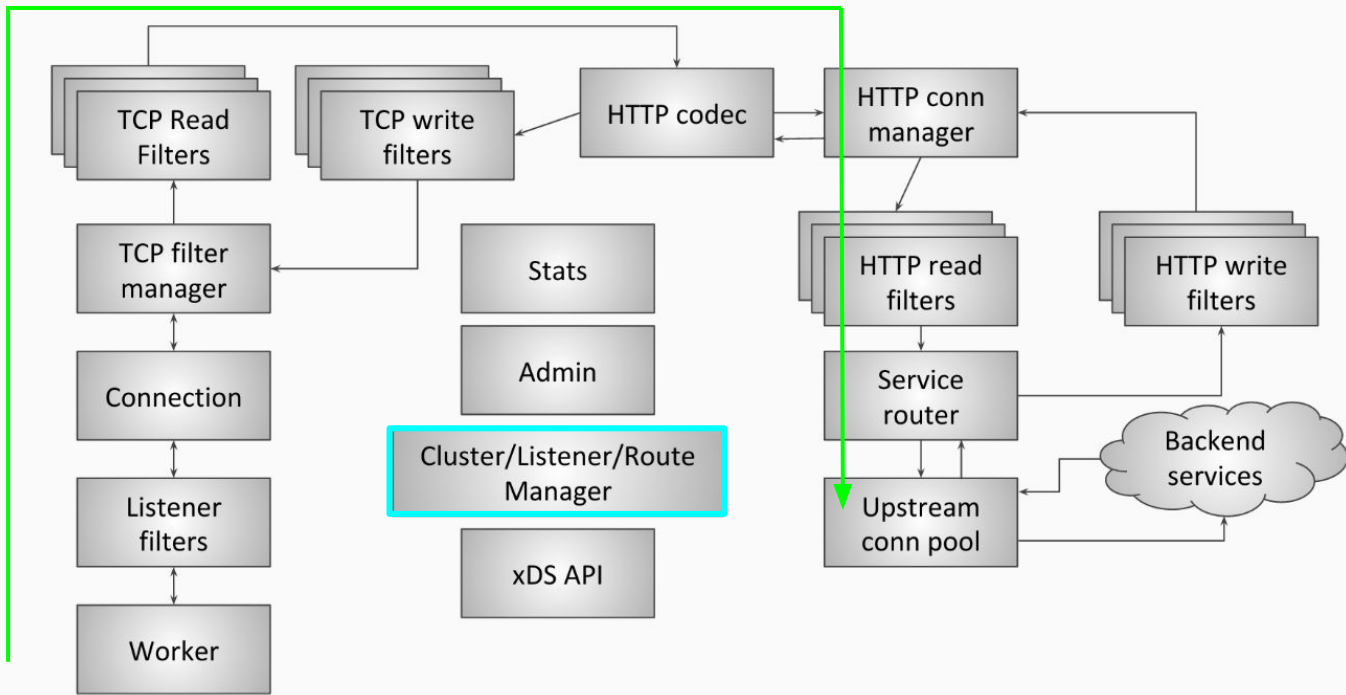
envoy - operation - cluster manager

Envoy architecture diagram

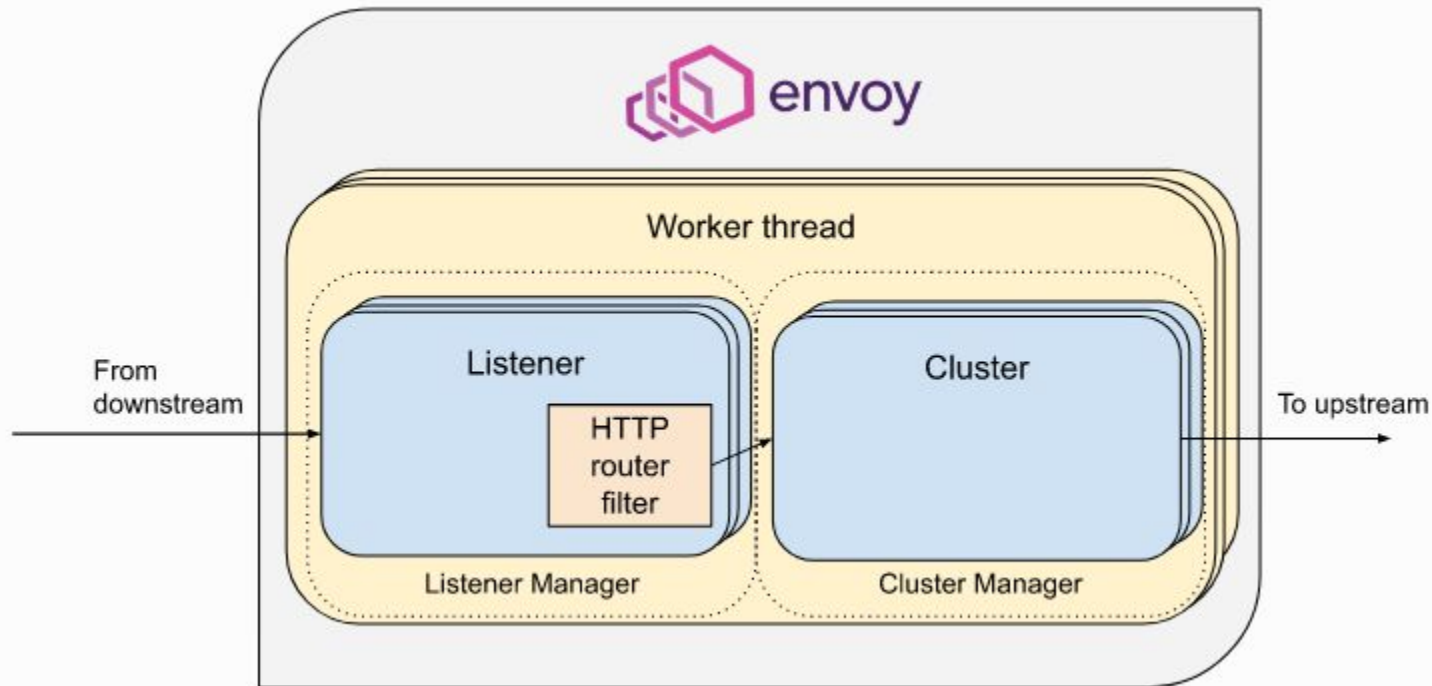


envoy - operation - cluster manager

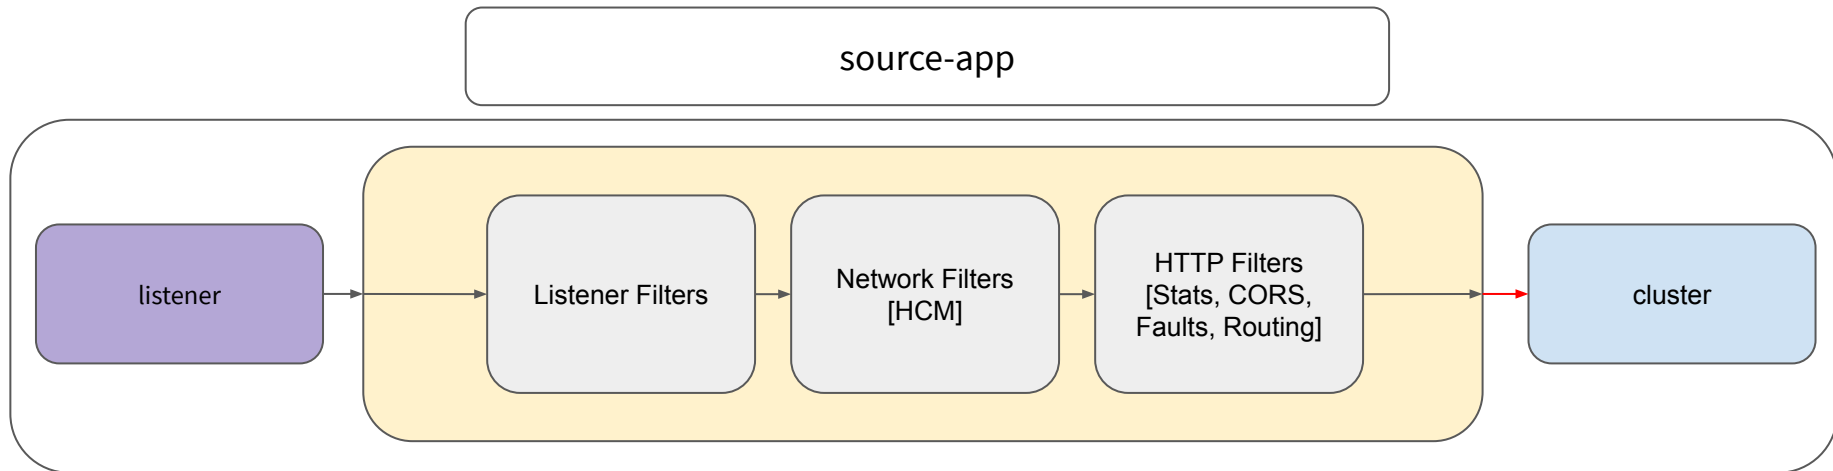
Envoy architecture diagram



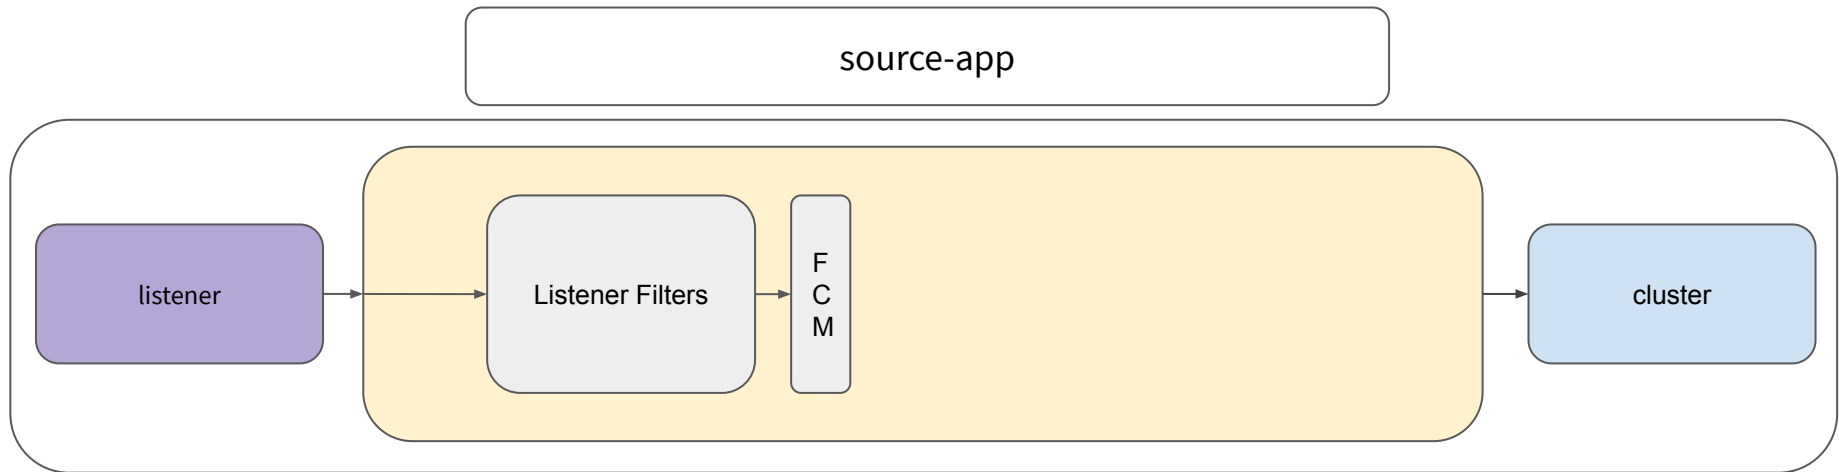
envoy - operation - cluster manager



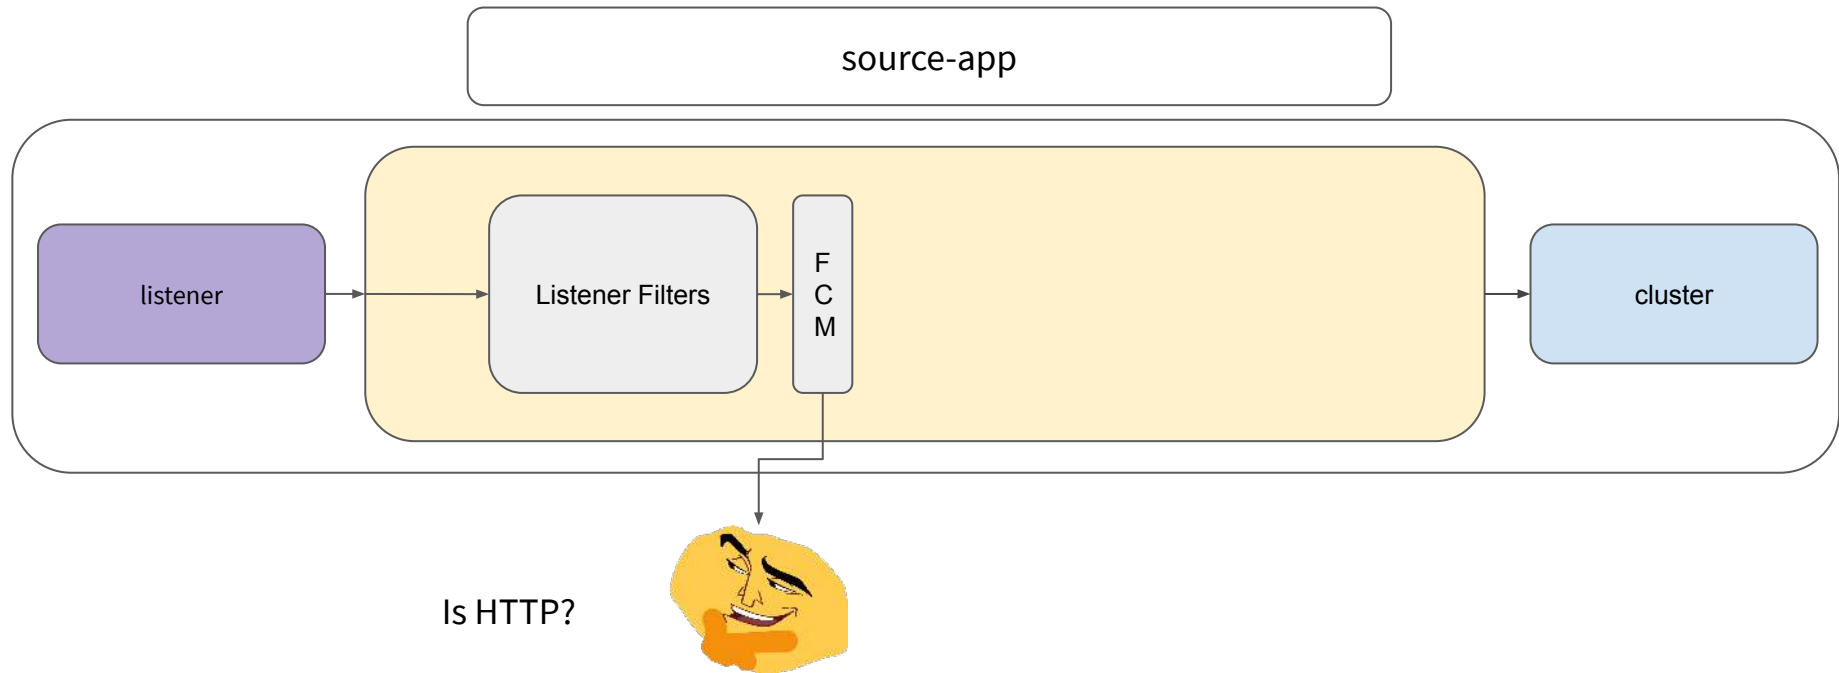
envoy - operation - filter chains



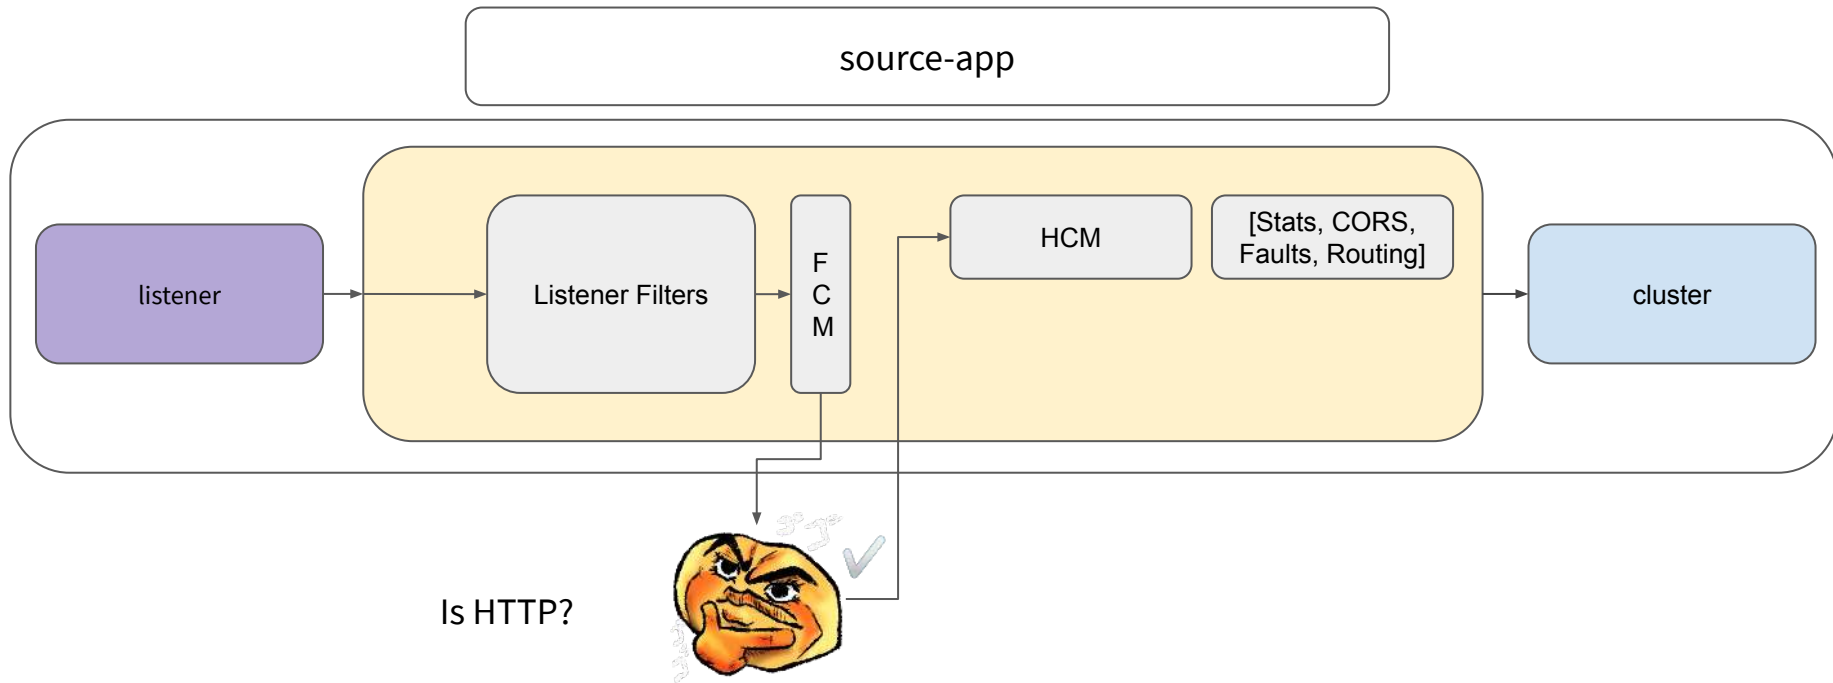
envoy - operation - filter chain match



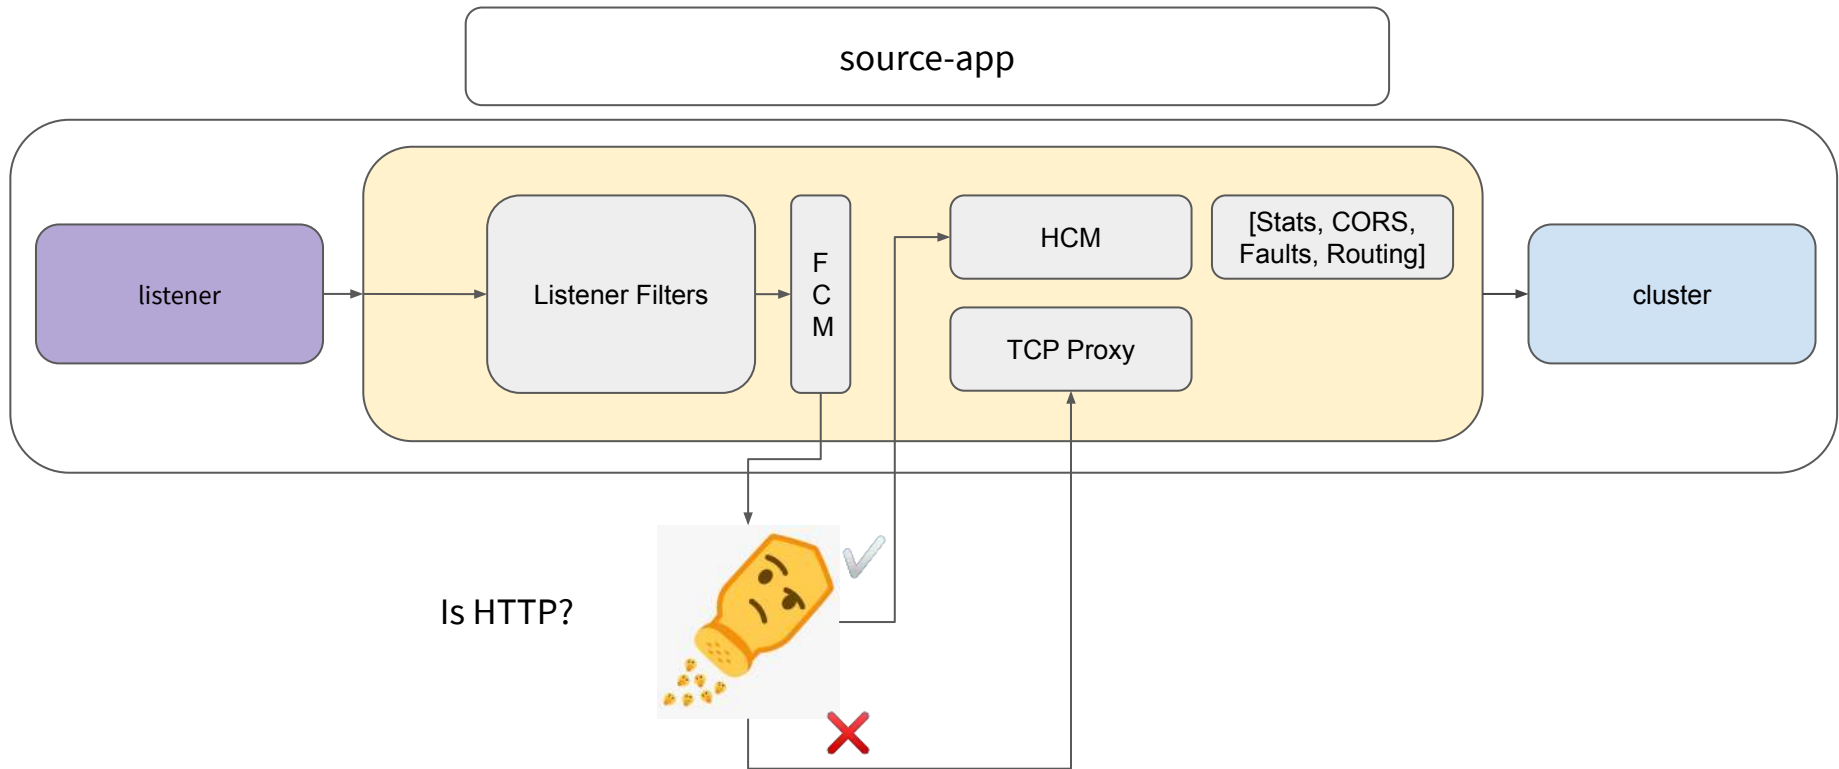
envoy - operation - filter chain match



envoy - operation - filter chain match



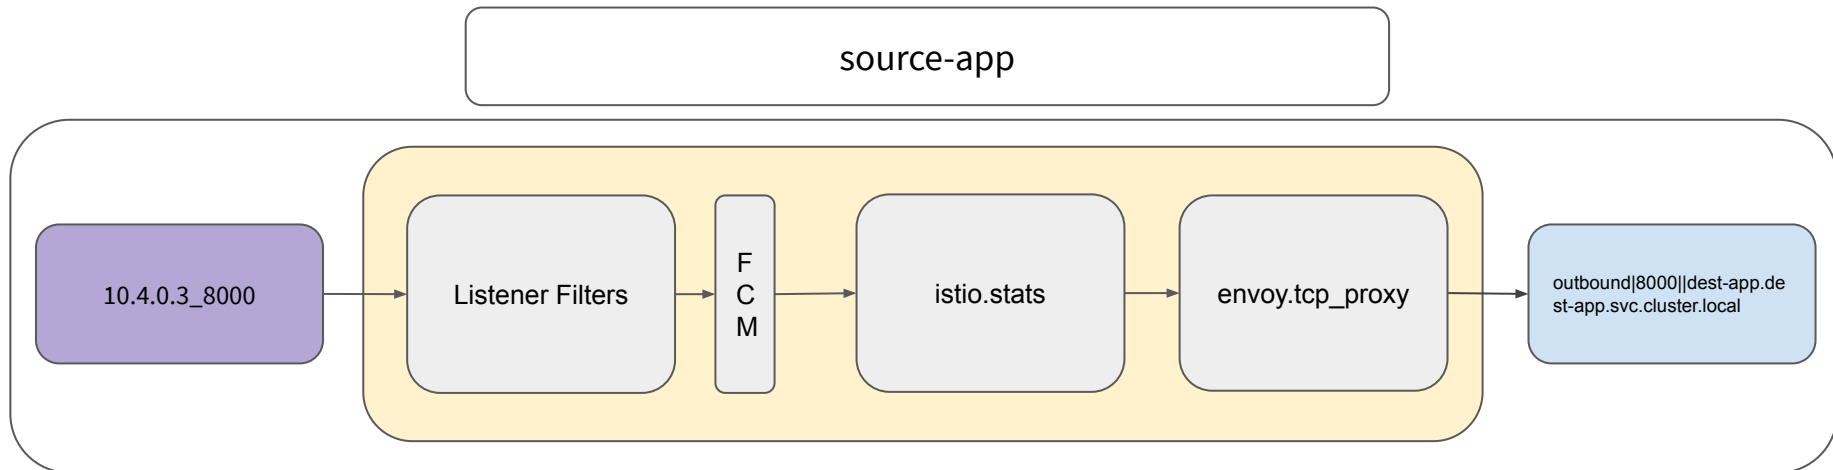
envoy - operation - filter chain match



envoy - operation - filter chain match

```
{
  "name": "0.0.0.0_15010",
  "active_state": {
    "version_info": "2020-12-05T22:58:30Z/8",
    "listener": {
      "@type": "type.googleapis.com/envoy.api.v2.Listener",
      "name": "0.0.0.0_15010",
      "address": {
        "socket_address": {
          "address": "0.0.0.0",
          "port_value": 15010
        }
      },
      "filter_chains": [
        {
          "filter_chain_match": {
            "application_protocols": [
              "http/1.0",
              "http/1.1",
              "h2c"
            ]
          },
          "filters": [...] // 1 item
        },
        {
          "filter_chain_match": {},
          "filters": [...], // 2 items
          "metadata": {...}, // 1 item
          "name": "PassthroughFilterChain"
        }
      ],
      "deprecated_v1": {...}, // 1 item
      "listener_filters": [...], // 2 items
      "listener_filters_timeout": "0.100s",
      "traffic_direction": "OUTBOUND",
      "continue_on_listener_filters_timeout": true
    },
    "last_updated": "2020-12-05T23:00:19.805Z"
  },
}
```

envoy - operation



envoy - operation

```
gke-compute-stage-preemptible-n1-stan-bddbfc3b-thtc /home/rob # iptables -t nat -S
-P PREROUTING ACCEPT
-P INPUT ACCEPT
-P OUTPUT ACCEPT
-P POSTROUTING ACCEPT
-N ISTIO_INBOUND
-N ISTIO_IN_REDIRECT
-N ISTIO_OUTPUT
-N ISTIO_REDIRECT
-A PREROUTING -p tcp -j ISTIO_INBOUND
-A OUTPUT -p tcp -j ISTIO_OUTPUT
-A ISTIO_INBOUND -p tcp -m tcp --dport 22 -j RETURN
-A ISTIO_INBOUND -p tcp -m tcp --dport 15090 -j RETURN
-A ISTIO_INBOUND -p tcp -m tcp --dport 15021 -j RETURN
-A ISTIO_INBOUND -p tcp -m tcp --dport 15020 -j RETURN
-A ISTIO_INBOUND -p tcp -j ISTIO_IN_REDIRECT
-A ISTIO_IN_REDIRECT -p tcp -j REDIRECT --to-ports 15006
-A ISTIO_OUTPUT -s 127.0.0.6/32 -o lo -j RETURN
-A ISTIO_OUTPUT ! -d 127.0.0.1/32 -o lo -m owner --uid-owner 1337 -j ISTIO_IN_REDIRECT
-A ISTIO_OUTPUT -o lo -m owner ! --uid-owner 1337 -j RETURN
-A ISTIO_OUTPUT -m owner --uid-owner 1337 -j RETURN
-A ISTIO_OUTPUT ! -d 127.0.0.1/32 -o lo -m owner --gid-owner 1337 -j ISTIO_IN_REDIRECT
-A ISTIO_OUTPUT -o lo -m owner ! --gid-owner 1337 -j RETURN
-A ISTIO_OUTPUT -m owner --gid-owner 1337 -j RETURN
-A ISTIO_OUTPUT -d 127.0.0.1/32 -j RETURN
-A ISTIO_OUTPUT -j ISTIO_REDIRECT
-A ISTIO_REDIRECT -p tcp -j REDIRECT --to-ports 15001
```

envoy - operation

```
gke-compute-stage-preemptible-n1-stan-bddbfc3b-thtc /home/rob # iptables -t nat -S
-P PREROUTING ACCEPT
-P INPUT ACCEPT
-P OUTPUT ACCEPT
-P POSTROUTING ACCEPT
-N ISTIO_INBOUND
-N ISTIO_IN_REDIRECT
-N ISTIO_OUTPUT
-N ISTIO_REDIRECT
-A PREROUTING -p tcp -j ISTIO_INBOUND
-A OUTPUT -p tcp -j ISTIO_OUTPUT
-A ISTIO_INBOUND -p tcp -m tcp --dport 22 -j RETURN
-A ISTIO_INBOUND -p tcp -m tcp --dport 15090 -j RETURN
-A ISTIO_INBOUND -p tcp -m tcp --dport 15021 -j RETURN
-A ISTIO_INBOUND -p tcp -m tcp --dport 15020 -j RETURN
-A ISTIO_INBOUND -p tcp -j ISTIO_IN_REDIRECT
-A ISTIO_IN_REDIRECT -p tcp -j REDIRECT --to-ports 15006
-A ISTIO_OUTPUT -s 127.0.0.6/32 -o lo -j RETURN
-A ISTIO_OUTPUT ! -d 127.0.0.1/32 -o lo -m owner --uid-owner 1337 -j ISTIO_IN_REDIRECT
-A ISTIO_OUTPUT -o lo -m owner ! --uid-owner 1337 -j RETURN
-A ISTIO_OUTPUT -m owner --uid-owner 1337 -j RETURN
-A ISTIO_OUTPUT ! -d 127.0.0.1/32 -o lo -m owner --gid-owner 1337 -j ISTIO_IN_REDIRECT
-A ISTIO_OUTPUT -o lo -m owner ! --gid-owner 1337 -j RETURN
-A ISTIO_OUTPUT -m owner --gid-owner 1337 -j RETURN
-A ISTIO_OUTPUT -d 127.0.0.1/32 -j RETURN
-A ISTIO_OUTPUT -j ISTIO_REDIRECT
-A ISTIO_REDIRECT -p tcp -j REDIRECT --to-ports 15001
```

envoy - operation

```
gke-compute-stage-preemptible-n1-stan-bddbfc3b-thtc /home/rob # iptables -t nat -S
-P PREROUTING ACCEPT
-P INPUT ACCEPT
-P OUTPUT ACCEPT
-P POSTROUTING ACCEPT
-N ISTIO_INBOUND
-N ISTIO_IN_REDIRECT
-N ISTIO_OUTPUT
-N ISTIO_REDIRECT
-A PREROUTING -p tcp -j ISTIO_INBOUND
-A OUTPUT -p tcp -j ISTIO_OUTPUT
-A ISTIO_INBOUND -p tcp -m tcp --dport 22 -j RETURN
-A ISTIO_INBOUND -p tcp -m tcp --dport 15090 -j RETURN
-A ISTIO_INBOUND -p tcp -m tcp --dport 15021 -j RETURN
-A ISTIO_INBOUND -p tcp -m tcp --dport 15020 -j RETURN
-A ISTIO_INBOUND -p tcp -j ISTIO_IN_REDIRECT
-A ISTIO_IN_REDIRECT -p tcp -j REDIRECT --to-ports 15006
-A ISTIO_OUTPUT -s 127.0.0.6/32 -o lo -j RETURN
-A ISTIO_OUTPUT ! -d 127.0.0.1/32 -o lo -m owner --uid-owner 1337 -j ISTIO_IN_REDIRECT
-A ISTIO_OUTPUT -o lo -m owner ! --uid-owner 1337 -j RETURN
-A ISTIO_OUTPUT -m owner --uid-owner 1337 -j RETURN
-A ISTIO_OUTPUT ! -d 127.0.0.1/32 -o lo -m owner --gid-owner 1337 -j ISTIO_IN_REDIRECT
-A ISTIO_OUTPUT -o lo -m owner ! --gid-owner 1337 -j RETURN
-A ISTIO_OUTPUT -m owner --gid-owner 1337 -j RETURN
-A ISTIO_OUTPUT -d 127.0.0.1/32 -j RETURN
-A ISTIO_OUTPUT -j ISTIO_REDIRECT
-A ISTIO_REDIRECT -p tcp -j REDIRECT --to-ports 15001
```

envoy - operation

```
gke-compute-stage-preemptible-n1-stan-bddbfc3b-thtc /home/rob # iptables -t nat -S
-P PREROUTING ACCEPT
-P INPUT ACCEPT
-P OUTPUT ACCEPT
-P POSTROUTING ACCEPT
-N ISTIO_INBOUND
-N ISTIO_IN_REDIRECT
-N ISTIO_OUTPUT
-N ISTIO_REDIRECT
-A PREROUTING -p tcp -j ISTIO_INBOUND
-A OUTPUT -p tcp -j ISTIO_OUTPUT
-A ISTIO_INBOUND -p tcp -m tcp --dport 22 -j RETURN
-A ISTIO_INBOUND -p tcp -m tcp --dport 15090 -j RETURN
-A ISTIO_INBOUND -p tcp -m tcp --dport 15021 -j RETURN
-A ISTIO_INBOUND -p tcp -m tcp --dport 15020 -j RETURN
-A ISTIO_INBOUND -p tcp -j ISTIO_IN_REDIRECT
-A ISTIO_IN_REDIRECT -p tcp -j REDIRECT --to-ports 15006

-A ISTIO_OUTPUT -j ISTIO_IN_REDIRECT

-A ISTIO_OUTPUT -j ISTIO_IN_REDIRECT

-A ISTIO_OUTPUT -j ISTIO_REDIRECT
-A ISTIO_REDIRECT -p tcp -j REDIRECT --to-ports 15001
```


envoy - operation

```
gke-compute-stage-preemptible-n1-stan-bddbfc3b-thtc /home/rob # iptables -t nat -S
```

```
-P PREROUTING ACCEPT
-P INPUT ACCEPT
-P OUTPUT ACCEPT
-P POSTROUTING ACCEPT
-N ISTIO_INBOUND
-N ISTIO_IN_REDIRECT
-N ISTIO_OUTPUT
-N ISTIO_REDIRECT
-A PREROUTING -p tcp -j ISTIO_INBOUND
-A OUTPUT -p tcp -j ISTIO_OUTPUT
-A ISTIO_INBOUND -p tcp -m tcp --dport 22 -j RETURN
-A ISTIO_INBOUND -p tcp -m tcp --dport 15090 -j RETURN
-A ISTIO_INBOUND -p tcp -m tcp --dport 15021 -j RETURN
-A ISTIO_INBOUND -p tcp -m tcp --dport 15020 -j RETURN
-A ISTIO_INBOUND -p tcp -j ISTIO_IN_REDIRECT
-A ISTIO_IN_REDIRECT -p tcp -j REDIRECT --to-ports 15006
```

```
-A ISTIO_REDIRECT -p tcp -j REDIRECT --to-ports 15001
```

envoy - operation

```
gke-compute-stage-preemptible-n1-stan-bddbfc3b-thtc /home/rob # iptables -t nat -S
-P PREROUTING ACCEPT
-P INPUT ACCEPT
-P OUTPUT ACCEPT
-P POSTROUTING ACCEPT
-N ISTIO_INBOUND
-N ISTIO_IN_REDIRECT
-N ISTIO_OUTPUT
-N ISTIO_REDIRECT
-A PREROUTING -p tcp -j ISTIO_INBOUND ←
-A OUTPUT -p tcp -j ISTIO_OUTPUT

[REDACTED]

-A ISTIO_INBOUND -p tcp -j ISTIO_IN_REDIRECT
-A ISTIO_IN_REDIRECT -p tcp -j REDIRECT --to-ports 15006
-A ISTIO_OUTPUT -s 127.0.0.6/32 -o lo -j RETURN
-A ISTIO_OUTPUT ! -d 127.0.0.1/32 -o lo -m owner --uid-owner 1337 -j ISTIO_IN_REDIRECT
-A ISTIO_OUTPUT -o lo -m owner ! --uid-owner 1337 -j RETURN
-A ISTIO_OUTPUT -m owner --uid-owner 1337 -j RETURN
-A ISTIO_OUTPUT ! -d 127.0.0.1/32 -o lo -m owner --gid-owner 1337 -j ISTIO_IN_REDIRECT
-A ISTIO_OUTPUT -o lo -m owner ! --gid-owner 1337 -j RETURN
-A ISTIO_OUTPUT -m owner --gid-owner 1337 -j RETURN
-A ISTIO_OUTPUT -d 127.0.0.1/32 -j RETURN
-A ISTIO_OUTPUT -j ISTIO_REDIRECT
-A ISTIO_REDIRECT -p tcp -j REDIRECT --to-ports 15001
```

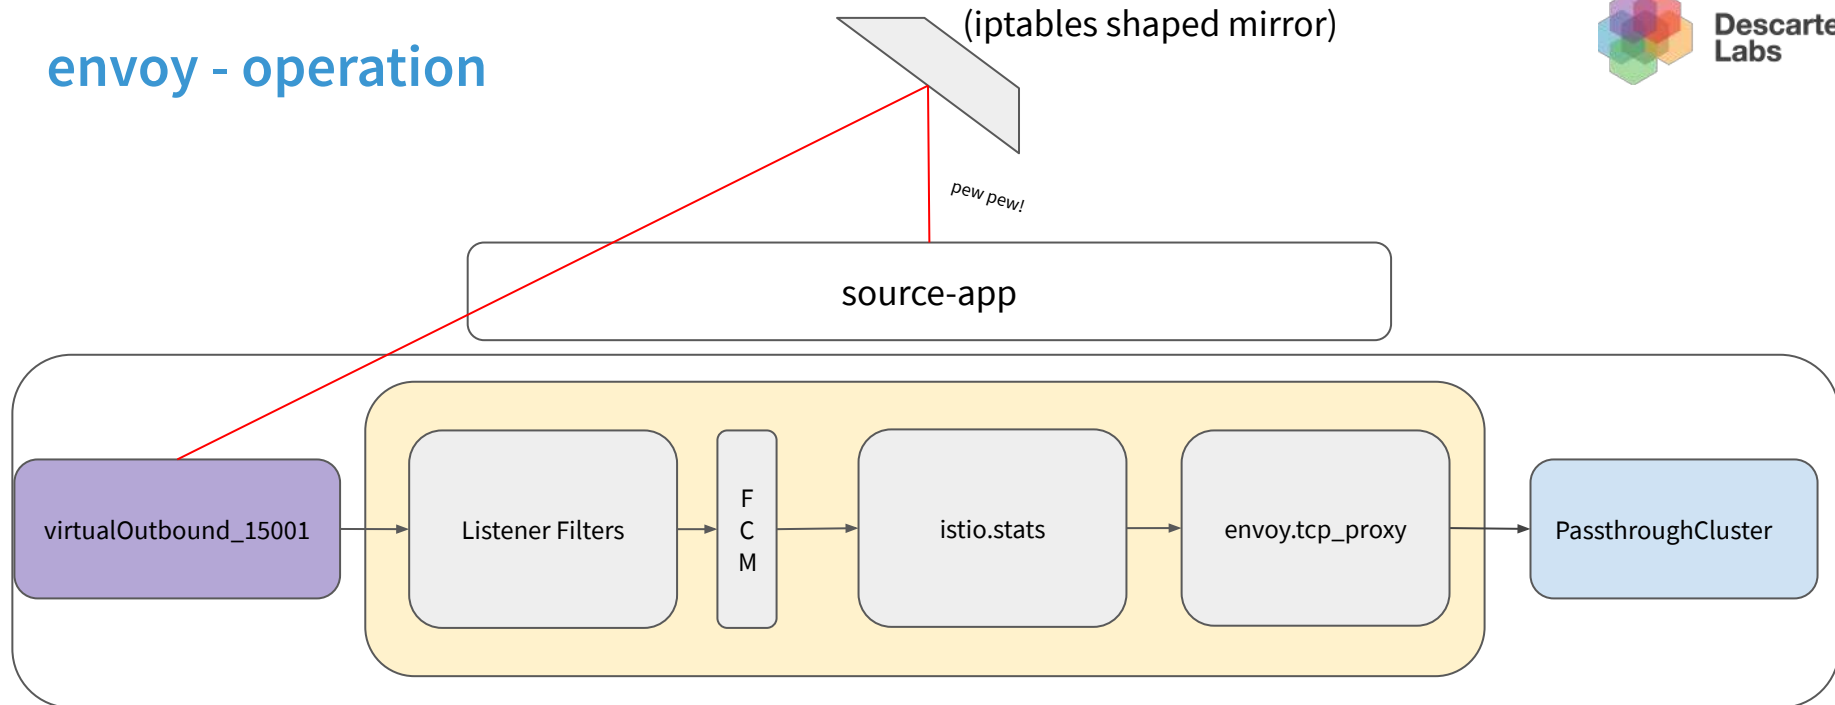
envoy - operation

```
$ ss -ltp
```

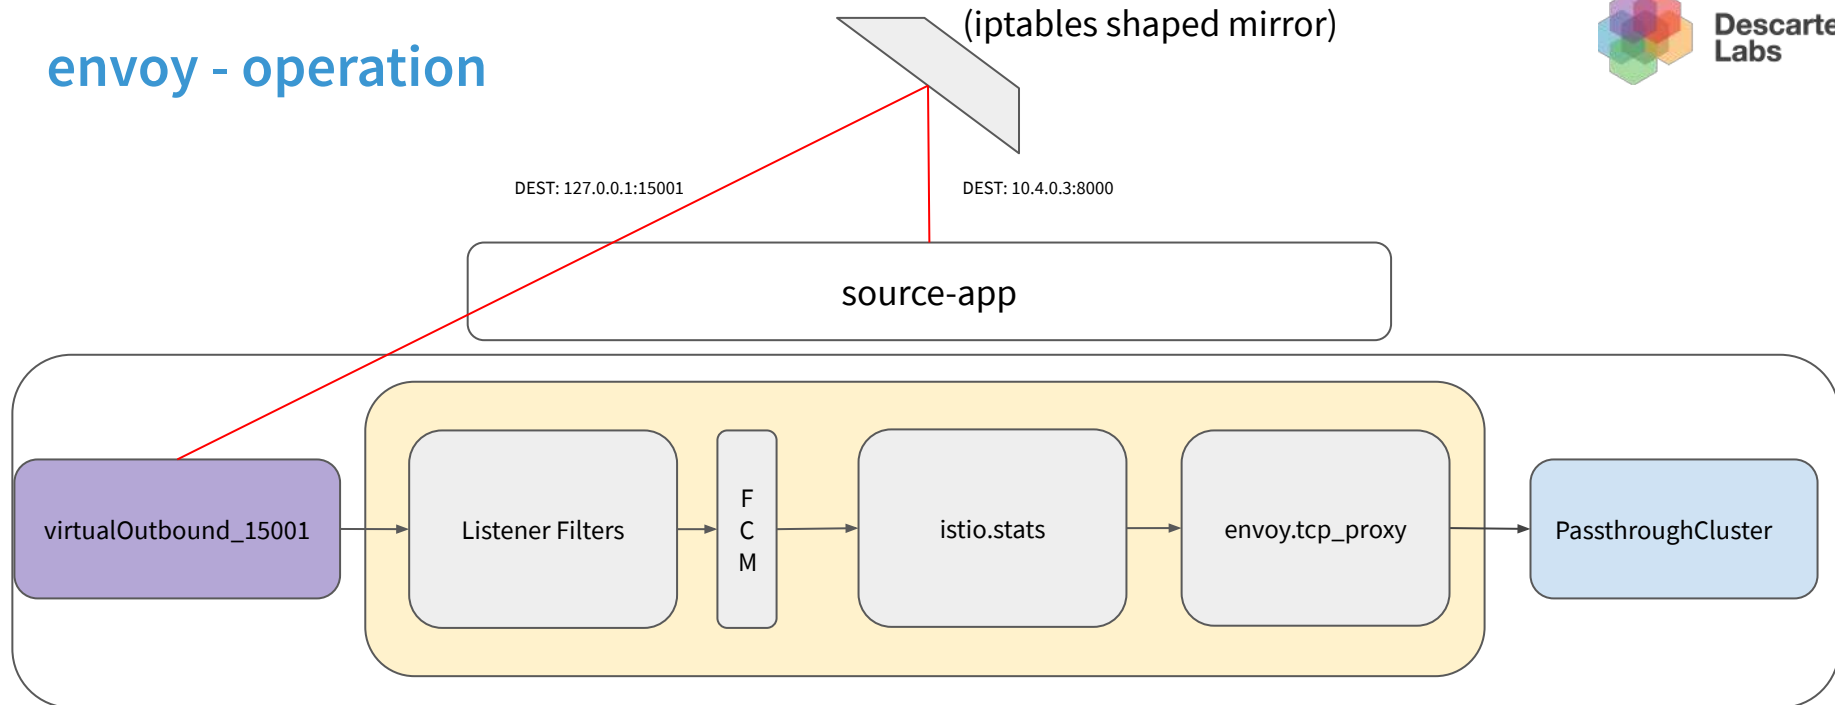
State	Recv-Q	Send-Q	Local Address:Port	Peer Address:Port	users:((("envoy",pid=14,fd=25))
LISTEN	0	128	0.0.0.0:15021	0.0.0.0:*	users:((("envoy",pid=14,fd=24))
LISTEN	0	128	0.0.0.0:15090	0.0.0.0:*	users:((("envoy",pid=14,fd=14))
LISTEN	0	128	127.0.0.1:15000	0.0.0.0:*	users:((("envoy",pid=14,fd=45))
LISTEN	0	128	0.0.0.0:15001	0.0.0.0:*	users:((("envoy",pid=14,fd=46))
LISTEN	0	128	0.0.0.0:15006	0.0.0.0:*	users:((("pilot-agent",pid=1,fd=7))
LISTEN	0	1024	*:15020	*:*	

```
$ █
```

envoy - operation

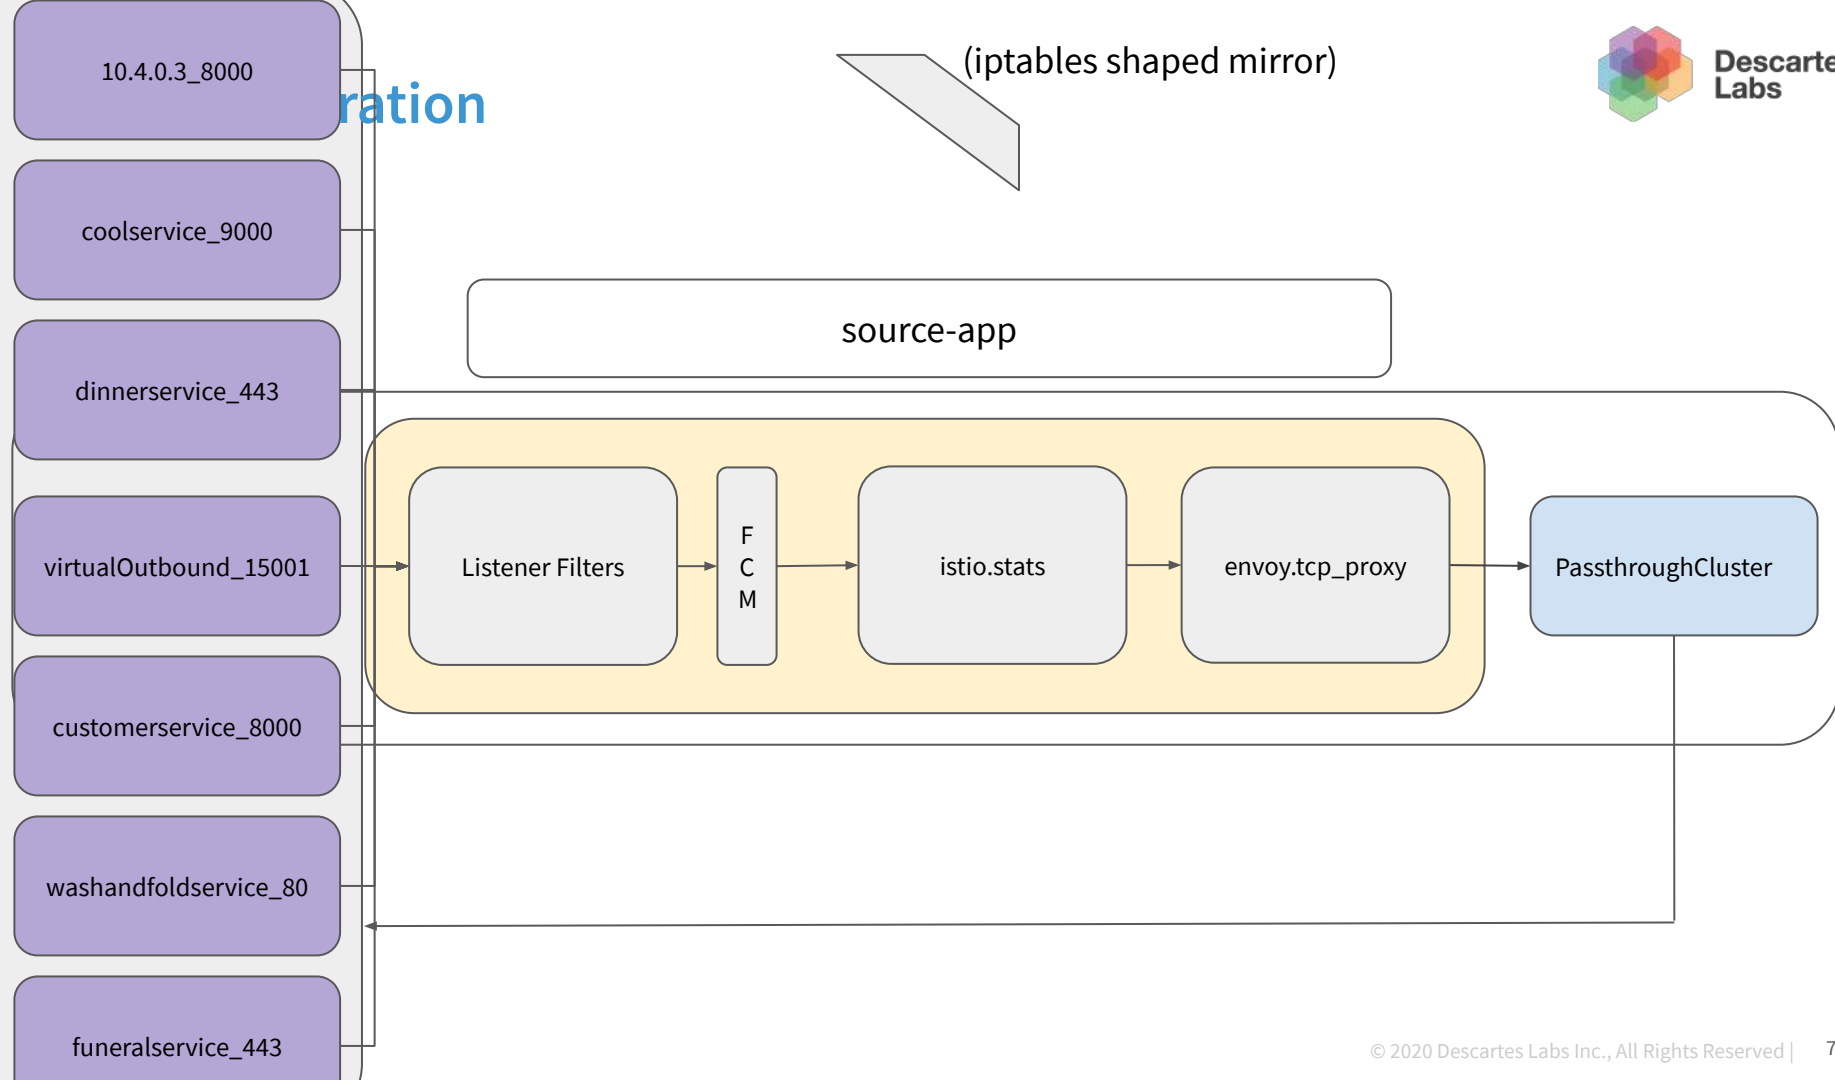


envoy - operation



ration

(iptables shaped mirror)



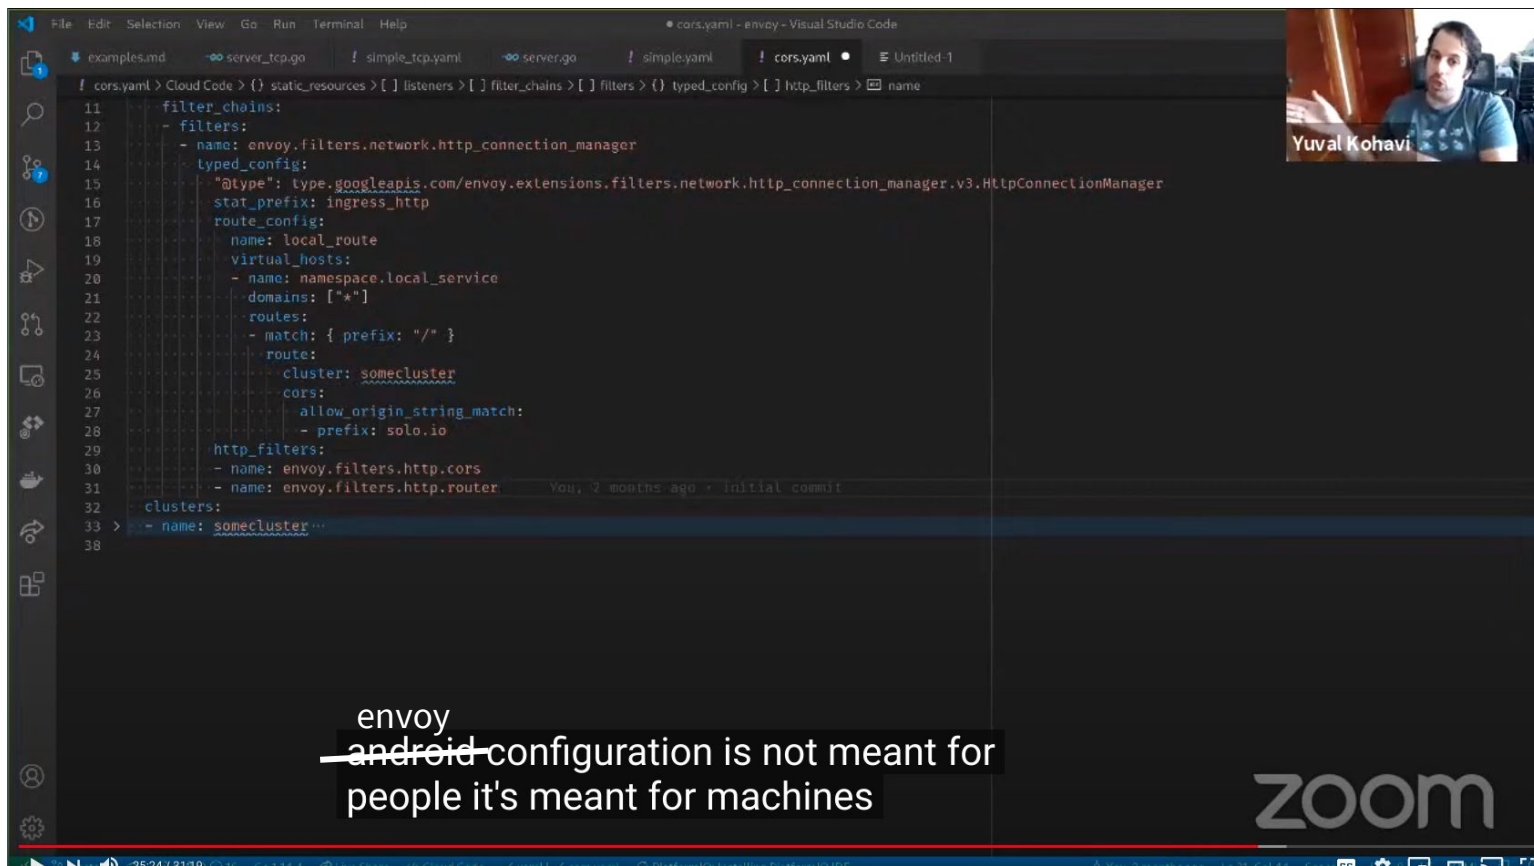
*professional SRE

Envoy: Configuration

Sidecar proxy with knobs on.



envoy - config



The screenshot shows a Zoom meeting window. On the left, a video feed of Yuval Kohavi is visible. The main area displays a VS Code editor with a file named 'cors.yaml' open. The editor shows the following YAML configuration:

```

11 filter_chains:
12   - filters:
13     - name: envoy.filters.network.http_connection_manager
14       typed_config:
15         "@type": type.googleapis.com/envoy.extensions.filters.network.http_connection_manager.v3.HttpConnectionManager
16         stat_prefix: ingress_http
17         route_config:
18           name: local_route
19           virtual_hosts:
20             - name: namespace.local_service
21               domains: ["*"]
22             routes:
23               - match: { prefix: "/" }
24                 route:
25                   cluster: somecluster
26             cors:
27               allow_origin_string_match:
28                 - prefix: solo.io
29             http_filters:
30               - name: envoy.filters.http.cors
31               - name: envoy.filters.http.router
32             clusters:
33               - name: somecluster
34 
```

Below the code editor, a text overlay reads: **envoy configuration is not meant for people it's meant for machines**. The Zoom logo is visible in the bottom right corner of the meeting window.

envoy - config


```
admin:
  access_log_path: /tmp/admin_access.log
  address:
    socket_address: { address: 127.0.0.1, port_value: 9901 }

static_resources:
  listeners:
  - name: listener_0
    address:
      socket_address: { address: 127.0.0.1, port_value: 10000 }
    filter_chains:
    - filters:
      - name: envoy.filters.network.http_connection_manager
        typed_config:
          "@type": type.googleapis.com/envoy.extensions.filters.network.http_connection_manager.v3.HttpConnectionManager
          stat_prefix: ingress_http
          codec_type: AUTO
          route_config:
            name: local_route
            virtual_hosts:
            - name: local_service
              domains: ["*"]
              routes:
              - match: { prefix: "/" }
                route: { cluster: some_service }
          http_filters:
          - name: envoy.filters.http.router
  clusters:
  - name: some_service
    connect_timeout: 0.25s
    type: STATIC
    lb_policy: ROUND_ROBIN
    load_assignment:
      cluster_name: some_service
      endpoints:
      - lb_endpoints:
        - endpoint:
            address:
              socket_address:
                address: 127.0.0.1
                port_value: 1234
```

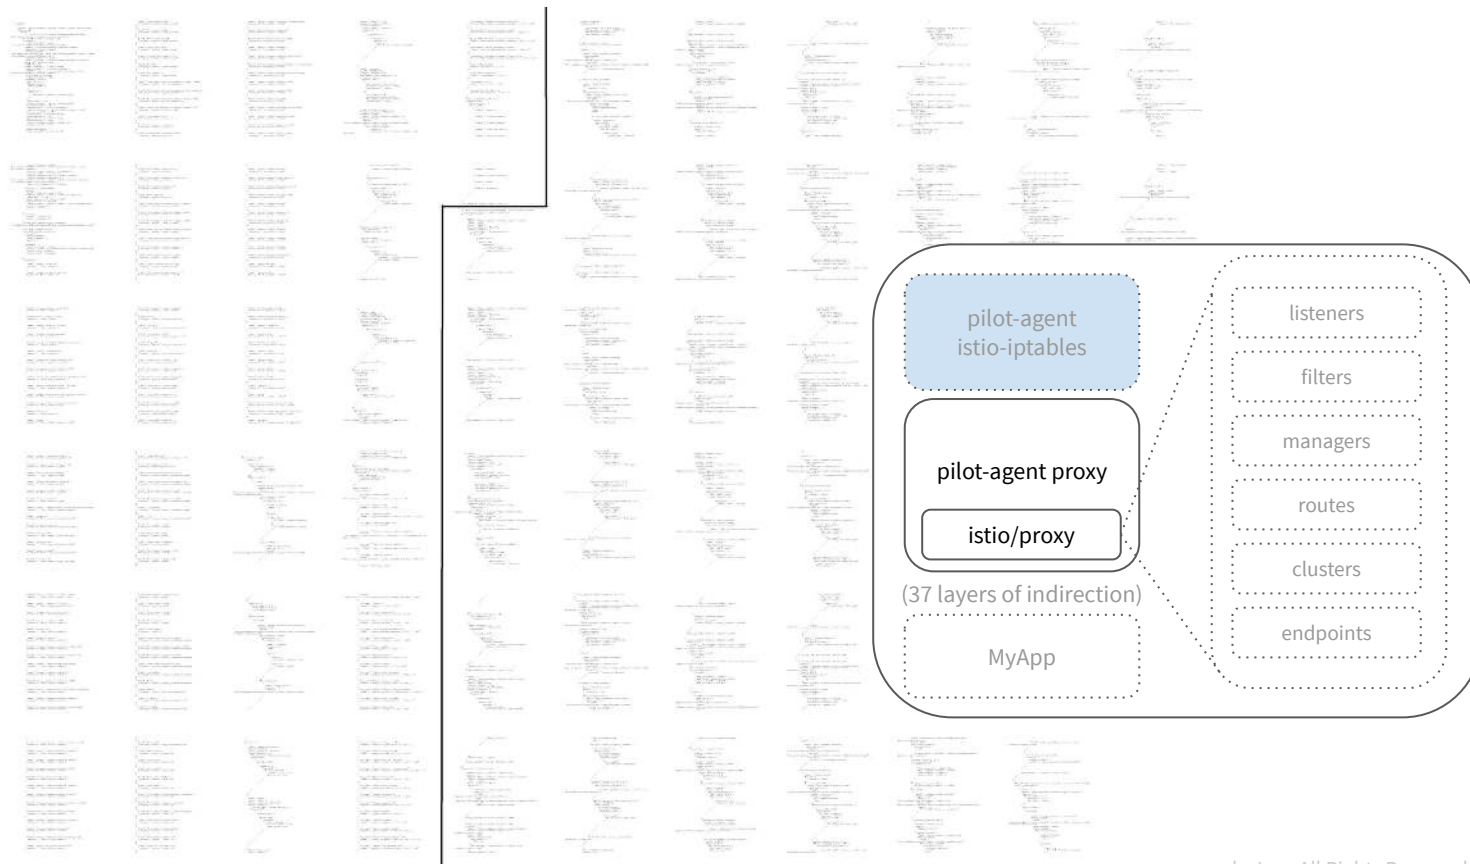
envoy - config




envoy - config - bootstrap




envoy - config - bootstrap



envoy - config - bootstrap



envoy - config - bootstrap





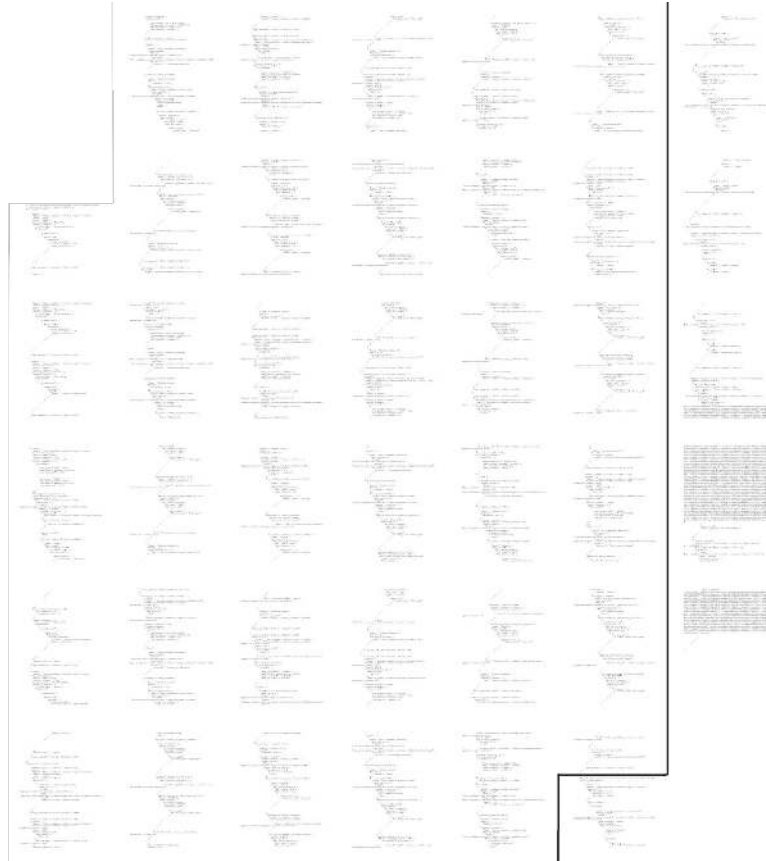
envoy - config - bootstrap




envoy - config - bootstrap



envoy - config



envoy - config - clusters



envoy - config



envoy - config - listeners




envoy - config - routes



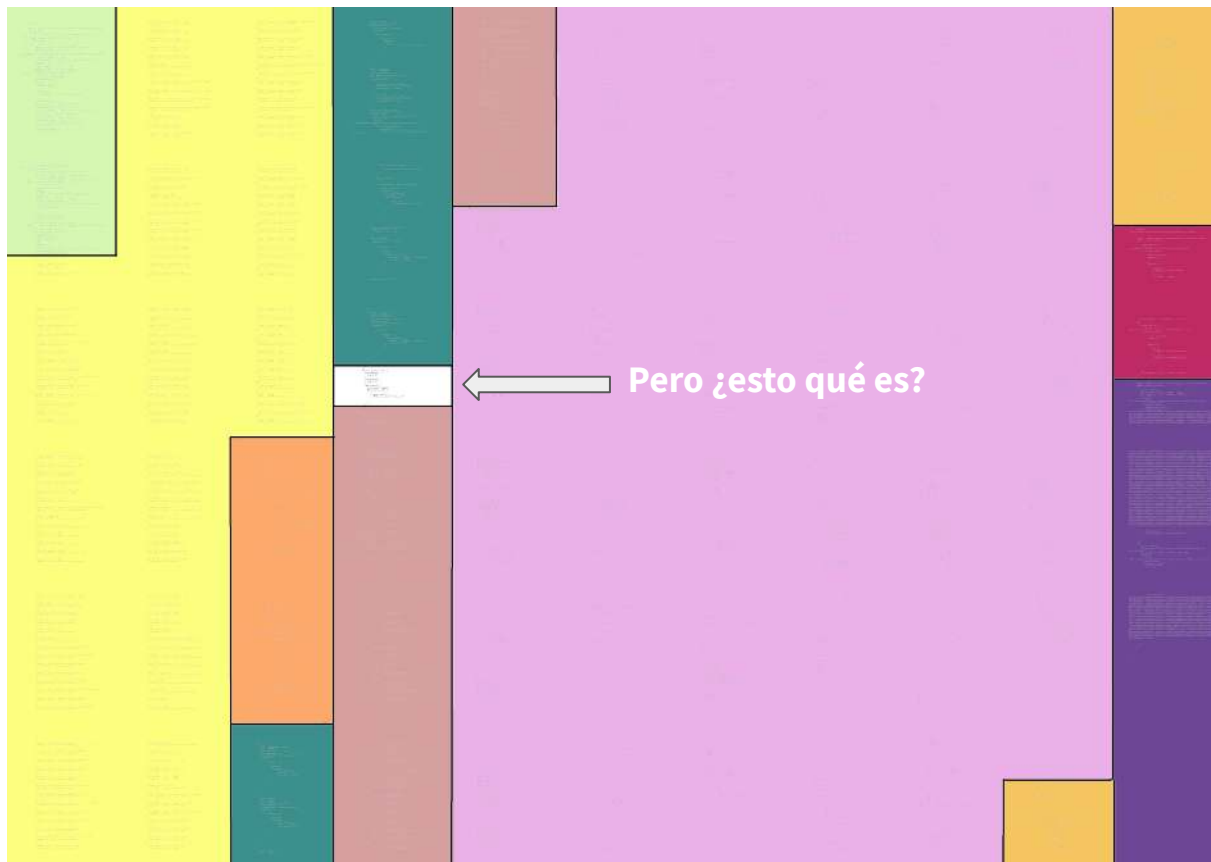
envoy - config - secrets



envoy - config



envoy - config



envoy - config

```
▼ "dynamic_resources": {  
  ▼ "lds_config": {  
    "ads": {}  
  },  
  ▼ "cds_config": {  
    "ads": {}  
  },  
  ▼ "ads_config": {  
    "api_type": "GRPC",  
    ▼ "grpc_services": [  
      ▼ {  
        ▼ "envoy_grpc": {  
          "cluster_name": "xds-grpc"  
        }  
      }  
    ]  
  }  
},
```

envoy - config

```
▼ "static_resources": {  
  ▶ "listeners": [...], // 2 items  
  ▼ "clusters": [  
    ▶ {...}, // 4 items  
    ▶ {...}, // 4 items  
    ▶ {...}, // 5 items  
    ▼ {  
      "name": "xds-grpc",  
      "type": "STRICT_DNS",  
      "connect_timeout": "1s",  
      "max_requests_per_connection": 1,  
      ▶ "circuit_breakers": {...}, // 1 item  
      "http2_protocol_options": {},  
      "dns_lookup_family": "V4_ONLY",  
      ▶ "transport_socket": {...}, // 2 items  
      "upstream_connection_options": {...}, // 1 item  
      ▼ "load_assignment": {  
        "cluster_name": "xds-grpc",  
        ▼ "endpoints": [  
          ▼ {  
            ▼ "lb_endpoints": [  
              ▼ {  
                ▼ "endpoint": {  
                  ▼ "address": {  
                    ▼ "socket_address": {  
                      "address": "istiod.istio-system.svc",  
                      "port_value": 15012  
                    }  
                  }  
                }  
              }  
            ]  
          }  
        ]  
      }  
    },  
    ],  
    "respect_dns_ttl": true  
  ],  
  ▶ {...} // 6 items  
},  
1  
},
```

[FIN, ACK]



Cuteness to soothe the spirit.

Metadata

This document URL: <https://bit.ly/3pFMT78>

References include links to [Istio Community documents in gdrive](#). To access these documents you need to be a member of [this google group](#).

Istio Acronyms

CR - Custom Resource - a specific instance of the defined spec

CRD - Custom Resource Definition - the spec for a Custom Resource

SD - Service Discovery

MX - Metadata Exchange - Istio specific Envoy extension providing “mixer like” attribute data without mixer

MCP - Mesh Configuration Protocol - how Istio components communicate

Envoy Acronyms

FCM - Filter Chain Match

ALS - Access Log Service

xDS - x Discovery Service (like the x in DirectX)

- Current version as of writing is v3, v2 supported but deprecated
- CDS cluster discovery service
- LDS listener discovery service
- RDS route discovery service
- SDS secret discovery service
- There are a few other more obscure ones in here too

[UDPA - Universal Data Plane API \(XDS v4\)](#)

UDPA-TP - UDPA Transport Protocol

UDPA-DM - UDPA Data Model

ORCA - Open Request Cost Aggregation - load reporting

DPLB - Data Plane Load Balancer

- This appears to refer to xDS / UDPA clients, like Envoy itself

Terminology

Upstream - where envoy sends requests to

Downstream - where envoy receives requests from

Management Server - xDS Server AKA control plane AKA istiod