



tetrate



THE ENTERPRISE SERVICE MESH COMPANY



WELCOME



Eitan Suez



Peter Jausovec





Meet the workshop instructors

ISTIO 0 TO 60

- Eitan Suez
- Peter Jausovec



About this workshop

ISTIO 0 TO 60

- Get up and running quickly with Istio
- Prerequisites
 - Basic understanding of Kubernetes, Docker, and Linux command line.



Logistics

ISTIO 0 TO 60

- Duration: 2.5 hours
 - 15 minute break half-way through the training
- Communication support through slack channel
 - <https://istio.slack.com/>
 - channel #istiocon-workshop-tetrate



Learn by doing

APPROACH

- *Lab-driven* training
- Minimize the use of slides
- We request your active participation
 - Please ask questions



Schedule



Lab environments

Configure access to your lab environment

What problems does Istio address?

A high-level introduction and overview of Istio

Install Istio

Get your cluster installed and configured with Istio

Sidecar injection and the app under test

Deploy a simple application to the mesh, and expose it with Ingress

Observability, Security, Traffic shifting

Three labs that cover the essential cross-cutting concerns that Istio addresses

Summary



Workshop labs

<https://tetratelabs.github.io/istio-0to60/>



Environments

IN A NUTSHELL

- BYOK (Bring Your Own Kubernetes) - if possible
- Fallback:
 - We have a few Kubernetes clusters provisioned in GCP that we can give you access to, "while supplies last."
 - Let us know via slack if you need a K8S cluster and we'll assign you a cluster.
 - The first lab contains instructions for accessing your cluster on GCP.
- We will begin with a lab to get your environment setup



LAB

Setup your environment

<https://tetratelabs.github.io/istio-0to60/environment/>



Meet Istio



PROBLEM

IT shift to a modern distributed architecture
has left enterprises unable to monitor,
connect, manage, and secure their services in
a consistent way.



Problem

MODERN DISTRIBUTED ARCHITECTURE

- Container based services
- Deployed into dynamic environments
- Composed via the network



Problem

MONITOR

- Understand what's actually happening in your deployment through basic tools:
 - Metrics
 - Logs
 - Tracing



Problem

CONNECT

- Get network out of the application
 - **Service discovery**
 - **Resiliency**
 - | Retries, outlier detection, circuit breaking, timeouts, etc.
 - **Load balancing**
 - | Client side



Problem

MANAGE

- Control which requests are allowed, and how and where they flow
 - **Fine-grained traffic control**
 - | L7, not L4! Route by headers, destination or source, etc.
 - **Policy on requests**
 - | Authentication, rate limiting, arbitrary policy based on L7 metadata



Problem

SECURE

- Elevate security out of the network
 - Service-to-service authentication
 - Workload identity (L7)



Service Mesh

WHAT IS SERVICE MESH

- Service mesh moves these facets out of the application for better division of labor and...
 - **Consistency across fleets**
 - **Centralized control**
 - **Ease of change**
 - | Update configurations without redeployment

What is Istio?

Istio is a platform to **monitor**, **secure**,
connect and **manage** services
consistently



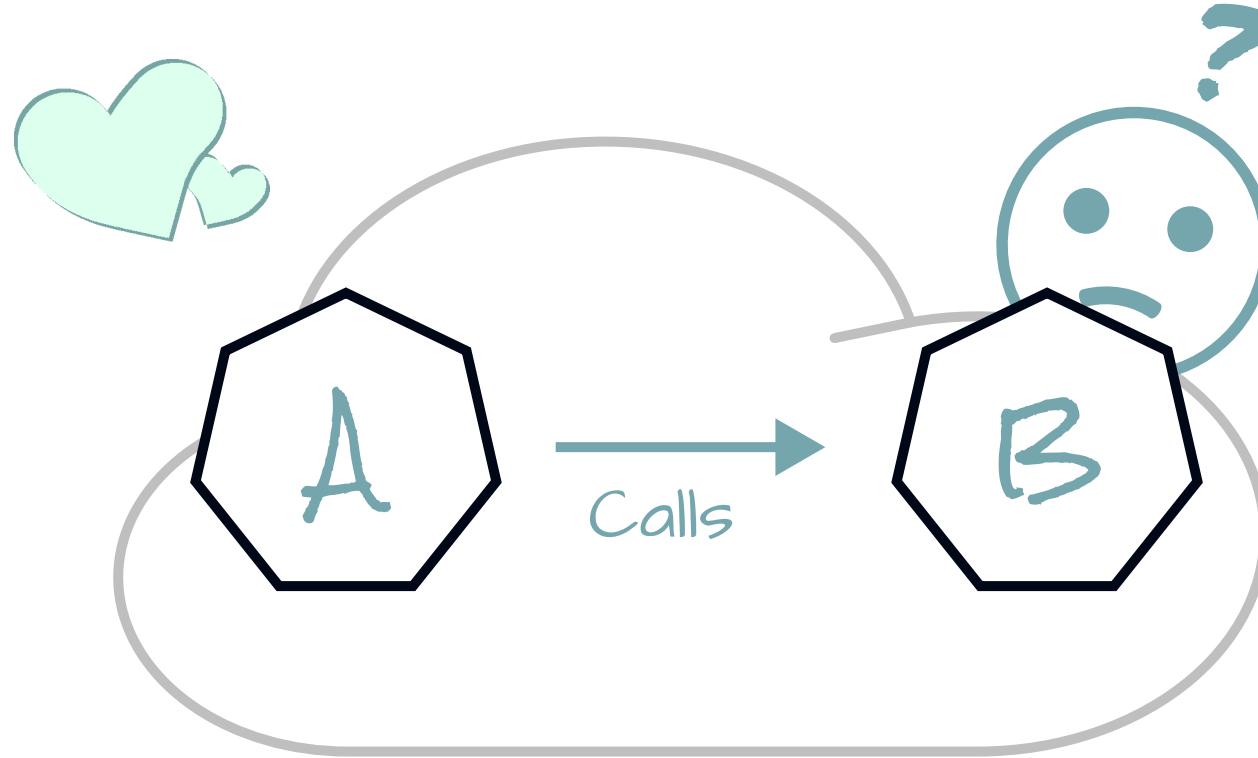
LAB

Installing Istio

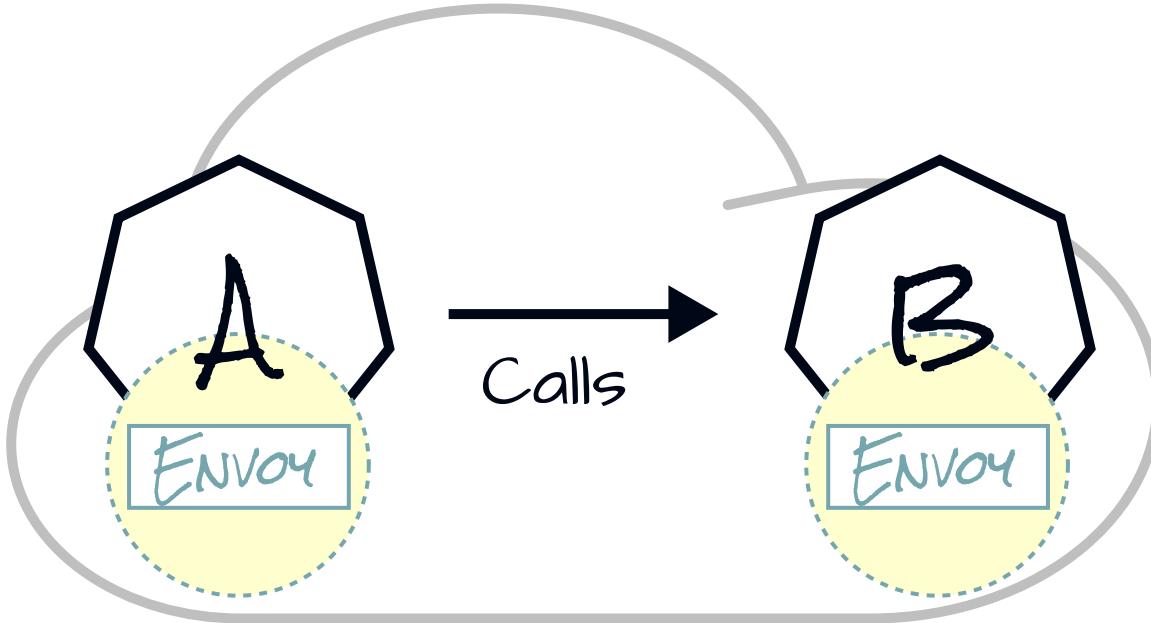
<https://tetratelabs.github.io/istio-0to60/install/>



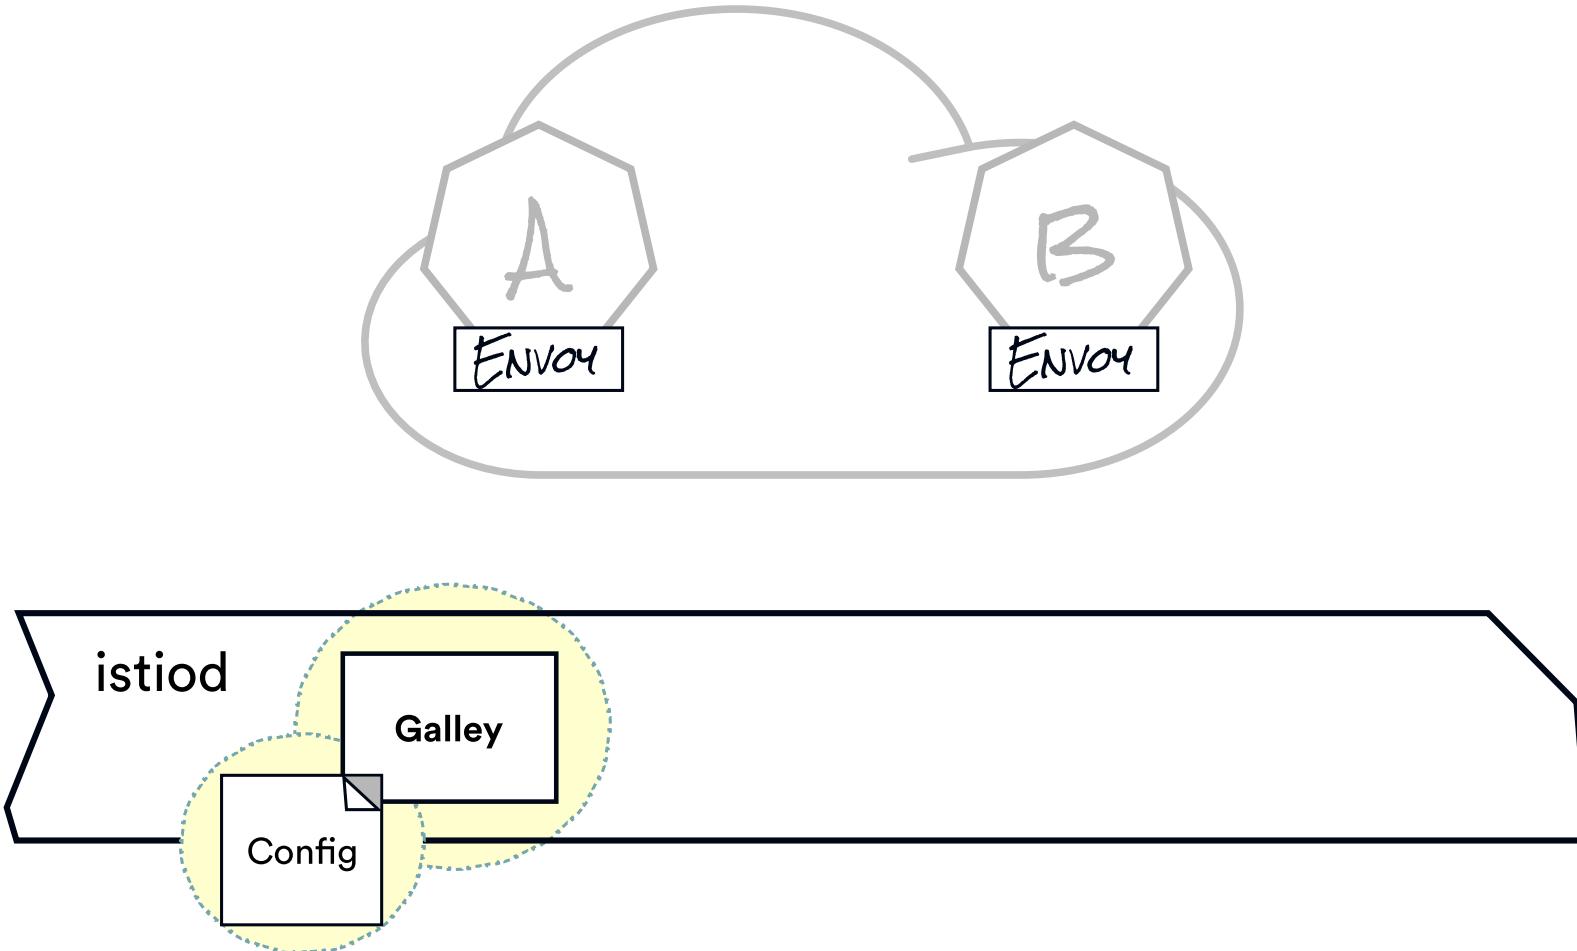
Istio Architecture



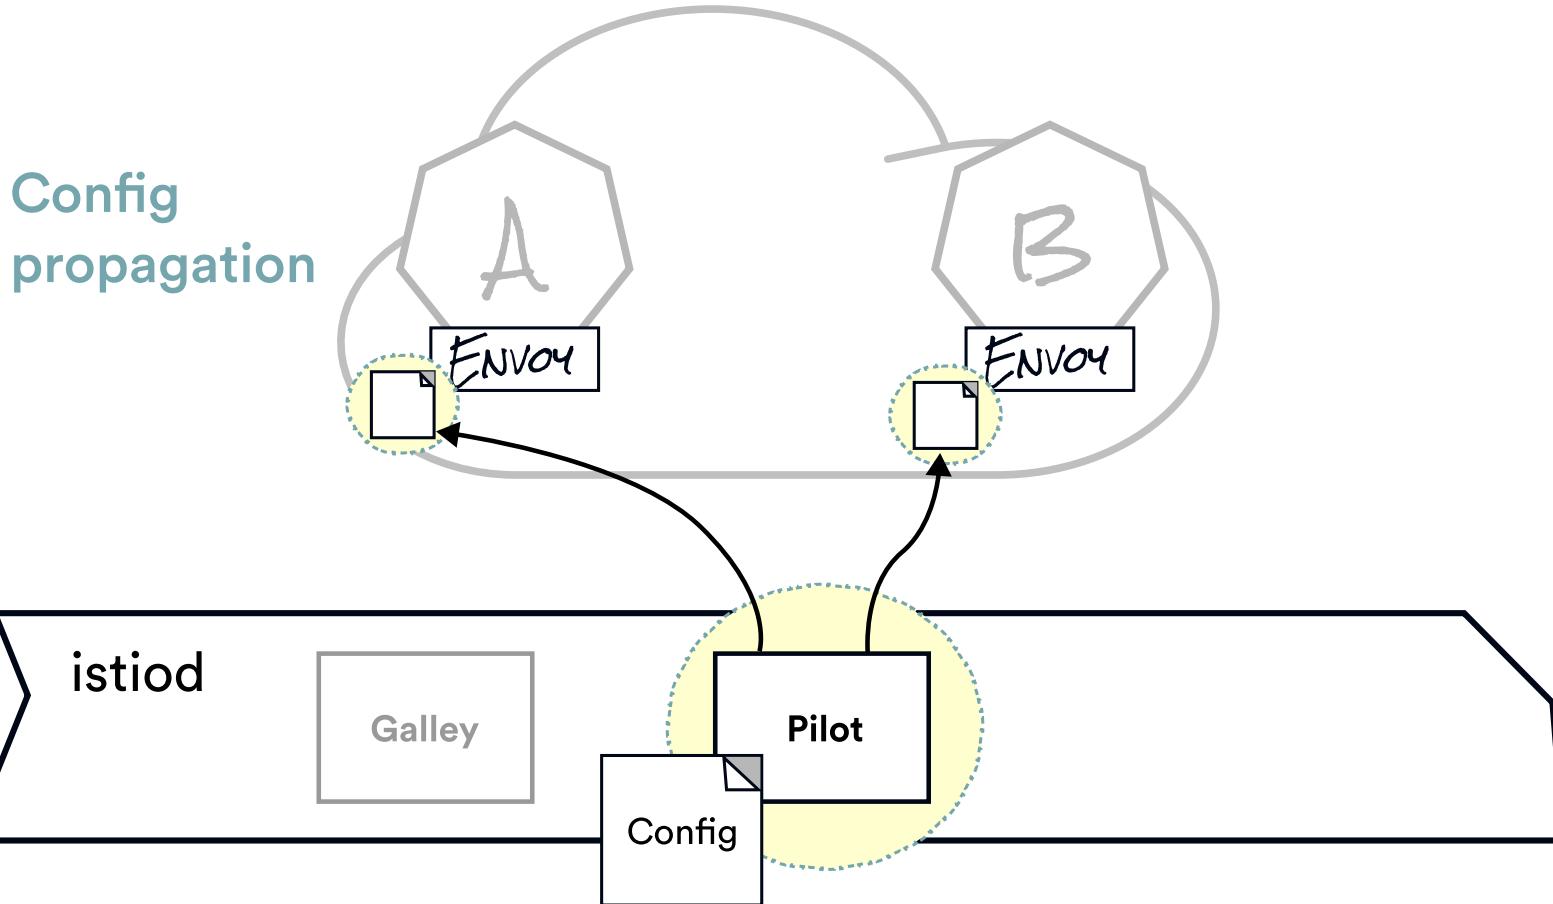
Story as old as time:
Service A meets service B...



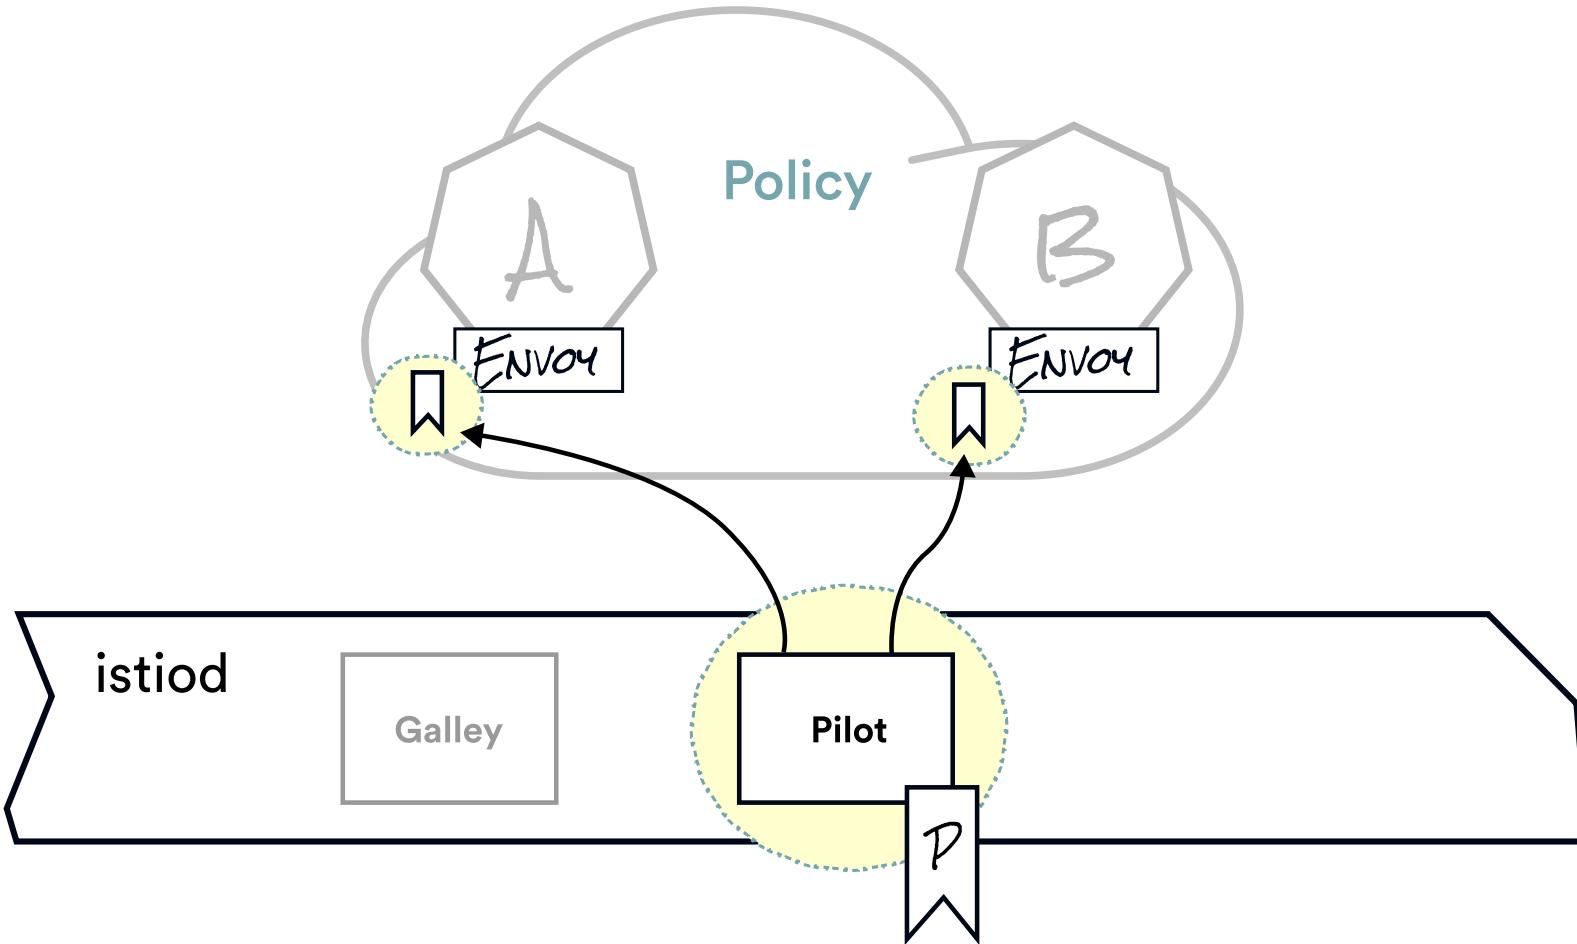
Deploy a proxy (Envoy) beside your application ("sidecar deployment")



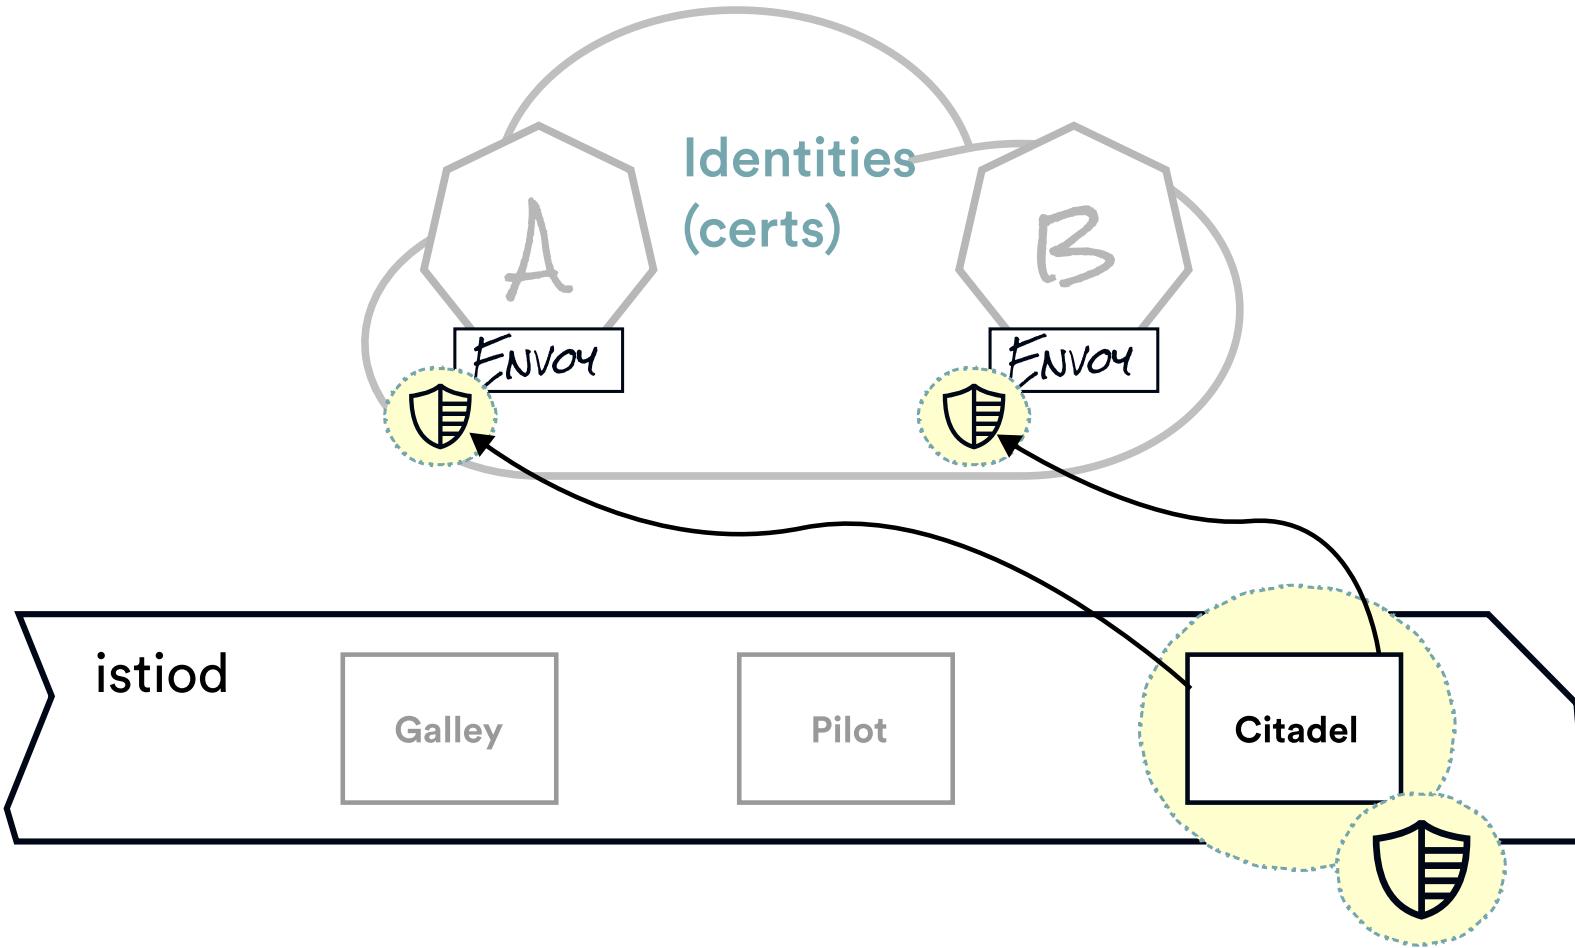
First logical component is Galley, which is responsible for validating incoming config.



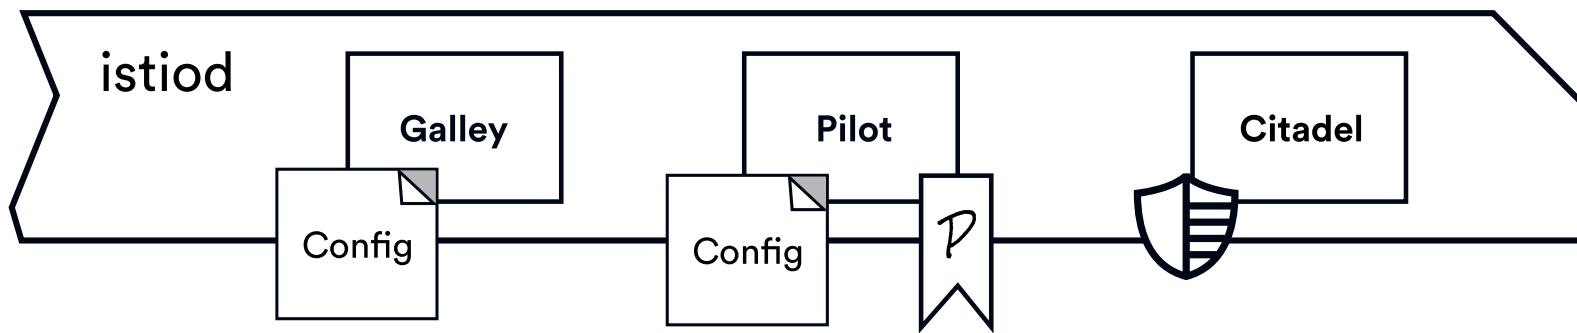
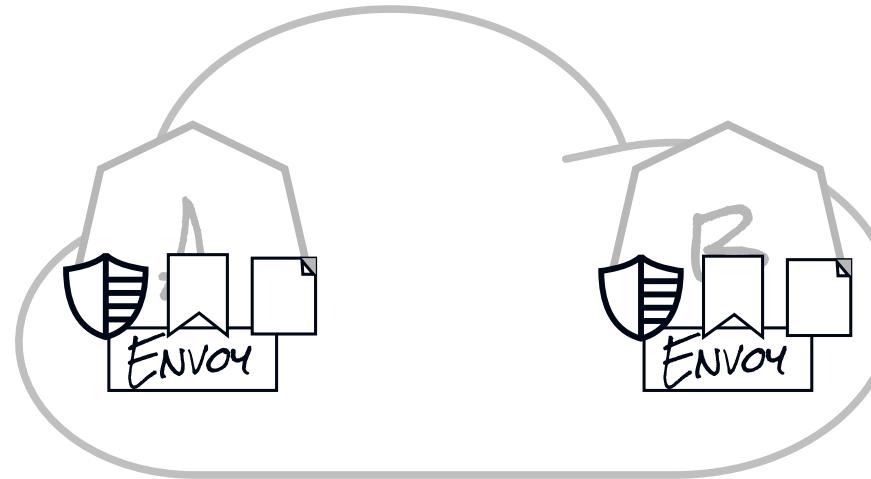
Pilot distributes the validated networking configuration to each Envoy



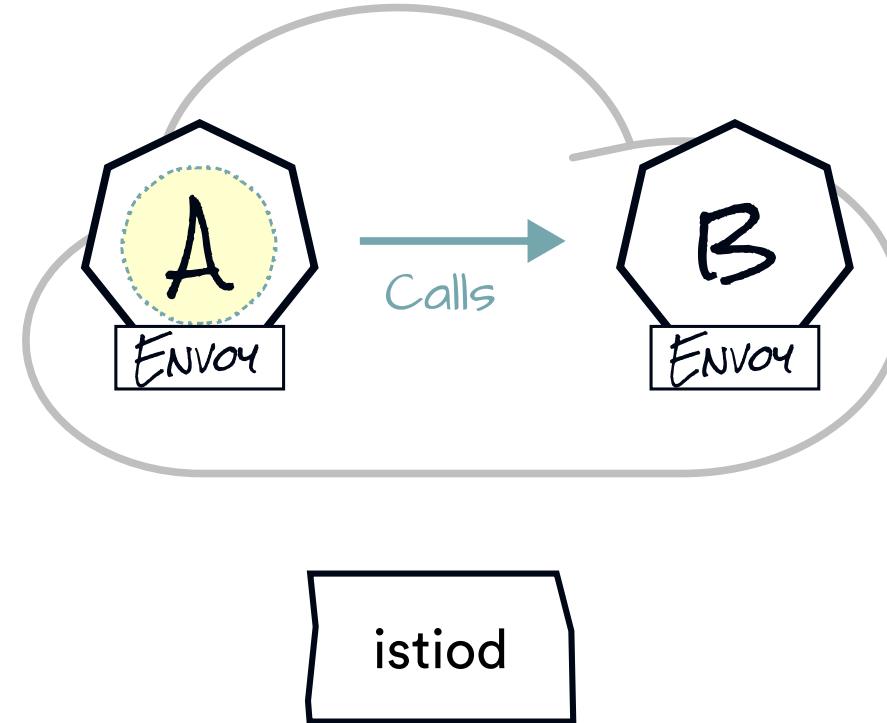
...and Pilot also distributes policy



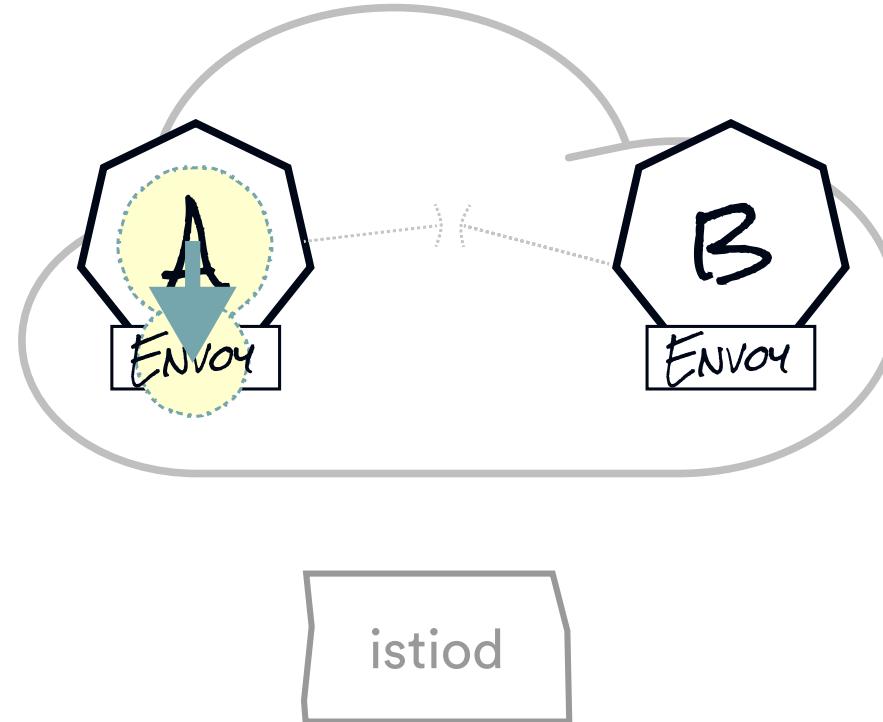
Citadel assigns SPIFFE identities to enable secure communication



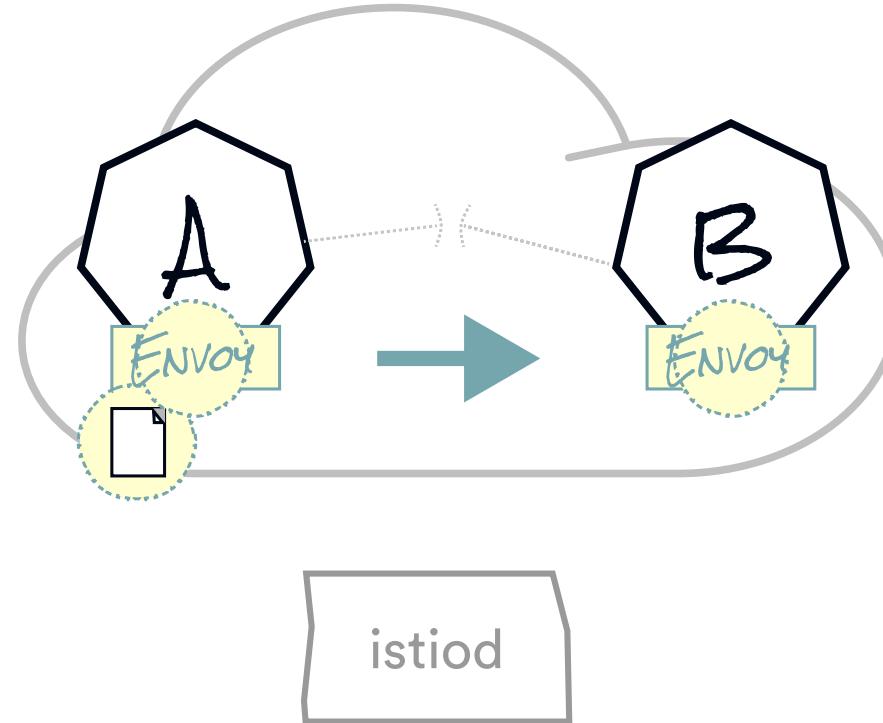
Control plane - Istiod



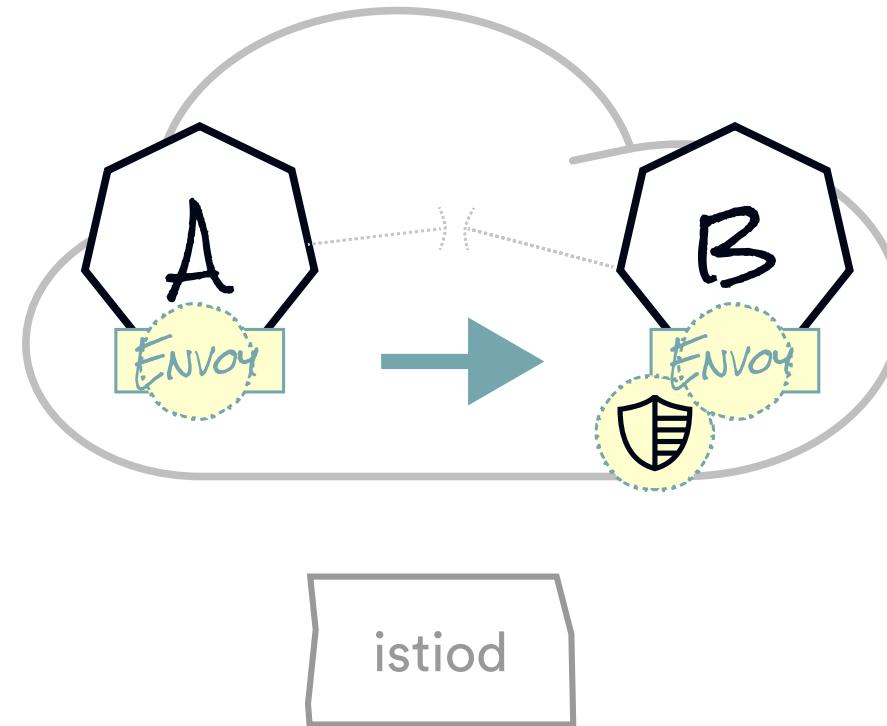
Now, let's track that call



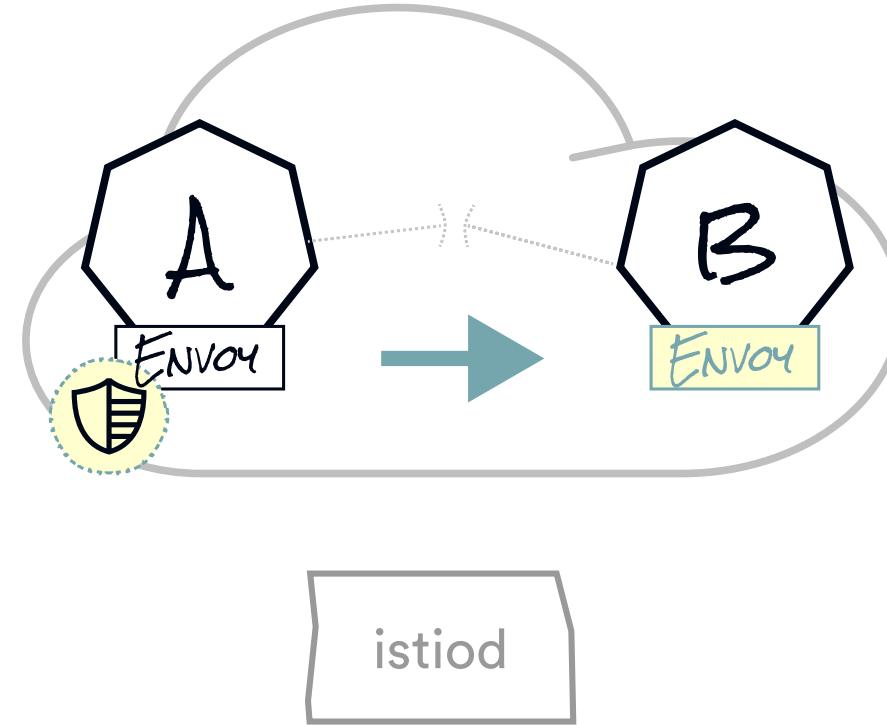
Envoy intercepts it



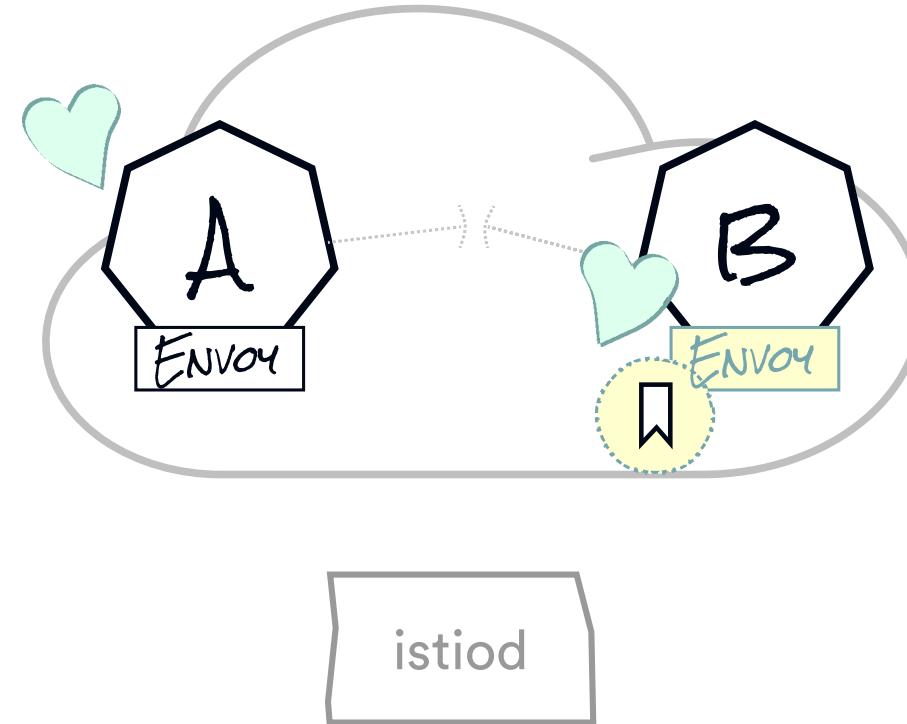
Uses the configuration to pick a new destination



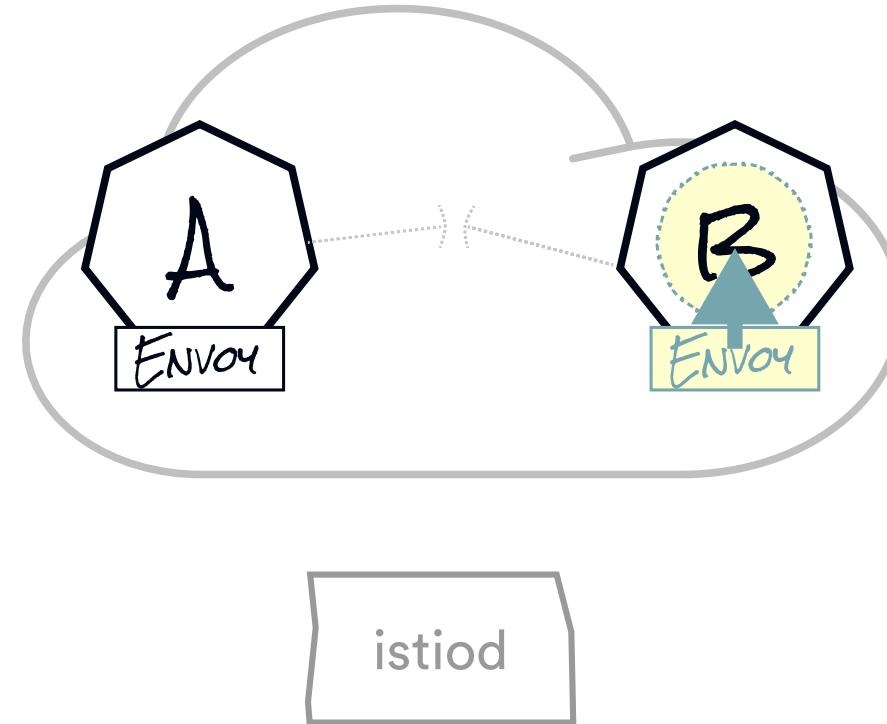
Verifies the destination's identity



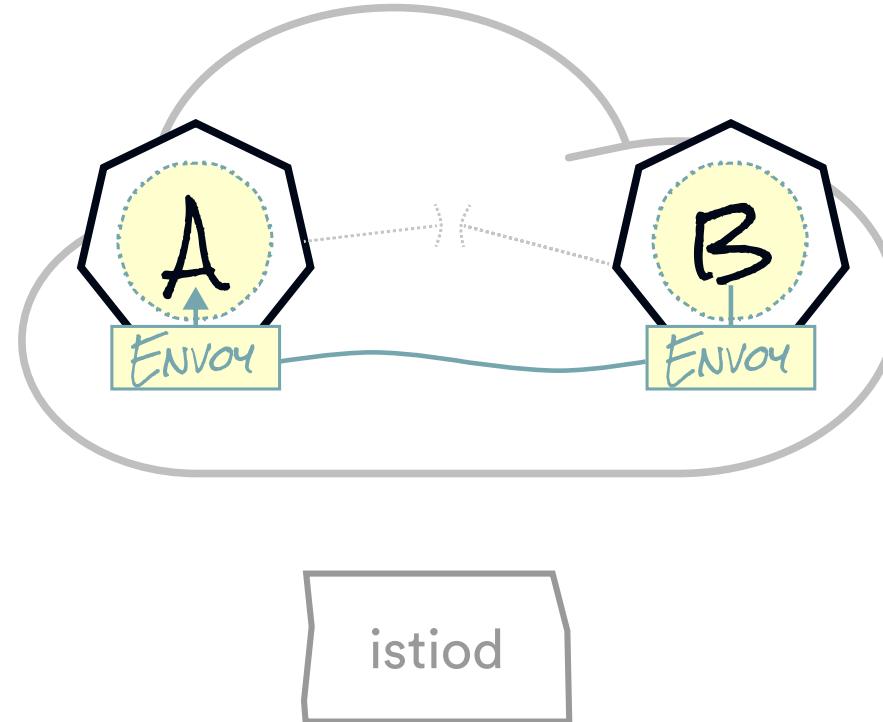
The receiving Envoy checks the sender's identity



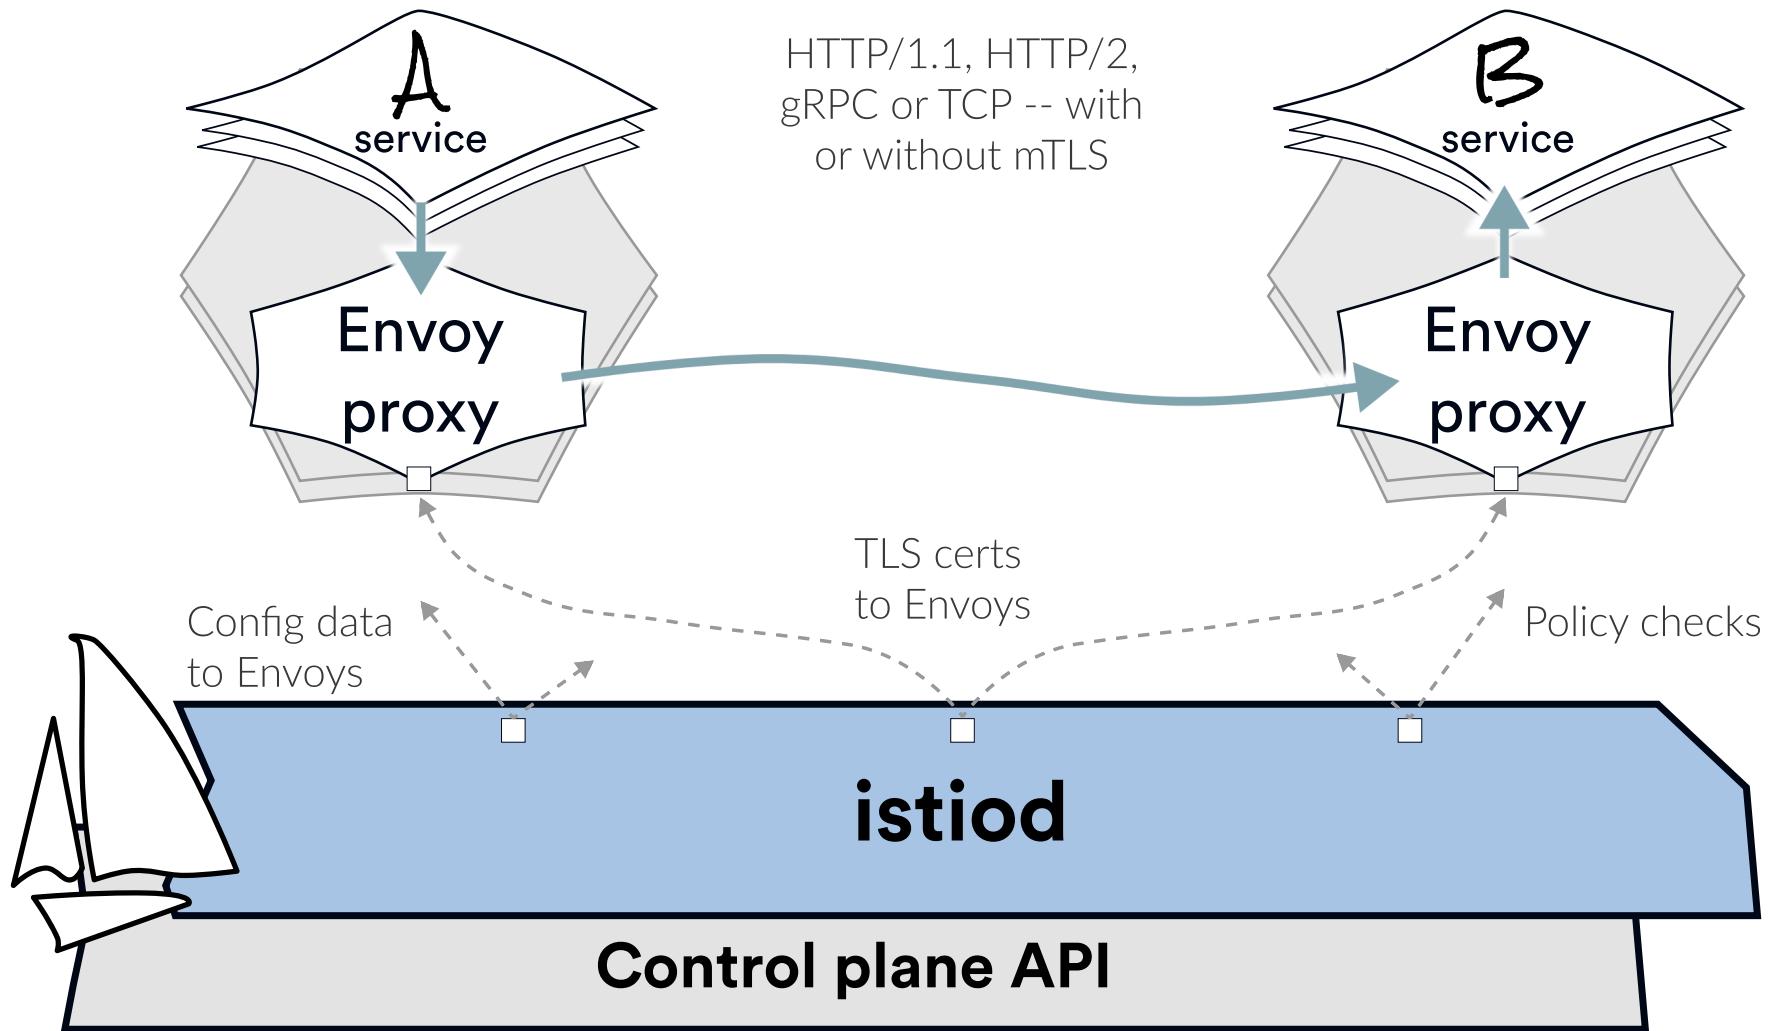
The receiving Envoy checks policy



Envoy hands the request to B



B answers



Istiod - Control plane Manager, handling



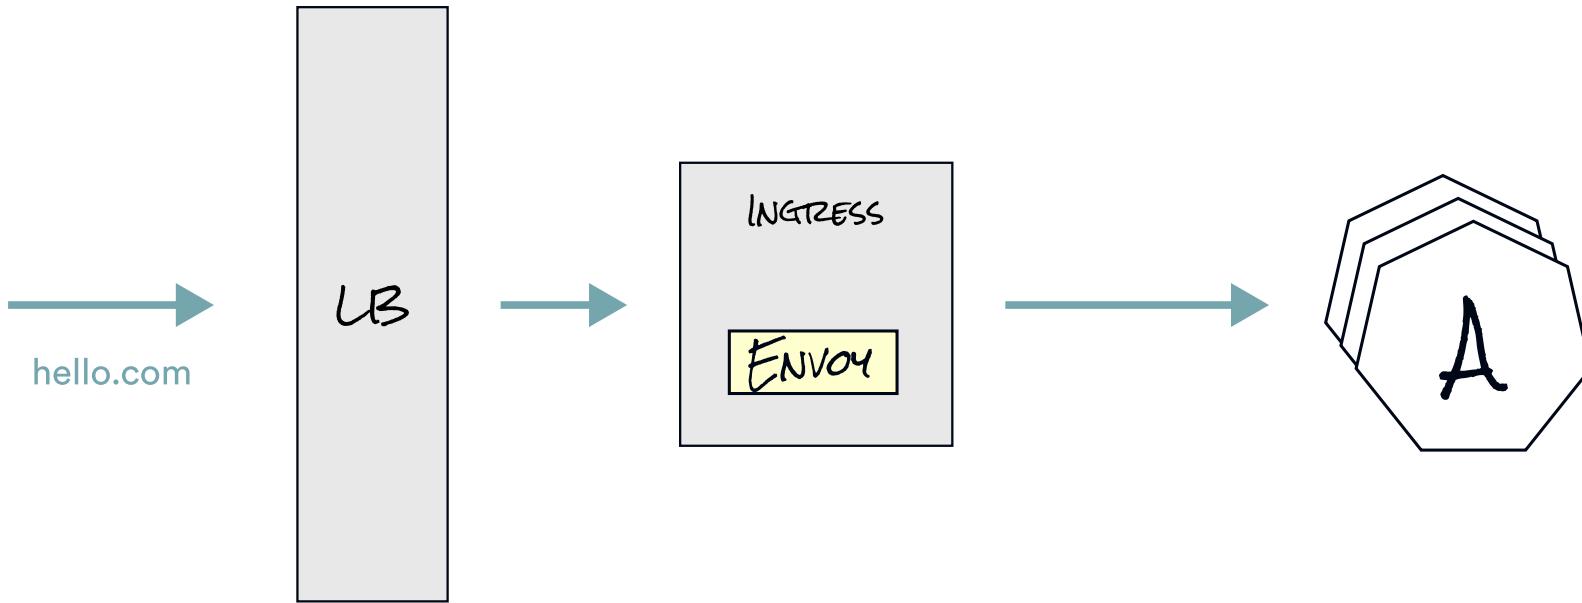
LAB

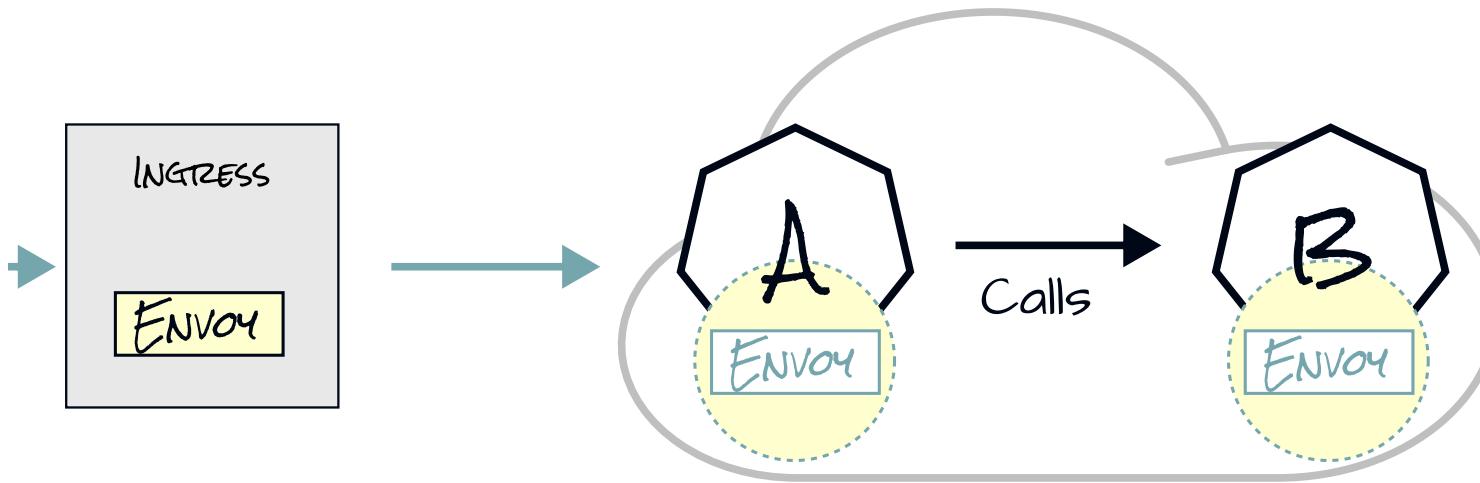
The application

<https://tetratelabs.github.io/istio-0to60/the-app/>



Ingress





What can we say about Ingress?

1. Ingress Pod runs in `istio-system` namespace
2. Configured logically in Kubernetes with the `Gateway` custom resource
3. Routing configured separately with `VirtualService` custom resource



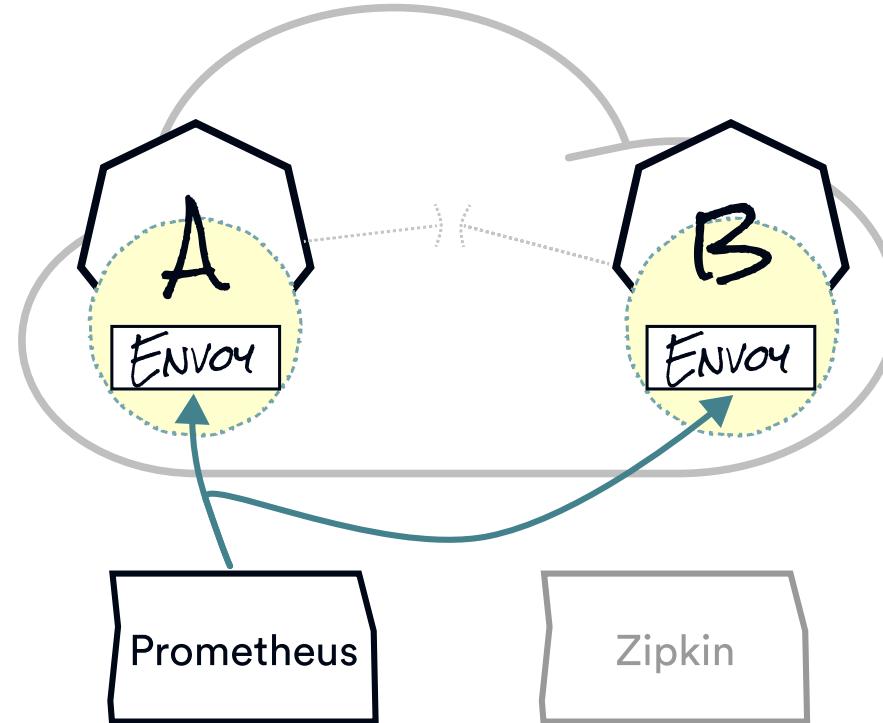
LAB

Ingress gateway

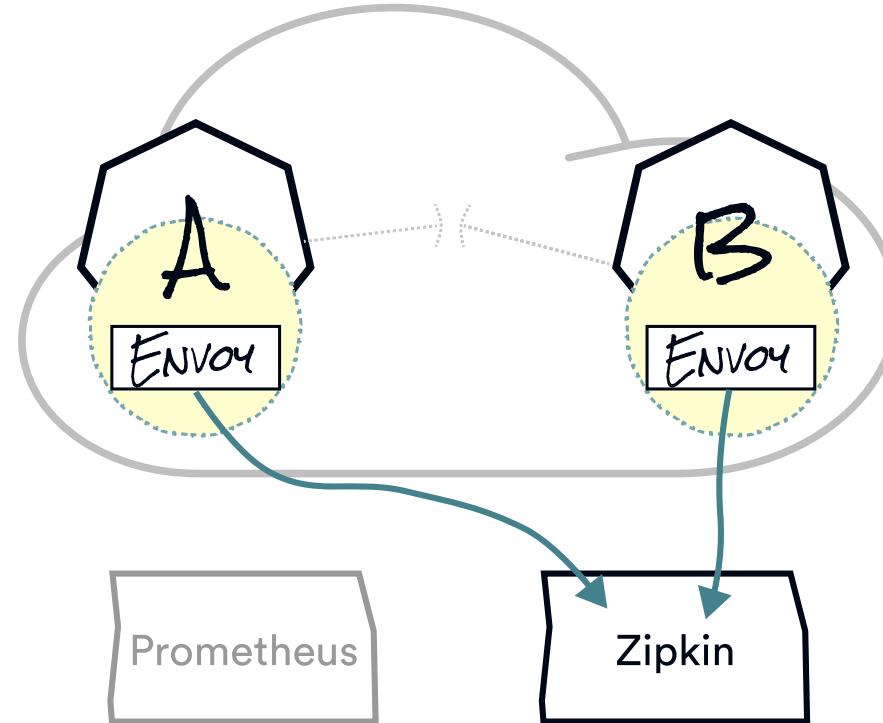
<https://tetratelabs.github.io/istio-0to60/ingress/>



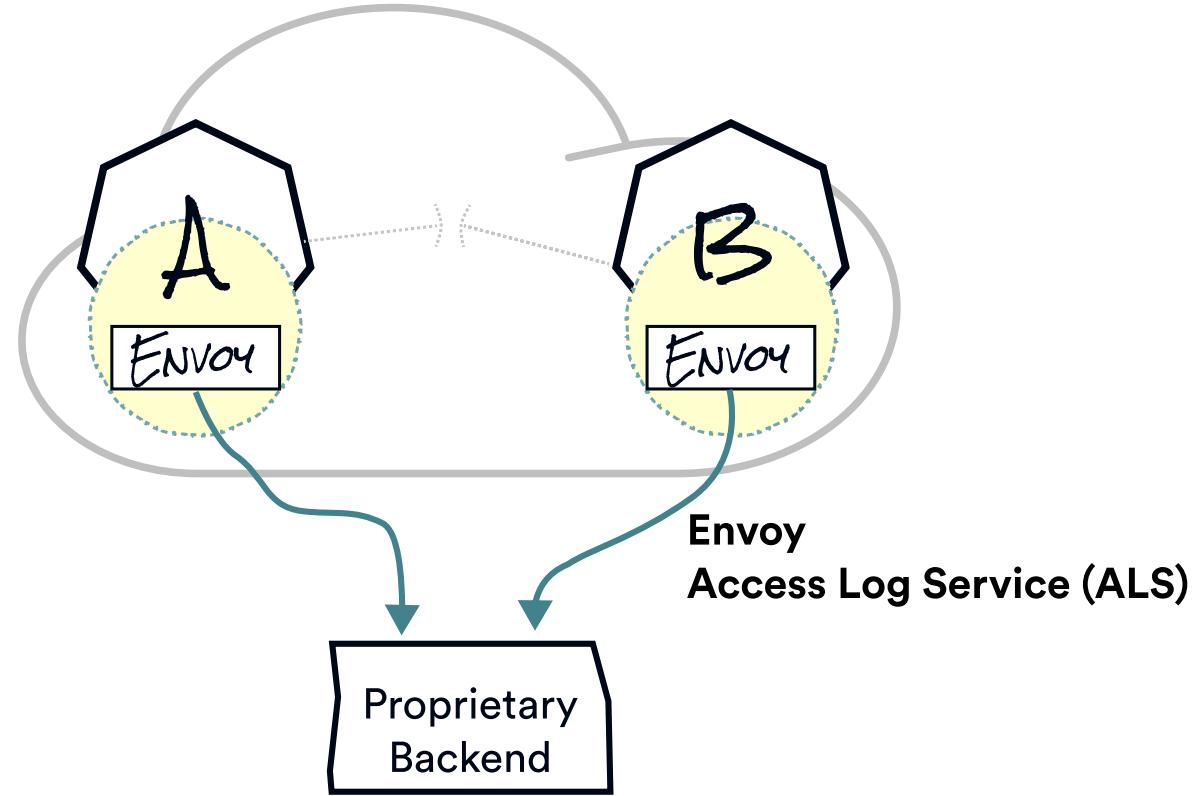
Observability



Metrics to Prometheus



Traces to Zipkin



Whatever you want via the Access Log Service



LAB

Observability

<https://tetratelabs.github.io/istio-0to60/dashboards/>



Secure your environment



Identifying workloads

AUTHENTICATION (AUTHN)

- Authn - all about the principal
- Each workload is assigned a unique identity that it uses to communicate with other workloads
 - Kubernetes = Istio uses service accounts



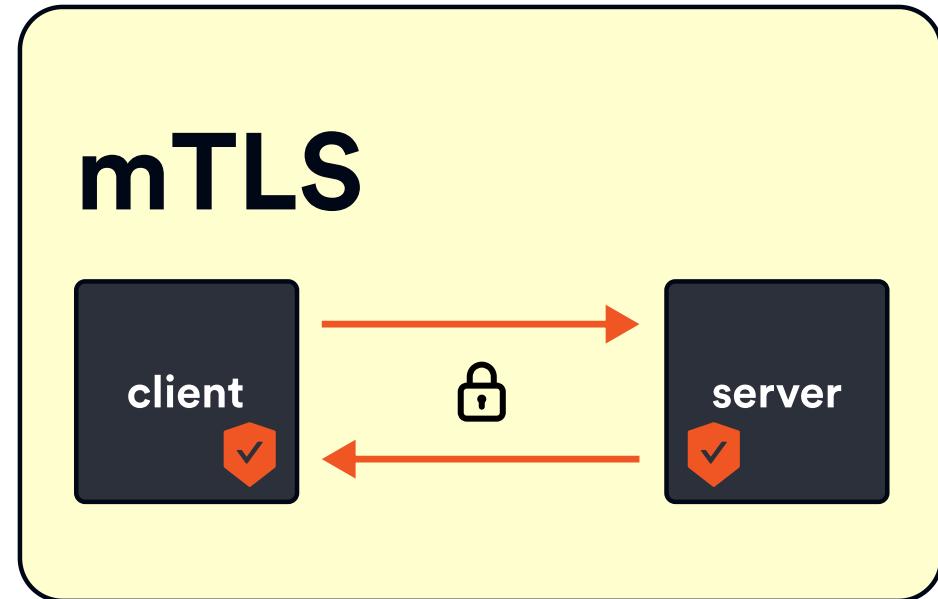
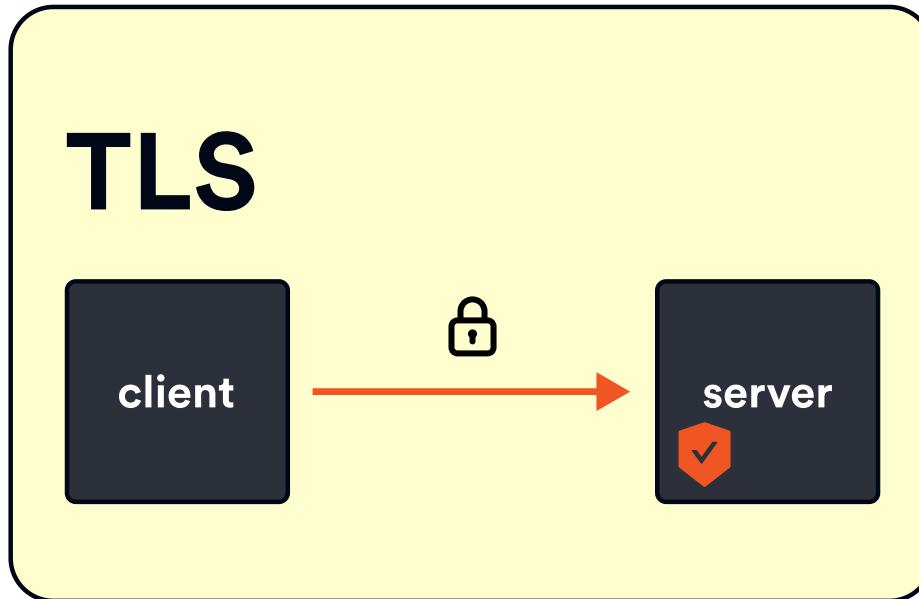
SPIFFE overview

SECURE PRODUCTION IDENTITY FRAMEWORK FOR EVERYONE

- X.509 certificate (from SA) + SPIFFE spec = **IDENTITY**
- SPIFFE is a spec that describes:
 - A **naming scheme** for workload identities
 - `spiffe://cluster.local/ns/default/sa/my-sa`
 - How to **encode those names** into a X.509 certificate
 - How a client **Validates an X.509** certificate to authenticate the SPIFFE identity inside of it

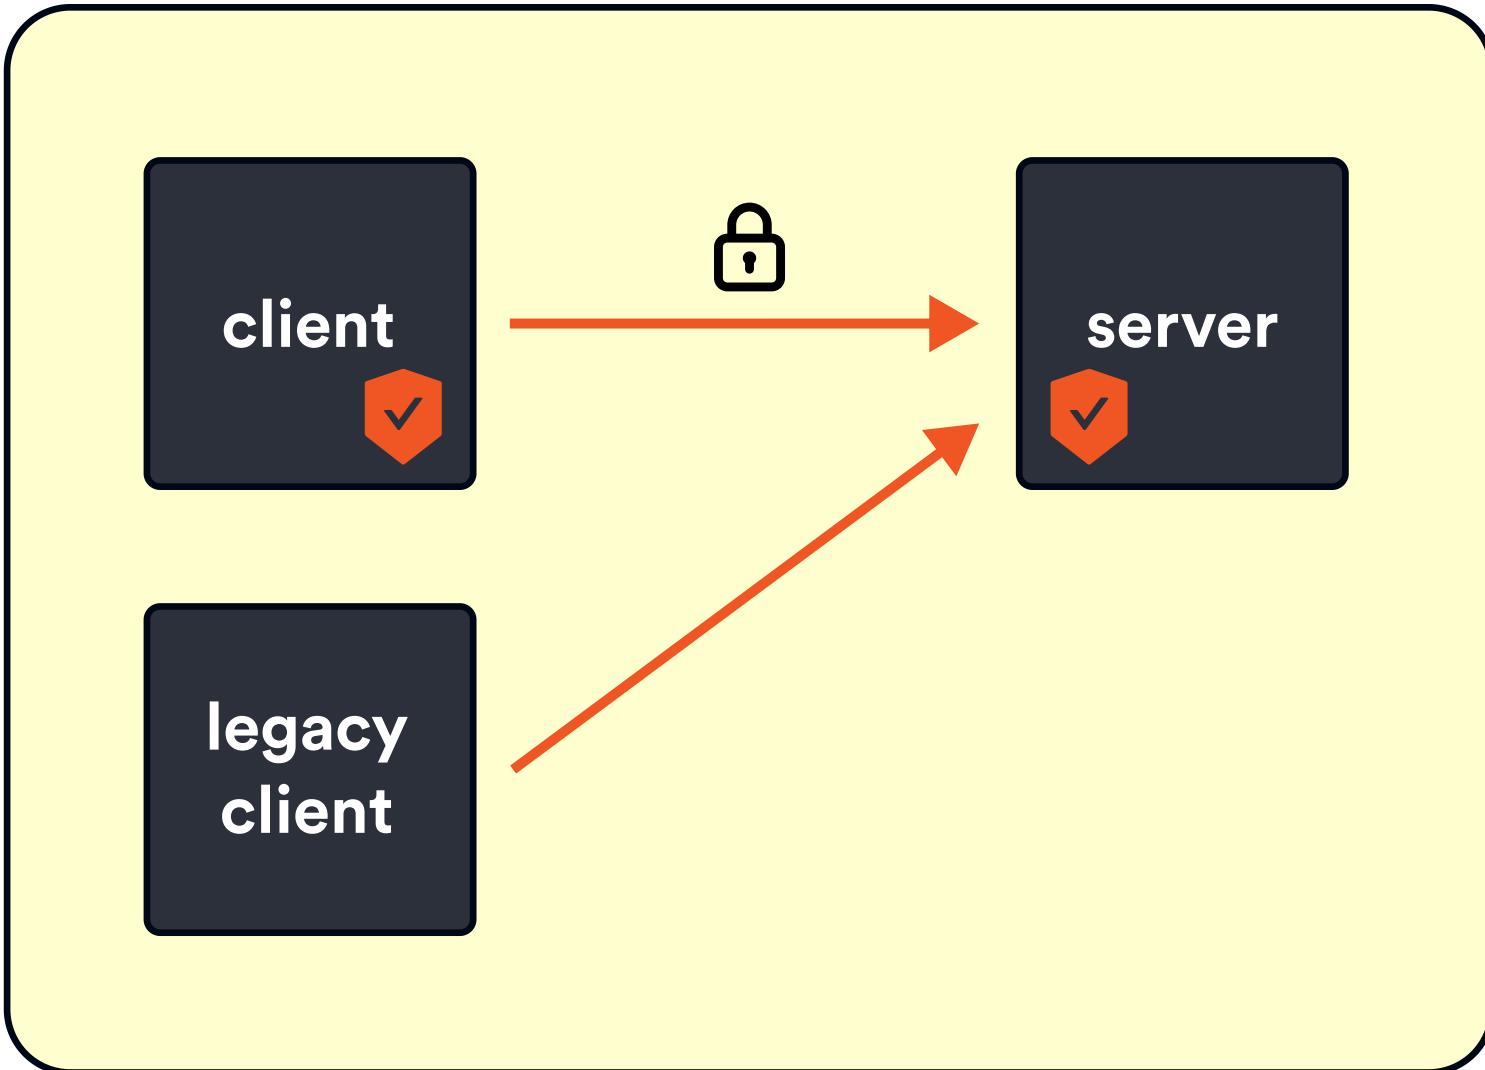


Mutual TLS (mTLS)



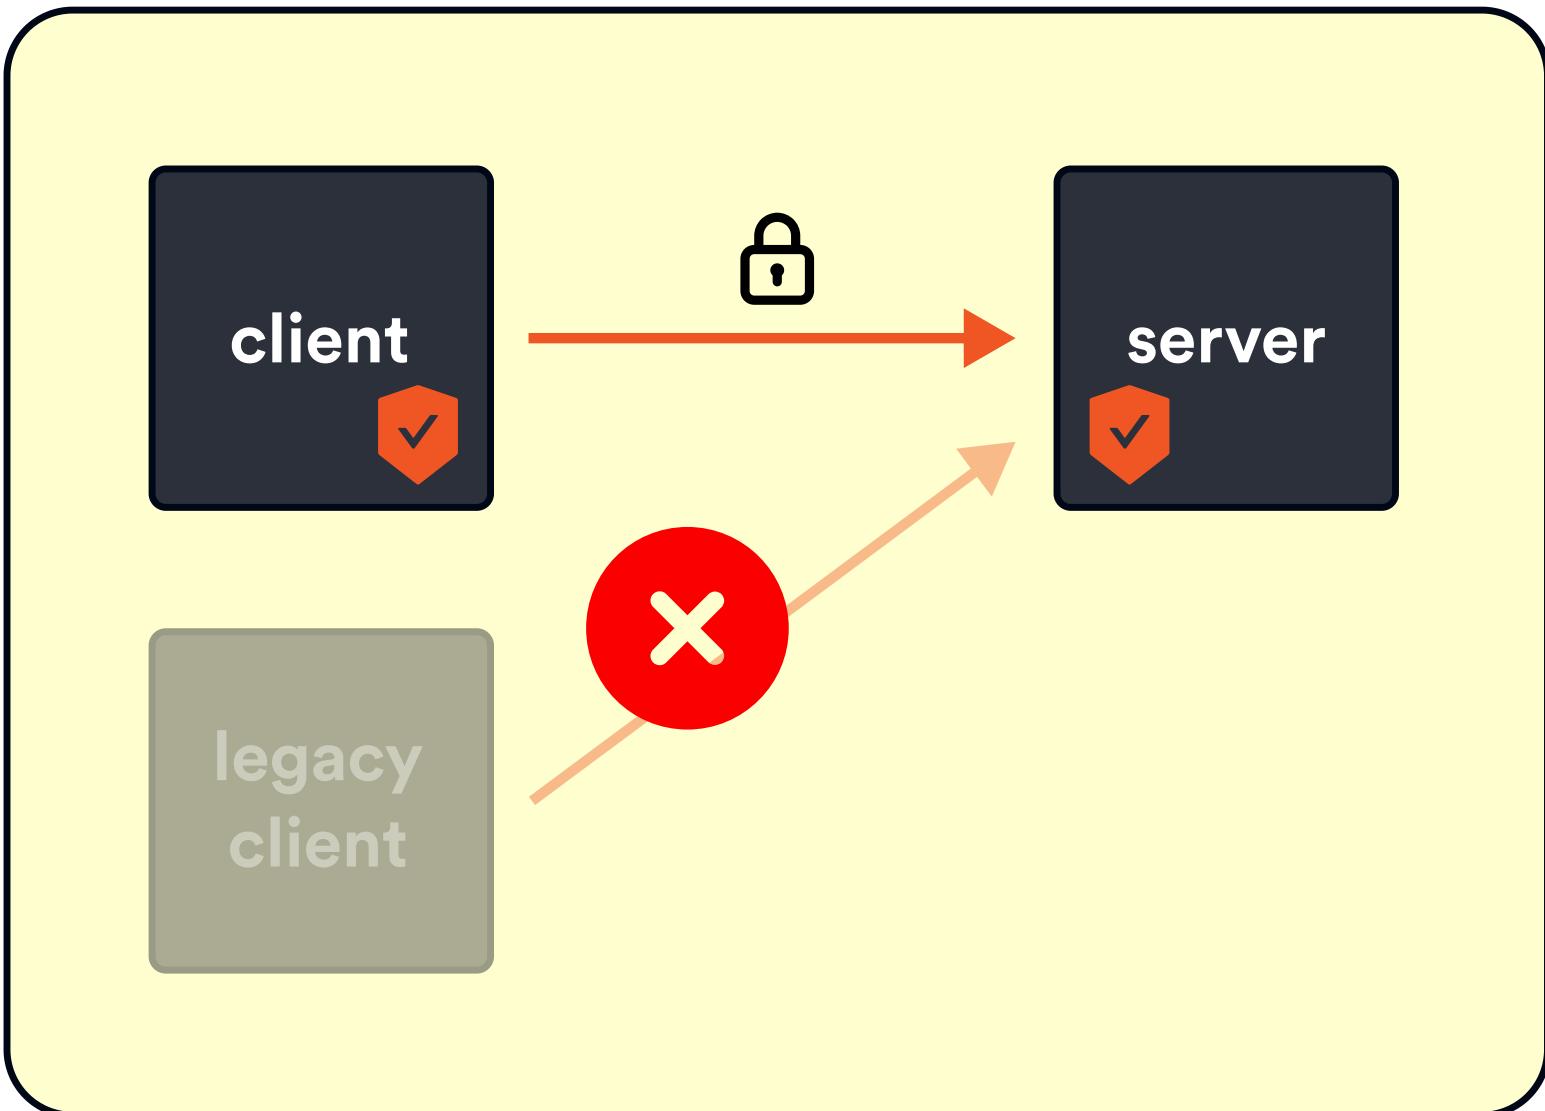


Permissive





Strict





Peer authentication

SERVICE-TO-SERVICE COMMUNICATION

- Controls communication between services
 - PERMISSIVE (default)
 - STRICT
- Mesh, namespace, workload, and port level



Peer authentication

NAMESPACE LEVEL

```
apiVersion: security.istio.io/v1beta1
kind: PeerAuthentication
metadata:
  name: default
  namespace: foo
spec:
  mtls:
    mode: STRICT
```



Peer authentication

WORKLOAD LEVEL

```
apiVersion: security.istio.io/v1beta1
kind: PeerAuthentication
metadata:
  name: default
  namespace: foo
spec:
  selector:
    matchLabels:
      app: prod
  mtls:
    mode: STRICT
```



Peer authentication

PORT LEVEL

```
apiVersion: security.istio.io/v1beta1
kind: PeerAuthentication
metadata:
  name: default
  namespace: foo
spec:
  mtls:
    mode: STRICT
  portLevelMtls:
    5000:
      mode: DISABLE
```



What about users?



Request authentication

USER AUTHENTICATION

- Uses JWT tokens
- Mesh/namespace/workload scope
- Also at ingress level:
 - forwardOriginalToken

```
apiVersion: security.istio.io/v1beta1
kind: RequestAuthentication
metadata:
  name: httpbin
  namespace: default
spec:
  selector:
    matchLabels:
      app: httpbin
  jwtRules:
    - issuer: "issuer-foo"
      jwksUri: "someuri"
```



JWT authentication filter

- Authn enforced by the filter
- Doesn't deny requests without JWT tokens (*allowMissing*)
 - Used together with AuthorizationPolicy

```
name: envoy.filters.http.jwt_authn
typedConfig:
  providers:
    origins-0:
      issuer: testing@secure.istio.io
      localJwks:
        inlineString: '...'
      payloadInMetadata: testing@secure.istio.io
    rules:
      - match:
          prefix: "/"
        requires:
          requiresAny:
            requirements:
              - providerName: origins-0
              - allowMissing: {}
```



JWT authentication filter

WHEN ARE THE REQUESTS APPROVED/DENIED

DENIED

- Mismatching issuers
- Token expired
- Invalid audience (if provided)
- Invalid signature

APPROVED

- Valid JWT
- No JWT
 - Use AuthorizationPolicy



Authorization (authz)

CAN A PRINCIPAL PERFORM AN ACTION?

- Can user **A** send a **GET request** to path **/hello** on service **B**?
- Authn without authz (and vice-versa) is useless
- Control authenticated principals with **AuthorizationPolicy**



Authorization policy

- Make use of identities extracted from:
 - **PeerAuthentication** -> *principals(service/peer)*
 - **RequestAuthentication** → *requestPrincipals(users)*



Example

```
apiVersion: security.istio.io/v1beta1
kind: AuthorizationPolicy
metadata:
  name: require-jwt
  namespace: default
spec:
  selector:
    matchLabels:
      app: prod
  rules:
  - from:
    - source:
        requestPrincipals: ["*"]
```



Authorization policy

RULES "FROM" FIELD

- source identities
- namespaces
- principals
- IP blocks and remote IP blocks

```
rules:  
- from:  
  - source:  
    principals: ["cluster.local/ns/default/sa/workload"]  
  - source:  
    namespaces: ["prod"]  
  - source:  
    requestPrincipals: ["tetrate.io/peterj"]
```



Authorization policy

RULES "TO" FIELD

- hosts
- ports
- methods
- paths

```
rules:  
- from:  
  - ...  
  to:  
    - operation:  
        methods: ["DELETE"]  
        paths: ["/logs*"]  
    - operation:  
        methods: ["GET"]  
        paths: ["/data"]  
    - operation:  
        hosts: ["request.host"]  
        ports: ["3000", "5000"]
```



Authorization policy

"WHEN" FIELD (CONDITIONS)

- Keys and values (or notValues)
- Request attributes:
 - `request.headers`
 - `source.ip, remote.ip`
 - `source.namespace | principal,`
 - ...

```
rules:  
- from:  
to:  
when:  
- key: request.auth.claims[iss]  
values: ["https://accounts.google.com"]  
- key: request.headers[my-header]  
values: ["some-value"]  
- key: source.namespace  
value: ["foo"]  
...
```



Authorization policy

"ACTION" FIELD

- CUSTOM
- DENY
- ALLOW
- AUDIT

```
spec:  
  action: DENY  
  rules:  
    - from:  
      to:  
      when:  
        ...
```



Recap

- Services → PeerAuthentication
- Users → RequestAuthentication
- Access control rules → AuthorizationPolicy
 - From, To, When



LAB

Security

<https://tetratelabs.github.io/istio-0to60/security/>



Istio Traffic Management

CUSTOM RESOURCES (1/2)

- **Virtual service**
 - Configure routing rules for each service
- **Destination rule**
 - Configure how to reach the target endpoint, applied after routing decision has been made



How to route the request?



K8S SERVICE

serviceA.example.cluster.local
app: svcA

K8S SERVICE

serviceB.example.cluster.local
app: svcB

K8S DEPLOYMENT

serviceA
app: svcA

Pod

app: svcA

svcA

Envoy

K8S DEPLOYMENT

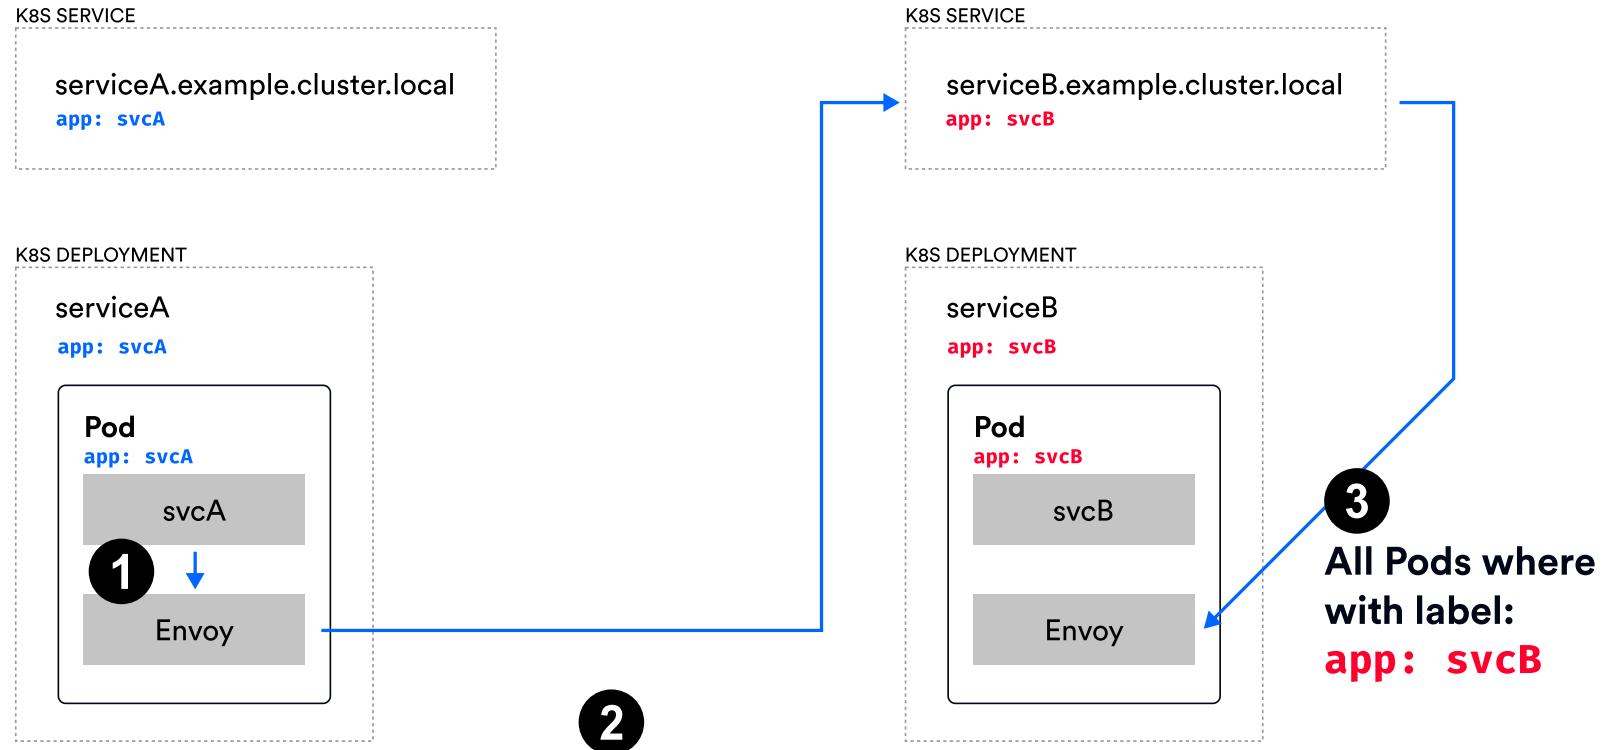
serviceB
app: svcB

Pod

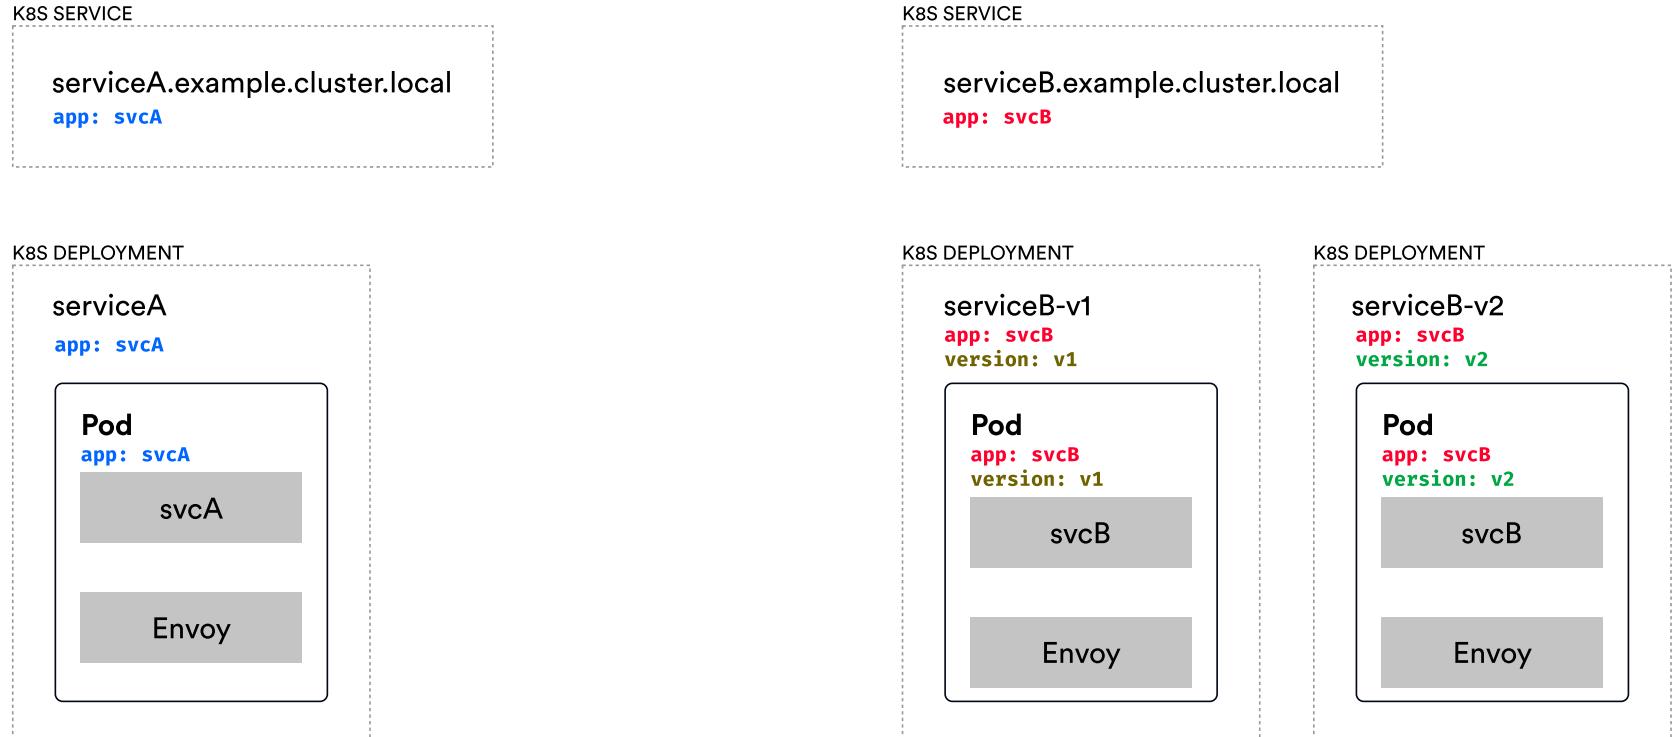
app: svcB

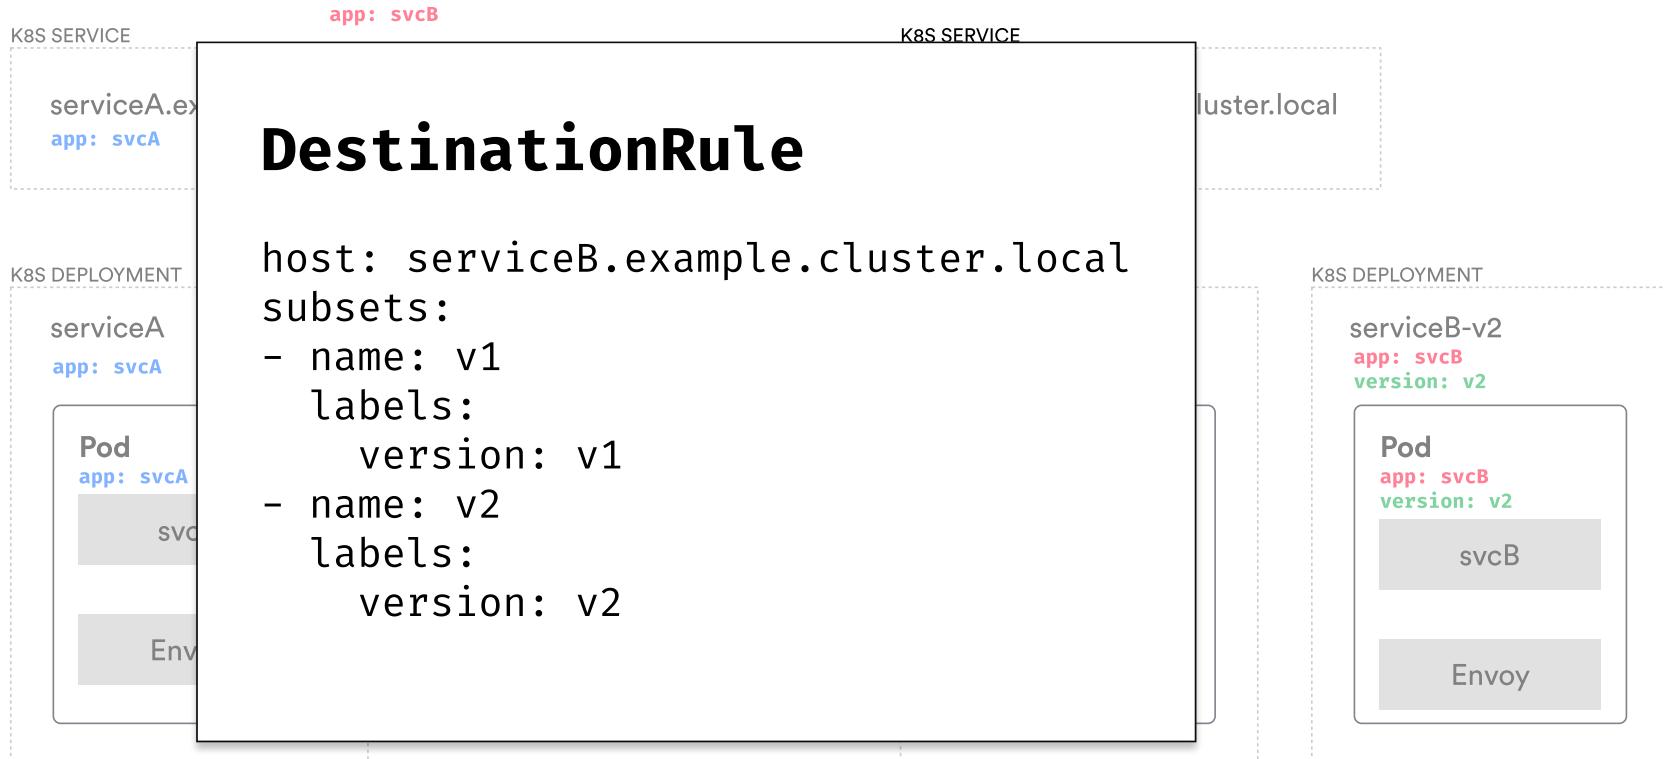
svcB

Envoy



<http://serviceb.example.cluster.local>







K8S SERVICE

app: svcB

VirtualService

```
hosts:  
  - serviceB.example.cluster.local  
http:  
  - route:  
    - destination:  
      host: serviceB.example ...  
      subset: v1  
      weight: 70  
    - destination:  
      host: serviceB.example ...  
      subset: v2  
      weight: 30
```

K8S SERVICE

DestinationRule

```
host: serviceB.example.cluster.local  
subsets:  
  - name: v1  
    labels:  
      version: v1  
  - name: v2  
    labels:  
      version: v2
```



K8S SERVICE

serviceA.example.cluster.local
app: svcA

K8S DEPLOYMENT

serviceA
app: svcA

Pod

app: svcA

svcA

Envoy

DestinationRule & virtualService



70% to subset v1
30% to subset v2

K8S SERVICE

serviceB.example.cluster.local
app: svcB

K8S DEPLOYMENT

serviceB-v1
app: svcB
version: v1

Pod

app: svcB
version: v1

svcB

Envoy

K8S DEPLOYMENT

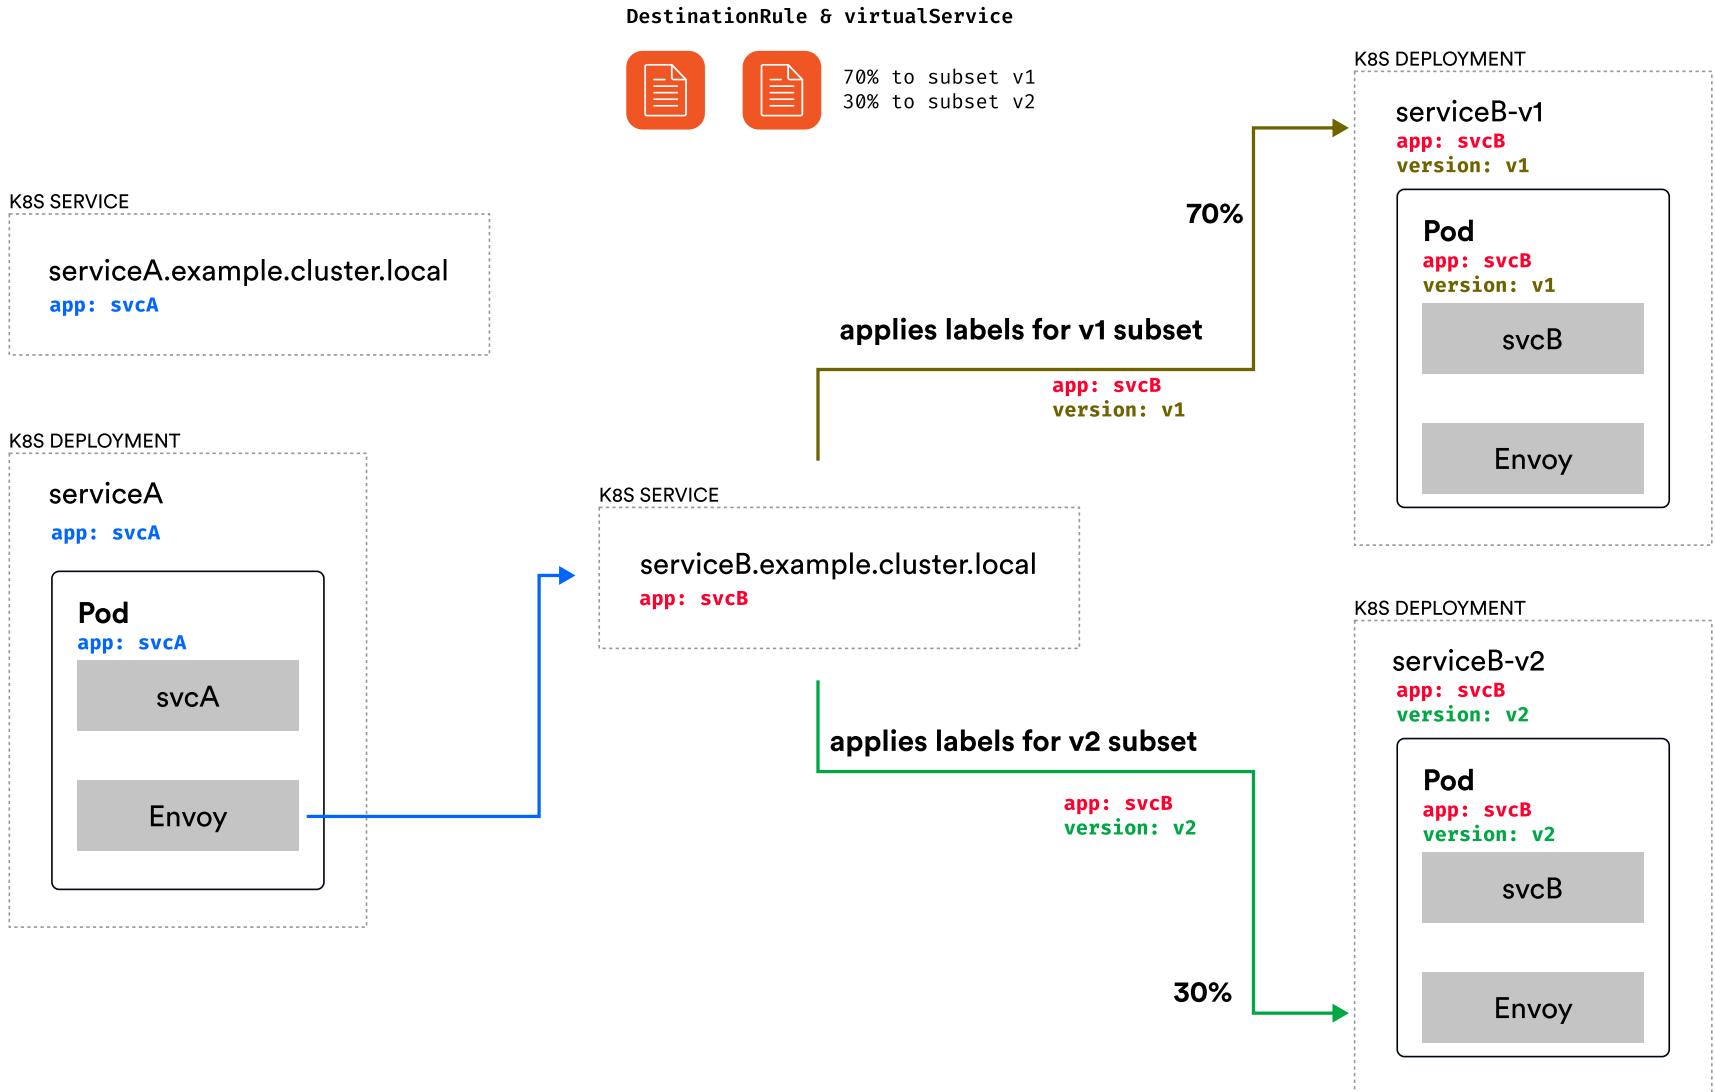
serviceB-v2
app: svcB
version: v2

Pod

app: svcB
version: v2

svcB

Envoy





DestinationRule

POLICIES APPLIED TO TRAFFIC FOR A SPECIFIC SERVICE

- Subsets = represent different service versions
- Traffic policies:
 - Load balancer settings (ROUND_ROBIN, LEAST_CONN, RANDOM, PASSTHROUGH)
 - Connection pool settings (TCP and HTTP)
 - Outlier detection
 - TLS



Connection pool settings

CONTROL THE VOLUME OF CONNECTIONS

- Applied to TCP and/or HTTP connections
- Timeouts
- Max connections/requests
- Max retries



Outlier detection

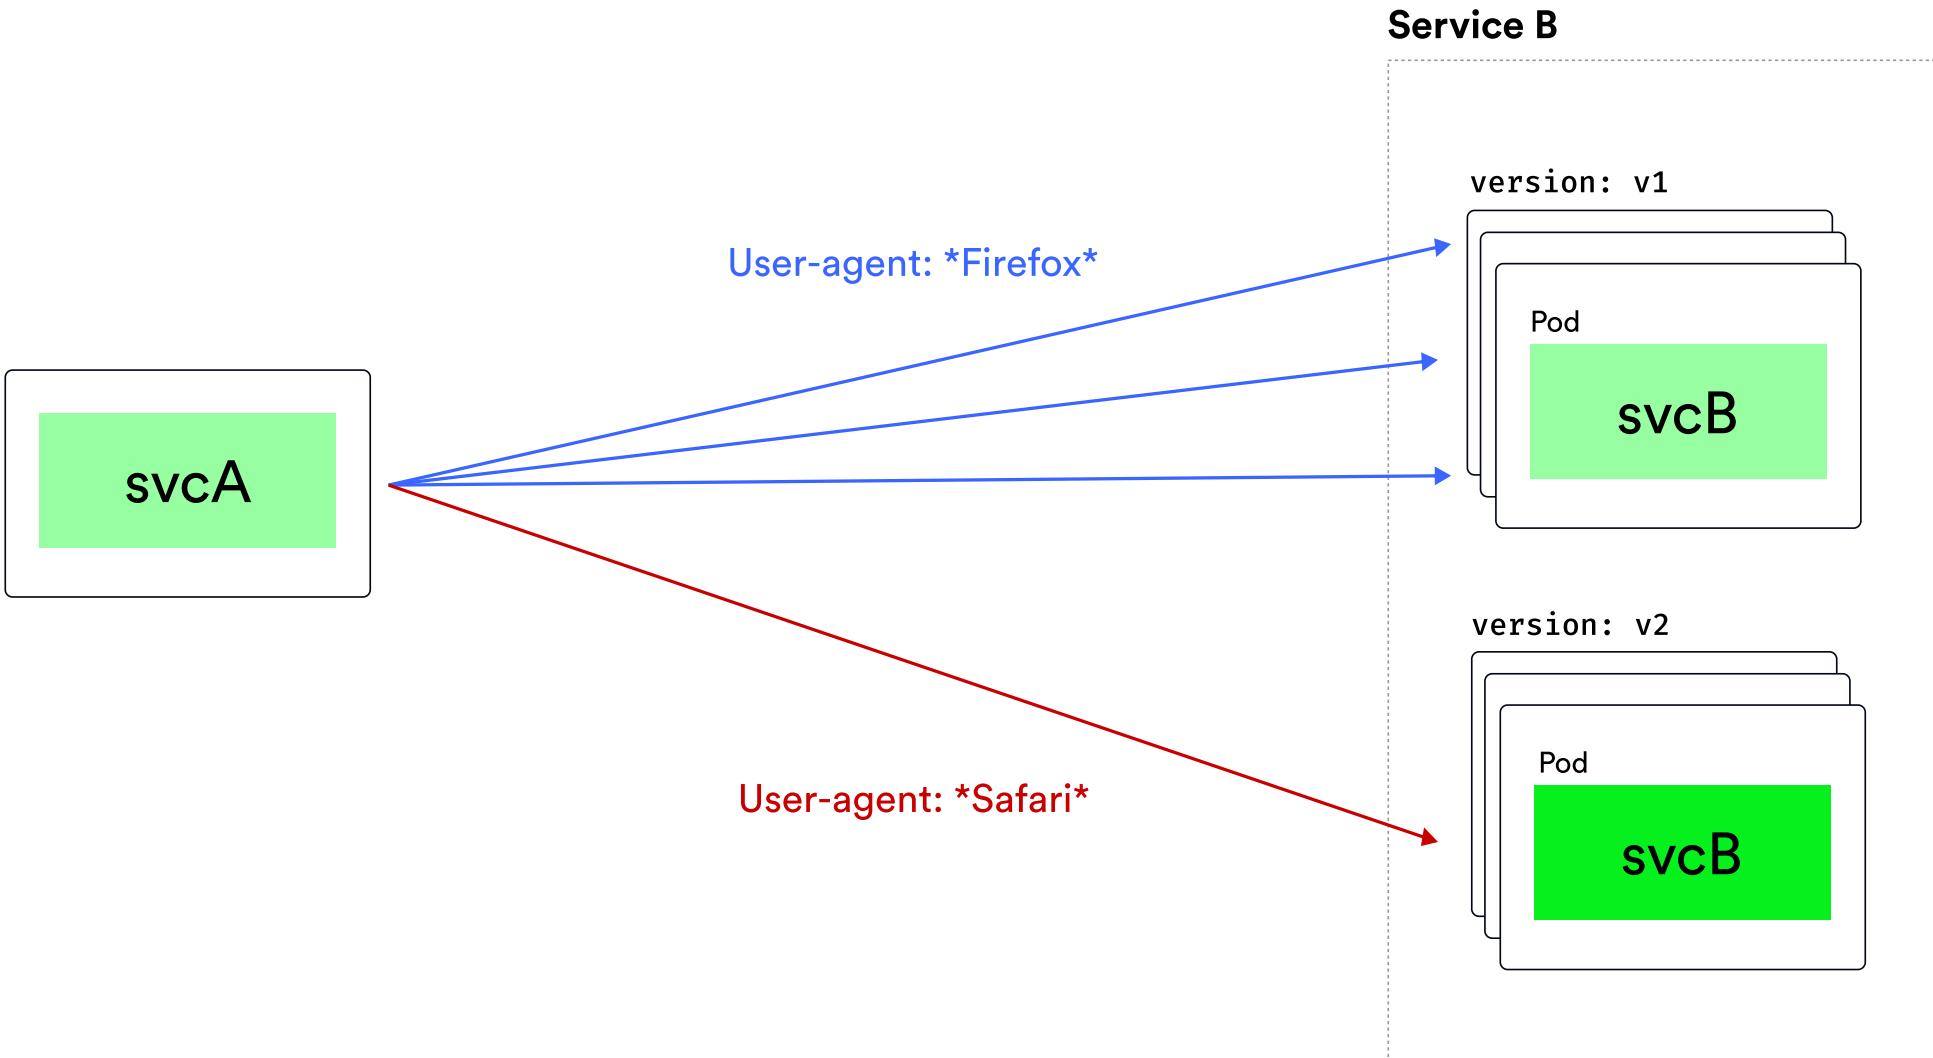
HOW TO EJECT UNHEALTHY HOSTS

- When to eject unhealthy hosts?
 - e.g. `consecutive5xxErrors`, `consecutiveGatewayErrors`
- How long to eject them for?
 - `baseEjectionTime`
- How many hosts can be ejected?
 - `maxEjectionPercent`
- When to enable ejection?
 - `minHealthyPercent`



VirtualService

- Routing rules (TCP, HTTP, non-terminated TLS/HTTPS traffic)
 - Match & route
 - Redirect
 - Rewrite
 - Mirroring
 - Cors, Timeouts, retries, and fault injection
- Header manipulation





Match on headers

```
hosts:
  - svcB.example.cluster.local
http:
- match:
  - headers:
    user-agent:
      regex: ".*Firefox.*"
route:
- destination:
  host: svcB.example.cluster.local
  subset: v1
- route:
  - destination:
    host: svcB.example.cluster.local
    subset: v2
```



AND semantics

```
hosts:
  - svcB.example.cluster.local
http:
- match:
  - headers:
    x-debug:
      exact: dev
  uri:
    prefix: /api/debug
route:
- destination:
  host: svcB.example.cluster.local
  subset: v1
- route:
  - destination:
    host: svcB.example.cluster.local
    subset: v2
```



OR semantics

```
hosts:
  - svcB.example.cluster.local
http:
- match:
  - headers:
    x-debug:
      exact: dev
  - uri:
    prefix: /api/debug
route:
- destination:
  host: svcB.example.cluster.local
  subset: v1
- route:
  - destination:
    host: svcB.example.cluster.local
    subset: v2
```



Timeout, retries

```
hosts:
  - svcB.example.cluster.local
http:
- route:
  - destination:
    host: svcB.example.cluster.local
    subset: v1
    weight: 30
    timeout: 5s
  - destination:
    host: svcB.example.cluster.local
    subset: v2
    weight: 70
    timeout: 0.5s
retries:
  attempts: 3
  perTryTimeout: 2s
  retryOn: connect-failure
```



LAB

Traffic shifting

<https://tetratelabs.github.io/istio-0to60/traffic-shifting/>